

LAB – OpenStack Installation

Hands on experience with OpenStack installation using JuJu

References:


- JuJu openstack documentation
<https://docs.openstack.org/project-deploy-guide/charm-deployment-guide/latest/install-openstack.html>

VM Configuration

- Labs from now on need to be carried out in groups of 4 or 5
- Each VM will have a different role in the system resulting in the following overall architecture

The architecture manager is a node that exposes to administrators a console to manage all the nodes of the platform

OpenStack controller (configured also as Network Node) and compute nodes will require with access to the public network (access network does not exist, we will emulate it)



Network Configuration

- Each VM has only one physical network and it is connected to only one physical network
- This physical network will be the private network exploited by the machines to communicate to implement their functions
- The public physical machines will be emulated on top of the physical network using VLAN
- To this aim we need to modify the network configuration of each machine
- By default a VM has a pre-defined network configuration that is injected by the cloud platform into the VM
- We need to modify the network configuration so each VM has a static IP address and create at bootstrap a virtual network on top of the physical one using VLAN

Netplan

Netplan is the tool adopted into Ubuntu 18.04 to manage the network configuration

network:

version: 2

renderer: networkd

ethernets:

eth0:

addresses:

- 172.16.1.59/16

gateway4: 172.16.0.1

nameservers:

addresses:

- 8.8.8.8

vlan:

vlan2:

id: 2

link: eth0

dhcp4: no

dhcp6: no

Vlan2 is a virtual network (manually created) to emulate the public network on top of the physical one

Apply the configuration to `/etc/netplan/01-netcfg.yaml` and run **netplan apply**

JuJu – Installation

- Juju is an open source application tool developed by Canonical to facilitate quick deployment and configuration of public and private cloud services along bare metal servers or VMs
- Connect to the manager and install juju

```
sudo snap install juju --classic
```

```
root@F0BVVYAMFIT3QIF:~# sudo snap install juju --classic
2020-02-21T16:02:32Z INFO Waiting for restart...
juju 2.7.2 from Canonical✓ installed
Channel latest/stable for juju is closed; temporarily forwarding to stable.
root@F0BVVYAMFIT3QIF:~# █
```

JuJu – Configuration

- Add new machines, since your VMs are existing machines that are like bare metal machines we perform manual installation, which is used every time machines are not managed by any cloud platform
<https://jaas.ai/docs/manual-cloud>
- The steps to be performed are the following:
 - Add cloud
 - Bootstrap
 - Create a controller
 - Check status
 - Add machines

Handle SSH keys

- Manual cloud installation require that the manager can connect to all the machines managed by the manager via ssh
- To this aim, ssh connection must be performed without password
- Create a set of ssh rsa key on the manager

ssh-keygen

- Add the key to all the machines (including the manager the host you are connected to!)

ssh-copy-id -i ~/.ssh/id_rsa.pub root@IP

```
root@F0BVVYAMFIT3QIF:~# ssh-keygen
Generating public/private rsa key pair.

Enter file in which to save the key (/root/.ssh/id_rsa): Enter passphrase (empty for no passphrase):
Enter same passphrase again:
Your identification has been saved in /root/.ssh/id_rsa.
Your public key has been saved in /root/.ssh/id_rsa.pub.
The key fingerprint is:
SHA256:nH1J0ffjR0ny4ruRiGtoXaFyn+GU0szH9/00HY0AY8 root@F0BVVYAMFIT3QIF
The key's randomart image is:
----[RSA 2048]-----
      ..
      o  +=..
      = *  .E+=..
      . o S o o. =o
      = o * o. =o=
      . o = 000*+
      + =... .+
      B+ .
-----[SHA256]-----
root@F0BVVYAMFIT3QIF:~#
```

Add cloud

- Create a new cloud of manual type on the manager
`juju add-cloud`
- **As IP of the controller select the Architecture manager** (the local machine you are connected to not the controller)

```
root@SNH0YM5GWPGE2L:~# juju add-cloud
Cloud Types
  lxd
  maas
  manual
  openstack
  vsphere

Select cloud type: manual

Enter a name for your manual cloud: mycloud

Enter the ssh connection string for controller, username@<hostname or IP> or <hostname or IP>: root@172.16.0.128

This operation can be applied to both a copy on this client and to the one on a controller.
No current controller was detected and there are no registered controllers on this client: either bootstrap one or register one.
Do you ONLY want to add cloud "mycloud" to this client? (Y/n): Y

Cloud "mycloud" successfully added to your local client.
root@SNH0YM5GWPGE2L:~#
```


Bootstrap

- Bootstrap the controller

It can take a while!

```
juju bootstrap mycloud manual-controller
```

```
root@SNH0YM5GWPGE2L:~# juju bootstrap mycloud manual-controller
Creating Juju controller "manual-controller" on mycloud/default
Looking for packaged Juju agent version 2.7.0 for amd64
Installing Juju agent on bootstrap instance
Fetching Juju GUI 2.15.0
Running machine configuration script...

Bootstrap agent now started
Contacting Juju controller at 172.16.0.128 to verify accessibility...

Bootstrap complete, controller "manual-controller" now is available
Controller machines are in the "controller" model
Initial model "default" added
```

```
juju status
```

```
root@SNH0YM5GWPGE2L:~# juju status
Model      Controller      Cloud/Region      Version  SLA          Timestamp
default    manual-controller  mycloud/default  2.7.0    unsupported   08:29:29Z

Model "admin/default" is empty.
root@SNH0YM5GWPGE2L:~#
```

Add machine

- Add all the remaining machines (controller, and compute nodes)

juju add-machine ssh:root@IP

It can take a while!

```
root@SNH0YM5GWPGME2L:~# juju add-machine ssh:root@172.16.0.110
ubuntu:x:1001:
created machine 1
root@SNH0YM5GWPGME2L:~# juju add-machine ssh:root@172.16.0.111
ubuntu:x:1001:
created machine 2
root@SNH0YM5GWPGME2L:~# juju status
```

Model	Controller	Cloud/Region	Version	SLA	Timestamp
default	manual-controller	mycloud/default	2.7.0	unsupported	08:36:28Z

Machine	State	DNS	Inst id	Series	AZ	Message
0	started	172.16.0.109	manual:172.16.0.109	bionic		Manually provisioned machine
1	started	172.16.0.110	manual:172.16.0.110	bionic		Manually provisioned machine
2	started	172.16.0.111	manual:172.16.0.111	bionic		Manually provisioned machine

Configure the network for containers

- Some components will be deployed as container
 - Some module is developed to run as the only software installed on a machine, those has to be installed on a single machine, a container in this case, otherwise they will conflict each other (e.g. openstack-dashboard and nova-controller cannot run on the same machine)
- In order to have containers running on different machine to communicate each other we have to change the juju network configuration for container in order to use FAN, which is a software component that creates an overlay network across different containers running on different machine

```
juju model-config fan-config=172.16.0.0/16=252.0.0.0/8
```

```
juju model-config container-networking-method=fan
```

```
juju model-config | egrep 'fan-config|container-networking-method'
```

```
root@F0BVYAMFIT3QIF:~# juju model-config | egrep 'fan-config|container-networking-method'
container-networking-method  model    fan
fan-config                   model    172.16.0.0/16=252.0.0.0/8
```

252.0.0.0/8 is the overlay network address

Use already provisioned machines

- If you want to reuse a machine that has been provisioned previously (using juju) you need to run the following steps:

```
sudo rm -rf /var/lib/juju
```

```
sudo rm -rf /lib/systemd/system/juju*
```

```
sudo rm -rf /run/systemd/units/invocation:juju*
```

```
sudo rm -rf /etc/systemd/system/juju*
```

Execute this only in case of errors, and
you need to perform again one of the
previous operations!

- If you need to reset the architecture manager run:

```
rm -rf /root/.local/share/juju
```

Storage

- Each VM has two hard drives attached, one for the OS (/dev/sda) and another additional hard drive for storage (/dev/sdb)
- Controller node will use the additional hard drive to install the platform software, compute nodes will use it as storage for VMs
- Using the hard drive on the **controller** node requires some manual steps (disks on the compute nodes will be managed automatically by the platform):
 - Create a new EXT4 partition
 - `fdisk /dev/sdb` (option n to create a new partition and then option w to write changes)
 - Format the partition
 - `mkfs.ext4 /dev/sdb1`
 - Mount the partition on the folder where JuJu will install the software running on the controller
 - `mkdir -p /var/lib/lxd/storage-pools/default/`
 - `mount /dev/sdb1 /var/lib/lxd/storage-pools/default/`
 - `rm -rf /var/lib/lxd/storage-pools/default/lost+found`
 - Add the following line to the file /etc/fstab to make the change permanent
 - `/dev/sdb1 /var/lib/lxd/storage-pools/default ext4 errors=remount-ro 0 1`

Storage

```
root@NACH6ARD4Y7ZDQI:~# fdisk /dev/sdb

Welcome to fdisk (util-linux 2.31.1).
Changes will remain in memory only, until you decide to write them.
Be careful before using the write command.

Device does not contain a recognized partition table.
Created a new DOS disklabel with disk identifier 0x3bde43a3.

Command (m for help): n
Partition type
   p   primary (0 primary, 0 extended, 4 free)
   e   extended (container for logical partitions)
Select (default p): p
Partition number (1-4, default 1): 1
First sector (2048-83886079, default 2048):
Last sector, +sectors or +size[K,M,G,T,P] (2048-83886079, default 83886079):

Created a new partition 1 of type 'Linux' and of size 40 GiB.

Command (m for help): w
The partition table has been altered.
Calling ioctl() to re-read partition table.
Syncing disks.

root@NACH6ARD4Y7ZDQI:~# █
```

```
root@NACH6ARD4Y7ZDQI:~# mkfs.ext4 /dev/sdb1
mke2fs 1.44.1 (24-Mar-2018)
Discarding device blocks: done
Creating filesystem with 10485504 4k blocks and 2621440 inodes
Filesystem UUID: 0918352d-0398-47a3-b505-5872b4731b24
Superblock backups stored on blocks:
    32768, 98304, 163840, 229376, 294912, 819200, 884736, 1605632, 2654208,
    4096000, 7962624

Allocating group tables: done
Writing inode tables: done
Creating journal (65536 blocks): done
Writing superblocks and filesystem accounting information: done

root@NACH6ARD4Y7ZDQI:~# mount /dev/sdb1 /var/lib/lxd/storage-pools/default/containers
mount: /var/lib/lxd/storage-pools/default/containers: mount point does not exist.
root@NACH6ARD4Y7ZDQI:~# mkdir -p /var/lib/lxd/storage-pools/default/containers
root@NACH6ARD4Y7ZDQI:~# mount /dev/sdb1 /var/lib/lxd/storage-pools/default/containers
mount: /var/lib/lxd/storage-pools/default/containers: mount point does not exist.
root@NACH6ARD4Y7ZDQI:~# █
```

Manual OpenStack Installation

- In the following steps we will install OpenStack manually component by component
- For each component juju require a configuration, the number of instances to deploy and the machines to which deploy the instances
- Configuration is specified though a configuration file yaml
- The syntax of the command in general is the following
`juju deploy --to ID_HOST --config config.yaml NAME_COMPONENT`
- To deploy an additional instance of the component
`juju add-unit --to ID_HOST NAME_COMPONENT`

Relations

- For each deployment juju allows to define relations
- A relation is a dependency between two components:
 - A configuration dependency, i.e. a parameter of a module is a configuration parameter of the other (e.g. the IP address of RabbitMQ is required by other modules like Nova)
 - An installation dependency, i.e. the installation of a module triggers automatically the installation of other modules
- The type of relation is specified by the module programmer
- A relation can be declared through this command:
juju add-relation COMPONENT1 COMPONENT2

Nova Compute

- Deploy nova compute on the compute nodes {1,2,3}
- compute.yaml

nova-compute:

enable-live-migration: True

enable-resize: True

migration-auth-type: ssh

virt-type: qemu

openstack-origin: cloud:bionic-train

Qemu (libvirt) will be the hypervisor



Version of the component

```
juju deploy --to 1 --config compute.yaml nova-compute
```

```
juju add-unit --to 2 nova-compute
```

Nova Compute

```
root@F0BVVYAMFIT3QIF:~/config# juju deploy --to 1 --config compute.yaml nova-compute
Located charm "cs:nova-compute-312".
Deploying charm "cs:nova-compute-312".
root@F0BVVYAMFIT3QIF:~/config# juju add-unit --to 2 nova-compute

root@F0BVVYAMFIT3QIF:~/config#
root@F0BVVYAMFIT3QIF:~/config# juju status
Model      Controller      Cloud/Region  Version  SLA          Timestamp
default    manual-controller  mycloud/default  2.7.2    unsupported   17:01:56Z

App                Version  Status  Scale  Charm          Store        Rev  OS   Notes
nova-compute              waiting    1/2   nova-compute   jujucharms   312  ubuntu

Unit              Workload  Agent    Machine  Public address  Ports  Message
nova-compute/0*   maintenance  executing    1        172.16.0.246    (install) installing charm software
nova-compute/1    waiting    allocating    2        172.16.0.248    agent initializing

Machine  State  DNS           Inst id              Series  AZ  Message
0        started  172.16.0.242  manual:172.16.0.242  bionic             Manually provisioned machine
1        started  172.16.0.246  manual:172.16.0.246  bionic             Manually provisioned machine
2        started  172.16.0.248  manual:172.16.0.248  bionic             Manually provisioned machine
```

Neutron - Controller

- Neutron is composed of different components, running both on the controller and on the compute nodes
- Let's start with the components running in the controller
- neutron.yaml

neutron-gateway:

```
bridge-mappings:      physnet1:br-ex
data-port:            br-ex:vlan2
openstack-origin: cloud:bionic-train
```

The port to
the external network

neutron-api:

```
flat-network-providers: physnet1
overlay-network-type: gre
neutron-security-groups: True
openstack-origin: cloud:bionic-train
```

Physnet1 won't
adopt any
virtualization
mechanism (flat
network)

Type of mechanism
to create virtual
networks over the
same physical
network (eth0 in this
case)

Neutron - Controller

- Deploy the components

```
juju deploy --to 0 --config neutron.yaml neutron-gateway
```

```
juju deploy --to lxd:0 --config neutron.yaml neutron-api
```

Neutron-api is deployed inside a container. Again, this component as others are developed starting from the assumption that they are installed on an entire system available only for them, if multiple components are installed on the same machine conflicts might arise

Neutron-gateway manages the connection with external networks. It must have direct access to the external interface, so it cannot be deployed in a container

Neutron - Controller

```
root@E8C4KA0NA4GYMWT:~/config# juju status
Model      Controller      Cloud/Region  Version  SLA          Timestamp
default    manual-controller  mycloud/default  2.7.3    unsupported   17:53:42Z

App                Version  Status      Scale  Charm          Store      Rev  OS    Notes
neutron-api        2.7.3    waiting     0/1    neutron-api    jujucharms  283  ubuntu
neutron-gateway    2.7.3    maintenance  1      neutron-gateway  jujucharms  279  ubuntu
nova-compute       2.7.3    maintenance  2      nova-compute    jujucharms  312  ubuntu

Unit                Workload  Agent      Machine  Public address  Ports  Message
neutron-api/0       waiting   allocating  0/lxd/0                  waiting for machine
neutron-gateway/0*  maintenance  executing  0         172.16.1.114    (install) Installing apt packages
nova-compute/0*     maintenance  executing  1         172.16.1.115    (install) Installing apt packages
nova-compute/1      maintenance  executing  2         172.16.1.117    (install) Installing apt packages

Machine  State  DNS           Inst id          Series  AZ  Message
0        started  172.16.1.114  manual:172.16.1.114  bionic  AZ  Manually provisioned machine
0/lxd/0  pending
1        started  172.16.1.115  manual:172.16.1.115  bionic  AZ  Manually provisioned machine
2        started  172.16.1.117  manual:172.16.1.117  bionic  AZ  Manually provisioned machine
3        started  172.16.1.118  manual:172.16.1.118  bionic  AZ  Manually provisioned machine
```

Neutron - Compute


- Each compute node has a neutron instance that control openvswitch
- neutron.yaml

neutron-openvswitch:

bridge-mappings:	physnet1:br-ex
data-port:	br-ex:vlan2
firewall-driver:	openvswitch

juju deploy --config neutron.yaml neutron-openvswitch

Deploy openvswitch and let juju select the nodes based on dependencies (it will be deployed on compute nodes)



Neutron - Relations

```
juju add-relation neutron-api:neutron-plugin-api  
neutron-gateway:neutron-plugin-api
```

```
juju add-relation neutron-api:neutron-plugin-api  
neutron-openvswitch:neutron-plugin-api
```

```
juju add-relation neutron-openvswitch:neutron-  
plugin nova-compute:neutron-plugin
```

MySQL

- OpenStack requires a shared database, so we have to deploy a MySQL on the controller node {0} inside a container
- mysql.yaml

mysql:

max-connections: 20000

source: cloud:bionic-train

juju deploy --to lxd:0 --config mysql.yaml percona-cluster mysql

juju add-relation neutron-api:shared-db mysql:shared-db

MySQL

```
root@E8C4KA0NA4GYMWT:~/config# juju status
Model      Controller      Cloud/Region  Version  SLA          Timestamp
default    manual-controller mycloud/default 2.7.3    unsupported   18:09:51Z

App                Version  Status      Scale  Charm          Store        Rev  OS   Notes
mysql              15.0.1   blocked     1      percona-cluster jujucharms   284  ubuntu
neutron-api        15.0.1   blocked     1      neutron-api    jujucharms   283  ubuntu
neutron-gateway    15.0.1   blocked     1      neutron-gateway jujucharms   279  ubuntu
neutron-openvswitch 15.0.1   blocked     2      neutron-openvswitch jujucharms 273  ubuntu
nova-compute       20.0.1   blocked     2      nova-compute   jujucharms   312  ubuntu

Unit                Workload  Agent    Machine  Public address  Ports  Message
mysql/0*            maintenance executing 0/lxd/1  252.1.114.82    (install) installing charm software
neutron-api/0*      maintenance executing 0/lxd/0  252.1.114.25    (install) Installing apt packages
neutron-gateway/0*  blocked   idle     0         172.16.1.114    Missing relations: messaging
nova-compute/0*     blocked   idle     1         172.16.1.115    Missing relations: image, messaging
neutron-openvswitch/0* blocked   idle     1         172.16.1.115    Missing relations: messaging
nova-compute/1      blocked   idle     2         172.16.1.117    Missing relations: image, messaging
neutron-openvswitch/1 blocked   idle     2         172.16.1.117    Missing relations: messaging

Machine  State  DNS                Inst id              Series  AZ  Message
0        started 172.16.1.114      manual:172.16.1.114  bionic  Manualy provisioned machine
0/lxd/0  started 252.1.114.25      juju-6dd12f-0-lxd-0 bionic  Container started
0/lxd/1  started 252.1.114.82      juju-6dd12f-0-lxd-1 bionic  Container started
1        started 172.16.1.115      manual:172.16.1.115  bionic  Manualy provisioned machine
2        started 172.16.1.117      manual:172.16.1.117  bionic  Manualy provisioned machine
3        started 172.16.1.118      manual:172.16.1.118  bionic  Manualy provisioned machine
```

Keystone

- Deploy keystone on the controller node {0} in a container
- keystone.yaml

keystone:

admin-password: openstack

openstack-origin: cloud:bionic-train

juju deploy --to lxd:0 --config keystone.yaml keystone

juju add-relation keystone:shared-db mysql:shared-db

juju add-relation keystone:identity-service neutron-api:identity-service

RabbitMQ

- Deploy RabbitMQ on the controller node {0} in a container

```
juju deploy --to lxd:0 rabbitmq-server
```

```
juju add-relation rabbitmq-server:amqp neutron-  
api:amqp
```

```
juju add-relation rabbitmq-server:amqp neutron-  
openvswitch:amqp
```

```
juju add-relation rabbitmq-server:amqp nova-  
compute:amqp
```

```
juju add-relation rabbitmq-server:amqp neutron-  
gateway:amqp
```

Nova Controller

- Deploy nova controller on the controller node {0} in a container
- controller.yaml

nova-cloud-controller:

network-manager: Neutron

console-access-protocol: novnc

console-proxy-ip: 172.16.3.26

openstack-origin: cloud:bionic-train

IP of the controller that will exposes the dashboard

Protocol for the console connection

juju deploy --to lxd:0 --config controller.yaml nova-cloud-controller

```
root@F0BVVYAMFIT3QIF:~/config# juju deploy --to lxd:0 --config controller.yaml nova-cloud-controller
Located charm "cs:nova-cloud-controller-341".
Deploying charm "cs:nova-cloud-controller-341".
```

Nova Controller

- Set the following dependencies

```
juju add-relation nova-cloud-controller:shared-db  
mysql:shared-db
```

```
juju add-relation nova-cloud-controller:identity-service  
keystone:identity-service
```

```
juju add-relation nova-cloud-controller:amqp rabbitmq-  
server:amqp
```

```
juju add-relation nova-cloud-controller:quantum-network-  
service neutron-gateway:quantum-network-service
```

```
juju add-relation nova-cloud-controller:neutron-api  
neutron-api:neutron-api
```

```
juju add-relation nova-cloud-controller:cloud-compute nova-  
compute:cloud-compute
```

Placement

- Placement is a recent component that implement the nova scheduling functionalities
- Deploy a placement instance on the controller node {0} in a container

```
placement.yaml
```

```
placement:
```

```
    openstack-origin: cloud:bionic-train
```

```
juju deploy --to lxd:0 --config placement.yaml placement
```

```
juju add-relation placement:shared-db mysql:shared-db
```

```
juju add-relation placement:identity-service keystone:identity-  
service
```

```
juju add-relation placement:placement nova-cloud-  
controller:placement
```

Horizon - Dashboard

- Deploy the dashboard on the controller node {0}. Deploy on the physical node in order to ensure that the dashboard is reachable from the network through the IP address on the controller node


dashboard.yaml

openstack-dashboard:

openstack-origin: cloud:bionic-train

os-public-hostname: 172.16.3.26

IP of the controller that will exposes the dashboard



juju deploy --to 0 --config dashboard.yaml openstack-dashboard

**juju add-relation openstack-dashboard:identity-service
keystone:identity-service**

Glance

- Deploy glance on the controller node {0} in a container

```
glance.yaml
```

```
glance:
```

```
    openstack-origin: cloud:bionic-train
```

```
juju deploy --to lxd:0 --config glance.yaml glance
```

```
juju add-relation glance:image-service nova-cloud-controller:image-service
```

```
juju add-relation glance:image-service nova-compute:image-service
```

```
juju add-relation glance:shared-db mysql:shared-db
```

```
juju add-relation glance:identity-service keystone:identity-service
```

```
juju add-relation glance:amqp rabbitmq-server:amqp
```


Cinder

- Deploy cinder on the controller node {0} in a container
- cinder.yaml

cinder:

glance-api-version: 2

block-device: None

openstack-origin: cloud:bionic-train

```
juju deploy --to lxd:0 --config cinder.yaml cinder
```

```
juju add-relation cinder:cinder-volume-service nova-cloud-  
controller:cinder-volume-service
```

```
juju add-relation cinder:shared-db mysql:shared-db
```

```
juju add-relation cinder:identity-service keystone:identity-service
```

```
juju add-relation cinder:amqp rabbitmq-server:amqp
```

```
juju add-relation cinder:image-service glance:image-service
```

Ceph

- Ceph is a specific module to create a distributed file system across different machines, to create a cloud file system
- Since OpenStack requires a file system to work (to store the images of the OS and the virtual disks of the VMs) we also deploy ceph
- We will cover ceph specifically later on, so in this lab you have to take it as a black box
- Ceph is composed of two components (OSD and MON) that are installed on each compute node
- Ceph will use the secondary hard drive of each compute node to create a distributed file system

CephOSD

- Deploy the component CephOSD on each compute node {1,2,3}, the component has to be deployed on the bare metal as it needs access to the secondary drive

ceph-osd:

osd-devices: /dev/sdb

source: cloud:bionic-train

juju deploy --to 1 --config ceph-osd.yaml ceph-osd

juju add-unit --to 2 ceph-osd

juju add-unit --to 3 ceph-osd

CephMON

- Deploy the component CephMON on each compute node {1,2,3}

`ceph-mon:`

`source: cloud:bionic-train`

`juju deploy --to lxd:1 --config ceph-mon.yaml ceph-mon`

`juju add-unit --to lxd:2 ceph-mon`

`juju add-unit --to lxd:3 ceph-mon`

`juju add-relation ceph-mon:osd ceph-osd:mon`

`juju add-relation ceph-mon:client nova-compute:ceph`

`juju add-relation ceph-mon:client glance:ceph`

`juju add-relation cinder ceph-mon`

NTP

- Ceph requires NTP to synchronize the time across different nodes of the distributed file system

```
juju deploy ntp
```

```
juju add-relation ceph-osd:juju-info ntp:juju-info
```

Final configuration

Retrieve the IP address of the dashboard via
`juju status --format=yaml`
`openstack-dashboard | grep public-`
`address | awk '{print $2}'`

```

root@F0BVVYAMFIT3QIF:~/config# juju status
Model      Controller      Cloud/Region  Version  SLA      Timestamp
default    manual-controller mycloud/default 2.7.2    unsupported 13:57:51Z

App                Version  Status  Scale  Charm          Store  Rev  OS   Notes
ceph-mon           14.2.4   active   3      ceph-mon       jujucharms 45   ubuntu
ceph-osd           14.2.4   active   3      ceph-osd       jujucharms 298   ubuntu
cinder             15.0.1   active   1      cinder         jujucharms 300   ubuntu
glance             19.0.0   active   1      glance         jujucharms 294   ubuntu
keystone           16.0.0   active   1      keystone       jujucharms 310   ubuntu
mysql              5.7.20   active   1      percona-cluster jujucharms 284   ubuntu
neutron-api        15.0.1   active   1      neutron-api    jujucharms 283   ubuntu
neutron-gateway    15.0.1   active   1      neutron-gateway jujucharms 279   ubuntu
neutron-openvswitch 15.0.1   active   3      neutron-openvswitch jujucharms 273   ubuntu
nova-cloud-controller 20.0.1   active   1      nova-cloud-controller jujucharms 341   ubuntu
nova-compute       20.0.1   active   3      nova-compute   jujucharms 312   ubuntu
ntp                3.2      active   3      ntp            jujucharms 38    ubuntu
openstack-dashboard 16.0.0   active   1      openstack-dashboard jujucharms 299   ubuntu
placement          2.0.0    active   1      placement      jujucharms 4    ubuntu
rabbitmq-server    3.6.10   active   1      rabbitmq-server jujucharms 99   ubuntu

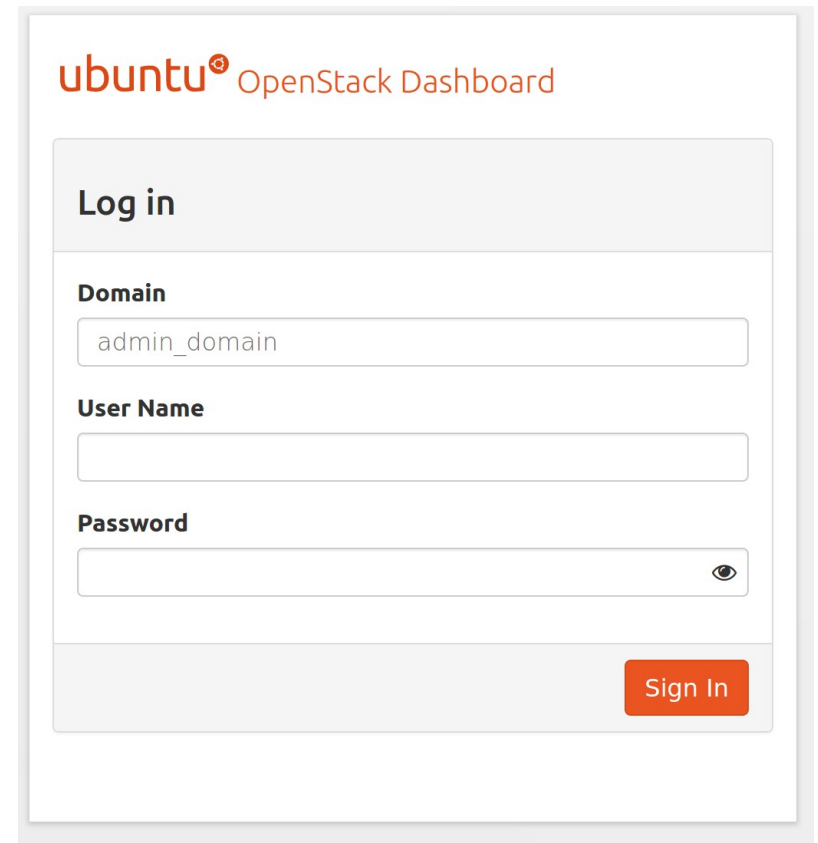
Unit                                Workload  Agent  Machine  Public address  Ports  Message
ceph-mon/3*                        active    idle   31        172.16.1.88      Unit is ready and clustered
ceph-mon/4                        active    idle   33        172.16.1.99      Unit is ready and clustered
ceph-mon/5                        active    idle   34        172.16.1.97      Unit is ready and clustered
ceph-osd/0*                        active    idle   31        172.16.1.88      Unit is ready (1 OSD)
  ntp/0*                            active    idle   31        172.16.1.88      chrony: Ready
ceph-osd/1                        active    idle   33        172.16.1.99      Unit is ready (1 OSD)
  ntp/1                            active    idle   33        172.16.1.99      chrony: Ready
ceph-osd/2                        active    idle   34        172.16.1.97      Unit is ready (1 OSD)
  ntp/2                            active    idle   34        172.16.1.97      chrony: Ready
cinder/9*                         active    idle   30/lxd/16  10.205.237.61    8776/tcp  Unit is ready
glance/12*                        active    idle   30/lxd/15  10.205.237.213  9292/tcp  Unit is ready
keystone/12*                      active    idle   30/lxd/10  10.205.237.26   5000/tcp  Unit is ready
mysql/11*                         active    idle   30/lxd/9   10.205.237.240  3306/tcp  Unit is ready
neutron-api/10*                   active    idle   30/lxd/8   10.205.237.86   9696/tcp  Unit is ready
neutron-gateway/11*               active    idle   30        172.16.1.93      Unit is ready
nova-cloud-controller/12*          active    idle   30/lxd/12  10.205.237.206  8774/tcp,8775/tcp Unit is ready
nova-compute/14*                  active    idle   31        172.16.1.88      Unit is ready
  neutron-openvswitch/8*           active    idle   31        172.16.1.88      Unit is ready
nova-compute/15                   active    idle   33        172.16.1.99      Unit is ready
  neutron-openvswitch/9            active    idle   33        172.16.1.99      Unit is ready
nova-compute/16                   active    idle   34        172.16.1.97      Unit is ready
  neutron-openvswitch/10           active    idle   34        172.16.1.97      Unit is ready
openstack-dashboard/12*            active    idle   30        172.16.1.93      80/tcp,443/tcp Unit is ready
placement/2*                       active    idle   30/lxd/13  10.205.237.105  8778/tcp  Unit is ready
rabbitmq-server/9*                 active    idle   30/lxd/11  10.205.237.95   5672/tcp  Unit is ready

Machine  State  DNS                Inst id          Series  AZ  Message
30       started 172.16.1.93        manual:172.16.1.93 bionic  Manualy provisioned machine
30/lxd/8 started 10.205.237.86      juju-f92316-30-lxd-8 bionic  Container started
30/lxd/9 started 10.205.237.240     juju-f92316-30-lxd-9 bionic  Container started
30/lxd/10 started 10.205.237.26      juju-f92316-30-lxd-10 bionic  Container started
30/lxd/11 started 10.205.237.95      juju-f92316-30-lxd-11 bionic  Container started
30/lxd/12 started 10.205.237.206     juju-f92316-30-lxd-12 bionic  Container started
30/lxd/13 started 10.205.237.105     juju-f92316-30-lxd-13 bionic  Container started
30/lxd/15 started 10.205.237.213     juju-f92316-30-lxd-15 bionic  Container started
30/lxd/16 started 10.205.237.61      juju-f92316-30-lxd-16 bionic  Container started
31       started 172.16.1.88        manual:172.16.1.88 bionic  Manualy provisioned machine
33       started 172.16.1.99        manual:172.16.1.99 bionic  Manualy provisioned machine
34       started 172.16.1.97        manual:172.16.1.97 bionic  Manualy provisioned machine

```

Connect to the dashboard

- Open the browser to the URL:
 - http://CONTROLLER_IP/horizon/auth/login/
- Domain: admin_domain
- User Name: admin
- Password: openstack




The screenshot shows the Ubuntu OpenStack Dashboard login interface. At the top, the text "ubuntu OpenStack Dashboard" is displayed, with "ubuntu" in orange and "OpenStack Dashboard" in grey. Below this is a "Log in" section with a light grey header. The form contains three input fields: "Domain" with the value "admin_domain", "User Name" which is empty, and "Password" which is empty and has a toggle icon (an eye) on the right. At the bottom right of the form is an orange "Sign In" button.

ubuntu[®] OpenStack Dashboard

Log in

Domain

User Name

Password
 

Sign In

Check that compute nodes are working

- Go to
 - Admin-> System -> System Information -> Compute Service

Admin / System / System Information

System Information

Services Compute Services Block Storage Services Network Agents

Filter

Displaying 4 items

Name	Host	Zone	Status	State	Last Updated
nova-conductor	juju-6dd12f-0-lxd-4	internal	Enabled	Up	0 minutes
nova-scheduler	juju-6dd12f-0-lxd-4	internal	Enabled	Up	0 minutes
nova-compute	ATU4IDDQBQFQ72H	nova	Enabled	Up	0 minutes
nova-compute	PCR2LY4JY8FPNH0	nova	Enabled	Up	0 minutes

Displaying 4 items