

Virtual Machine, and Memory Management

1

Virtual machines (VM)

Virtual machines (VM) were first developed in the mid-1960s, and they have remained an important part of mainframe computing over the years.

System Virtual Machines. IBM VM/370

Largely ignored in the single-user PC era in the 1980s and 1990s.

They have recently gained popularity due to

- The increasing importance of isolation and security in modern systems
- The sharing of a single computer among many unrelated users, for Cloud computing

A single computer runs multiple VMs and can support a number of different *operating systems, and also a number of different versions of the same operating systems.*

On a conventional platform, a single OS “owns” all the hardware resources, but with a VM, multiple OS all share the hardware resources.

2

Hypervisor, Host and Guest

Virtual machine monitor (VMM) or Hypervisor, the software that supports VMs, determines how to map virtual resources to physical resources

Host, the underlying hardware platform, and its resources are shared among the **guest** VMs.

VMs improve protection allow us to:

1. Managing software.

VMs provide an abstraction that can run the complete software stack, even including old operating systems like DOS.

2. Managing hardware.

One reason for multiple servers is to have each application running with the compatible version of the operating system on separate computers, as this separation can improve dependability.

Processor virtualization produces overhead.

3

As the VM works

To “virtualize” the processor, the VMM must control just about everything—access to privileged state, I/O, exceptions, and interrupts—even though the guest VM and OS presently running are temporarily using them.

For example, in the case of a timer interrupt, the VMM would suspend the currently running guest VM, save its state, handle the interrupt, determine which guest VM to run next, and then load its state. Guest VMs that rely on a timer interrupt are provided with a virtual timer and an emulated timer interrupt by the VMM.

If a guest OS attempts to access or modify information related to hardware resources via a privileged instruction—for example, reading or writing a status bit that enables interrupts—it will trap to the VMM. The VMM can then affect the appropriate changes to corresponding real resources.

To be in charge, the VMM must be at a higher privilege level than the guest VM, which generally runs in user mode; this also ensures that the execution of any privileged instruction will be handled by the VMM.

4

Overhead of processor virtualization

The overhead of processor virtualization depends on the workload.

User-level processor-bound programs have zero virtualization overhead, because the OS is rarely invoked, so everything runs at native speeds.

I/O-intensive workloads are generally also OS-intensive, executing many system calls and privileged instructions that can result in high virtualization overhead.

On the other hand, if the I/O-intensive workload is also I/O-bound, the cost of processor virtualization can be completely hidden, since the processor is often idle waiting for I/O.

5

Basic system requirements

The basic system requirements to support VMMs are:

- At least two processor modes, system and user.
- A privileged subset of instructions that is available only in system mode, resulting in

a trap if executed in user mode; all system resources must be controllable just via these instructions.

Hence, if any instruction that tries to read or write such sensitive information traps when executed in user mode, the VMM can intercept it and support a virtual version of the sensitive information, as the guest OS expects.

A way to remap the address to physical memory.

6

Virtual memory

The main memory can act as a “cache” for the secondary storage.

Differing historical roots have led to the use of different terminology. A virtual memory block is called a **page**, and a virtual memory miss is called a **page fault**.

7

Virtual Memory motivations

Historically, there were two major motivations:

1. to allow efficient and safe sharing of memory among several programs,
 - such as for the memory needed by multiple virtual machines for Cloud computing, and to
2. remove the programming burdens of a small, limited amount of main memory.

8

Virtual Memory protection and requirements

- Protection via virtual memory
 - Keeps processes in their own memory space
 - to allow multiple virtual machines to share the same physical memory, we must be able to protect the virtual machines from each other, ensuring that a program can just read and write the portions of main memory that have been assigned to it.
- Role of architecture:
 - Provide user mode and supervisor mode
 - Protect certain aspects of CPU state
 - Provide mechanisms for switching between user mode and supervisor mode
 - Provide mechanisms to limit memory accesses
 - Provide TLB to translate addresses


9

Relocation of addresses

Virtual memory also simplifies loading the program for execution by providing relocation.

This relocation allows us to load the program anywhere in main memory.

All virtual memory systems in use today relocate the program as a set of fixed-size blocks (pages), thereby eliminating the need to find a contiguous block of memory to allocate to a program; instead, the operating system needs only to find enough pages in main memory.



A virtual page may be absent from main.

In that case, the page resides on disk.

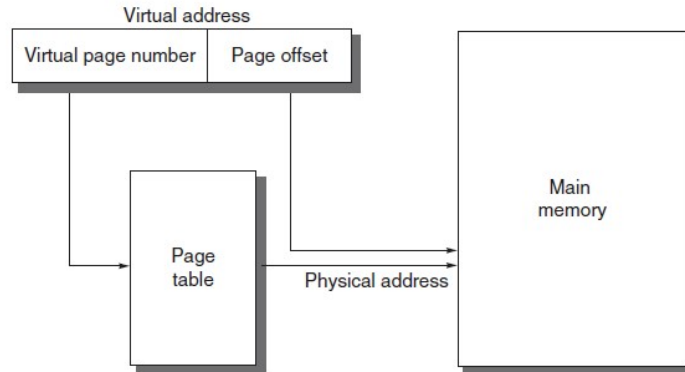
Physical pages can be shared by having two virtual addresses point to the same physical address. This capability is used to allow two different programs to share data or code.

10

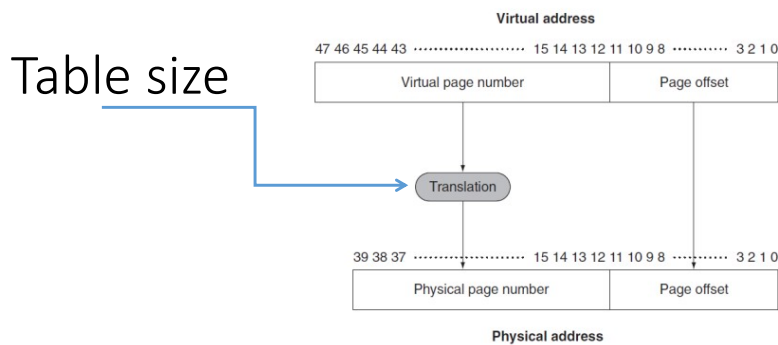
Virtual Memory

Main memory need contain only the active portions of the many virtual machines, just as a cache contains only the active portion of one program.

The principle of locality enables virtual memory as well as caches, and virtual memory allows us to share the processor efficiently as well as the main memory.



11



ARMv8

- It has a 64-bit **virtual address**, the upper 16 bits are not used, so the address to be mapped is 48 bits.
- The **physical memory** is 1 TiB, or 2^{40} bytes, which needs a 40-bit address.
- Page size = 2^{12} bytes = 4 kbytes
- Number of Virtual pages = $2^{(48-12)} = 2^{36} = 64$ G virtual pages
- Number of Virtual pages = $2^{(40-12)} = 2^{28} = 256$ M physical pages

Having a larger number of virtual pages than physical pages is the basis for the illusion of an essentially unbounded amount of virtual memory.

12

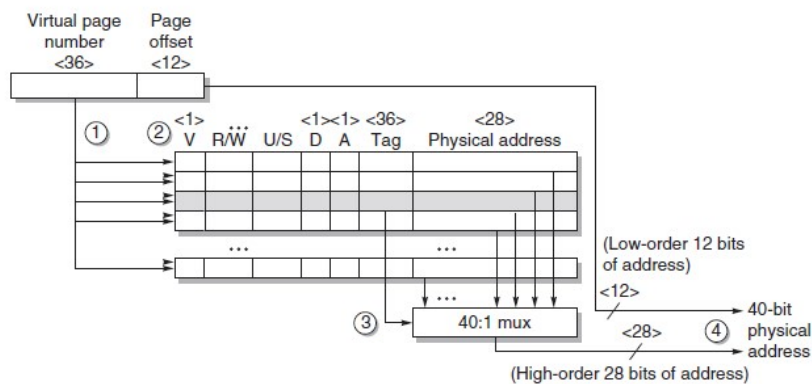
Page Fault cost

- A page fault to disk will take millions of clock cycles to process. This enormous miss penalty, dominated by the time to get the first word for typical page sizes, leads to several key decisions in designing virtual memory systems:
- Pages should be large (4 Kbytes to 64 Kbytes) enough to try to amortize the high access time.
 - Organizations that reduce the page fault rate are attractive.
 - Page faults can be handled in software because the overhead will be small compared to the disk access time.
 - In addition, software can afford to use smart algorithms for choosing how to place pages because even little reductions in the miss rate will pay for the cost of such algorithms.
 - Write-through will not work for virtual memory, since writes take too long.
 - Instead, virtual memory systems use write-back.

13

Virtual Memory


ARM8 supports three options for granule size: 4, 16, and 64 Kbyte.



14

Segmentation

In segmentation, an address consists of two parts: a segment number and a segment offset.



15

ARM Cortex-A8

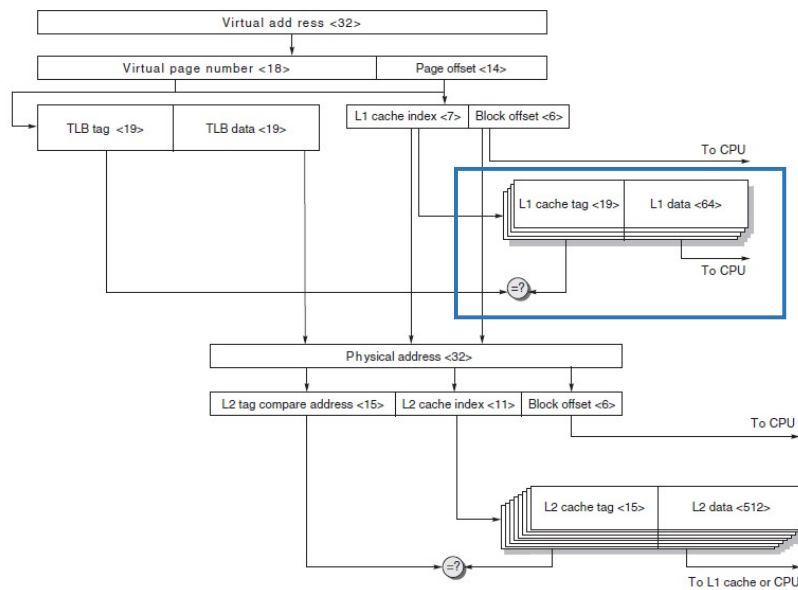
16

Organization of a two-level cache hierarchy using virtually indexed caches

- The Cortex-A8 is a configurable core that supports the ARMv7 instruction set architecture
- Cortex-A8 IP core is used in the Apple iPad and smartphones by several manufacturers including Motorola and Samsung.
- Two-level cache hierarchy with the first level being a pair of caches (for I & D), each 16 KB or 32 KB.
- The optional second-level cache when present is eight-way set associative and can be configured with 128 KB up to 1 MB

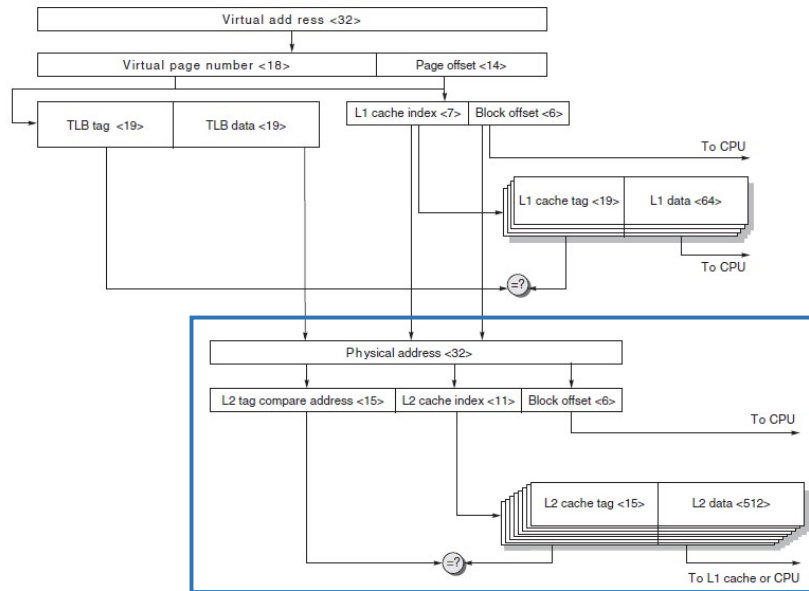
17

ARM Cortex-A8



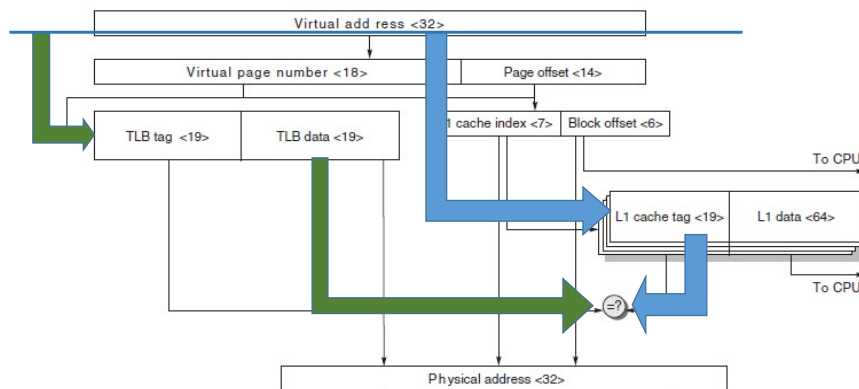
18

ARM Cortex-A8



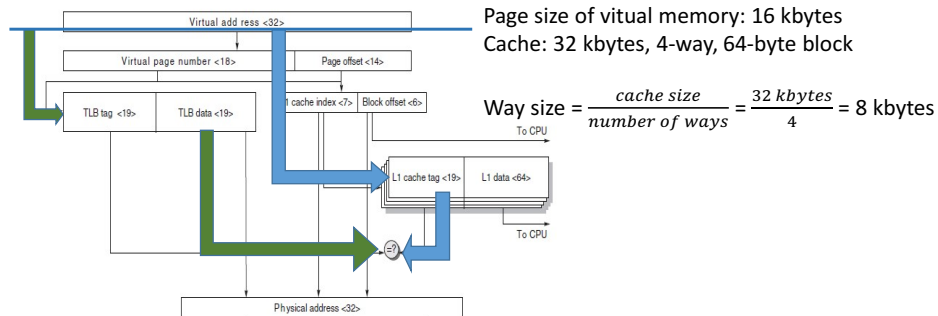
19

ARM Cortex-A8



20

ARM Cortex-A8

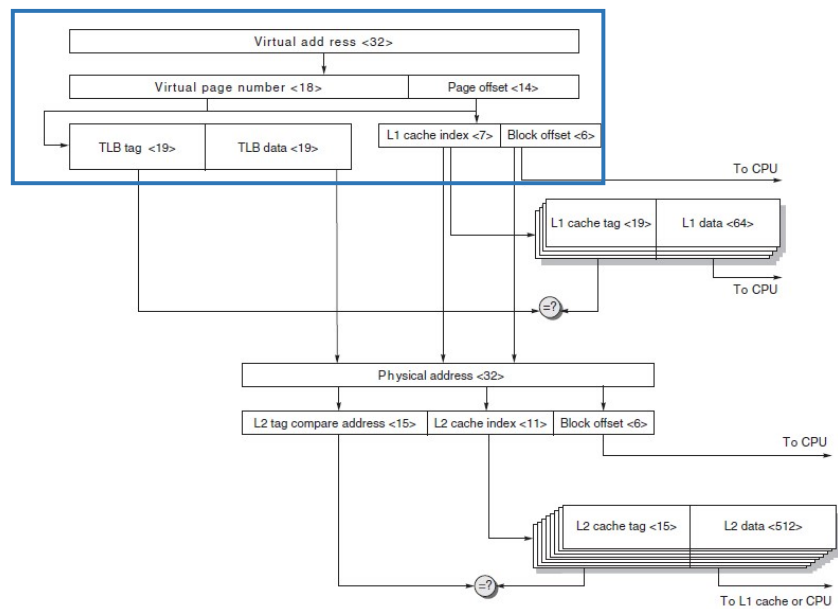


page size of virtual memory \geq size of the cache way

16 kbytes \geq 8 kbytes

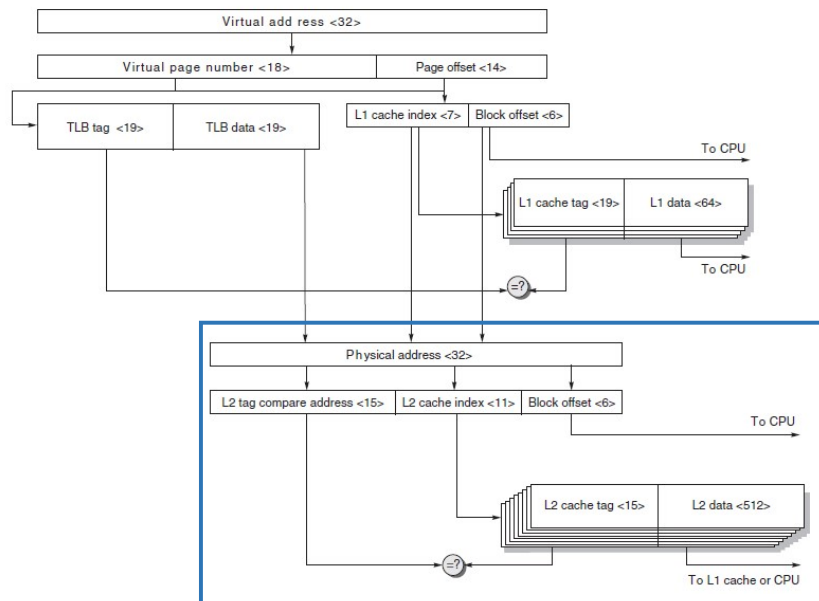
21

ARM Cortex-A8



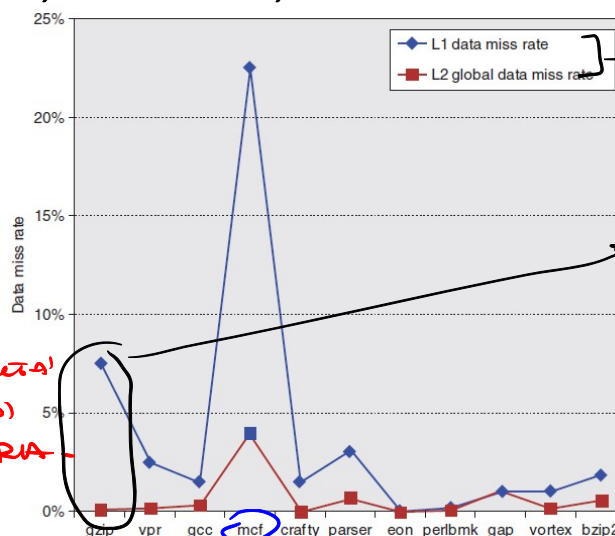
22

ARM Cortex-A8



23

Performance of the Cortex-A8 Memory Hierarchy



VOGLIAMO RIDURRE
L'OVERHEAD DOVUTO
AL COSTO DELL'ACCESSO

DEGGIO ANDARE SUL
PRIMO LIVELLO O
SUL SECONDO?

VOGLIAMO RIDURRE LA LATENZA
NELLA VERSIONE TEMPORALE
ACCESSO A MEMORIA

SUL PRIMO LIVELLO

24

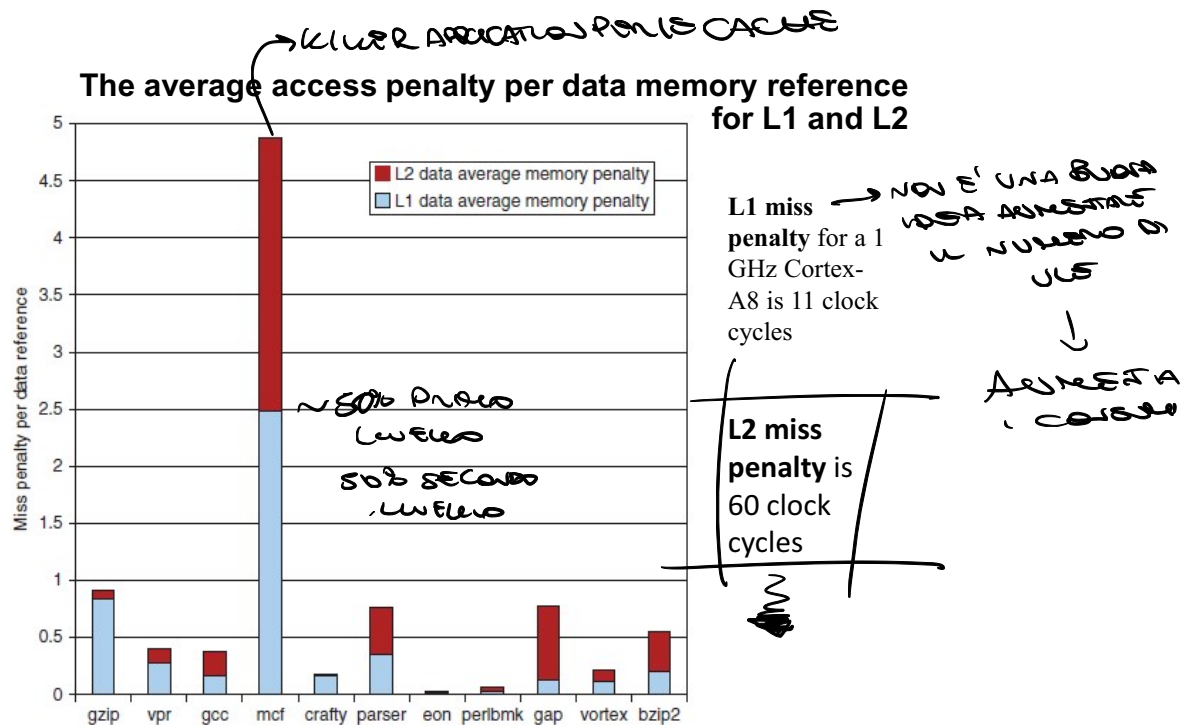
2 LIVELLI DI
CASH

PERCHÉ HA QUESTA
PERFORMANCE?

gzip
SOURCE DESTINATION

SPATIAL COHERENCY
TEMPORAL COHERENCY

LEGGERE SOLO
UNA VOLTA



25

The Processor Performance Equation

CPU time CPU = clock cycles for a program × Clock cycle time

$$CPI = \frac{\text{CPU clock cycles for a program}}{\text{Instruction count}}$$

$$\text{CPU time} = \left(\sum_{i=1}^n IC_i \times CPI_i \right) \times \text{Clock cycle time}$$

where IC_i represents the number of times instruction i is executed in a program and CPI_i represents the average number of clocks per instruction for instruction i .

CPU time = Instruction count × Cycles per instruction × Clock cycle time


Clock cycle time—Hardware technology and organization

CPI—Organization and instruction set architecture

Instruction count—Instruction set architecture and compiler technology

26

The average memory access penalty per data memory reference




27

The Intel Core i7

- i7 supports the x86-64 instruction set architecture, a 64-bit extension of the 80x86 architecture.
- i7 can execute up to four 80x86 instructions per clock cycle.
- i7 uses 48-bit virtual addresses and 36-bit physical addresses.
- Memory management is handled with a two level TLB

28

The Intel i7 memory hierarchy



29

i7: the caches

Characteristic	L1	L2	L3
Size	32 KB I/32 KB D	256 KB	2 MB per core
Associativity	4-way I/8-way D	8-way	16-way
Access latency	4 cycles, pipelined	10 cycles	35 cycles
Replacement scheme	Pseudo-LRU	Pseudo-LRU	Pseudo-LRU but with an ordered selection algorithm

PER RIDUZIONE IL NUMERO
DI BA NECESSARI

OBSERVO: RIDUZIONE CPI DOWTO A
INSTRUCTION ACCESS - 4 way
DATA ACCESS - 8 way
PER RIDUZIONE I PROBLEMI DOWTO A UNA SERIE AL PUNTO CAUSATO AUMENTO # WAY

30

Characteristics of the i7's TLB structure

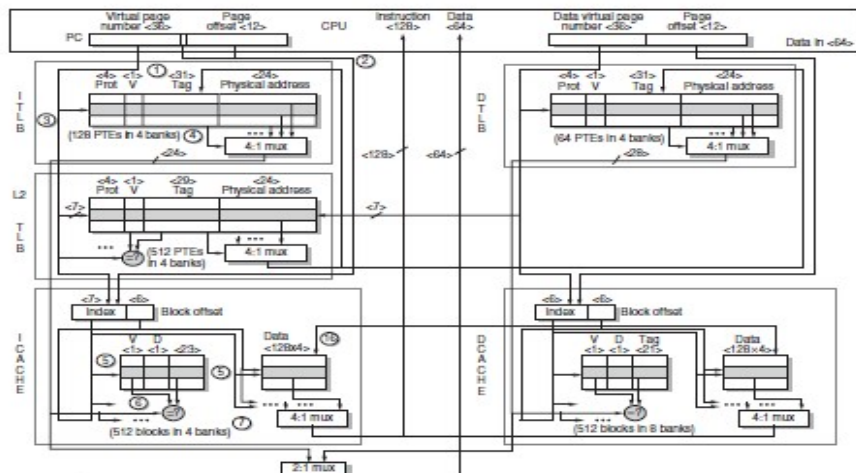
Handwritten notes:
 - *PER ACCESS AND ACCESS FROM 16 PAGE*
 - *NUMBER REGISTER = 16, FULL ASSOCIATIVE CACHES*

Characteristic	Instruction TLB	Data DLB	Second-level TLB
Size	128	64	512
Associativity	4-way	4-way	4-way
Replacement	Pseudo-LRU	Pseudo-LRU	Pseudo-LRU
Access latency	1 cycle	1 cycle	6 cycles
Miss	7 cycles	7 cycles	Hundreds of cycles to access page table

Handwritten notes:
 - *Cache RATIONE USES CACHES!*
 - *AUMENTO DEFICIT*
 - *CACHE USATA PER*
 - *WSPU BLOW*


31

The Intel i7 memory hierarchy




32

The L1 data cache miss rate for 17 SPECCPU2006 benchmarks



33

The L2 and L3 data cache miss rates for 17 SPECCPU2006 benchmarks



34