
Electronics Systems

Computer Engineering

Designing for Low Power

Luca Fanucci

[Adapted from Rabaey's *Digital Integrated Circuits*, Second Edition, ©2003
J. Rabaey, A. Chandrakasan, B. Nikolic]

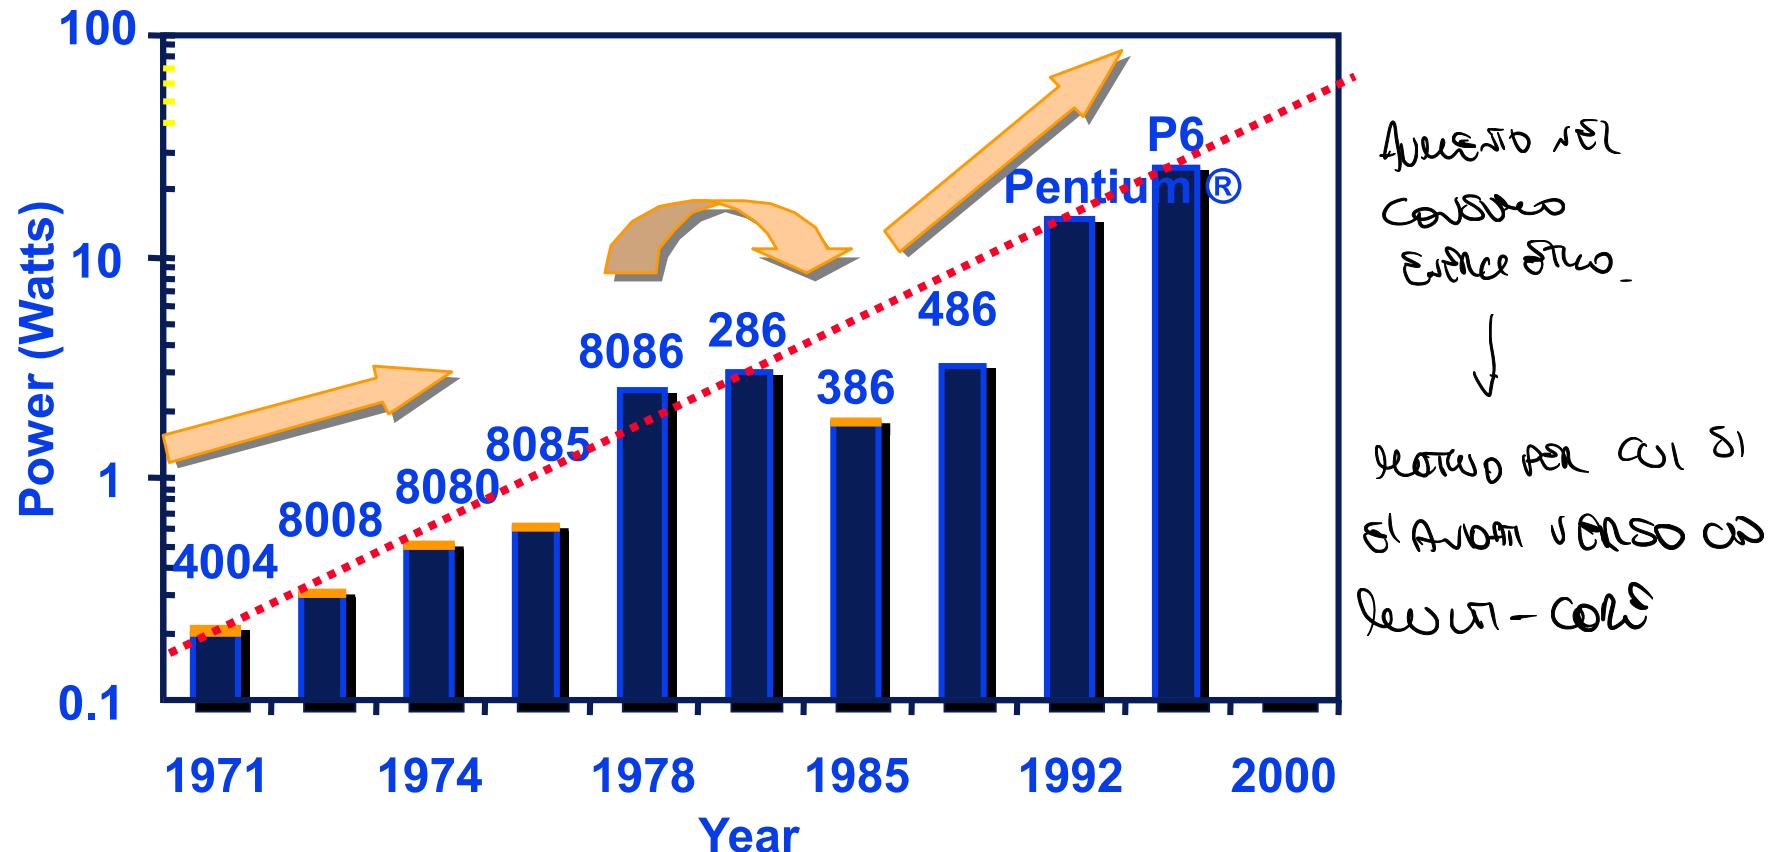
Why Power Matters

- ❑ Packaging costs
- ❑ Power supply rail design
- ❑ Chip and system cooling costs
- ❑ Noise immunity and system reliability
- ❑ Battery life (in portable systems)
- ❑ Environmental concerns
 - Office equipment (professional, government and banks) accounted for 14% of total US commercial energy usage in 2012*
 - *Energy Star* compliant systems. The *Energy Star* program is incorporating standby energy into its ratings. Standby energy in office equipment represents a significant hidden energy cost.

* U.S. Energy Information Administration, 2012 *Commercial Building Energy Consumption Survey: Energy Usage Summary*, Table 1 (March 2016)

Why worry about power? -- Power Dissipation

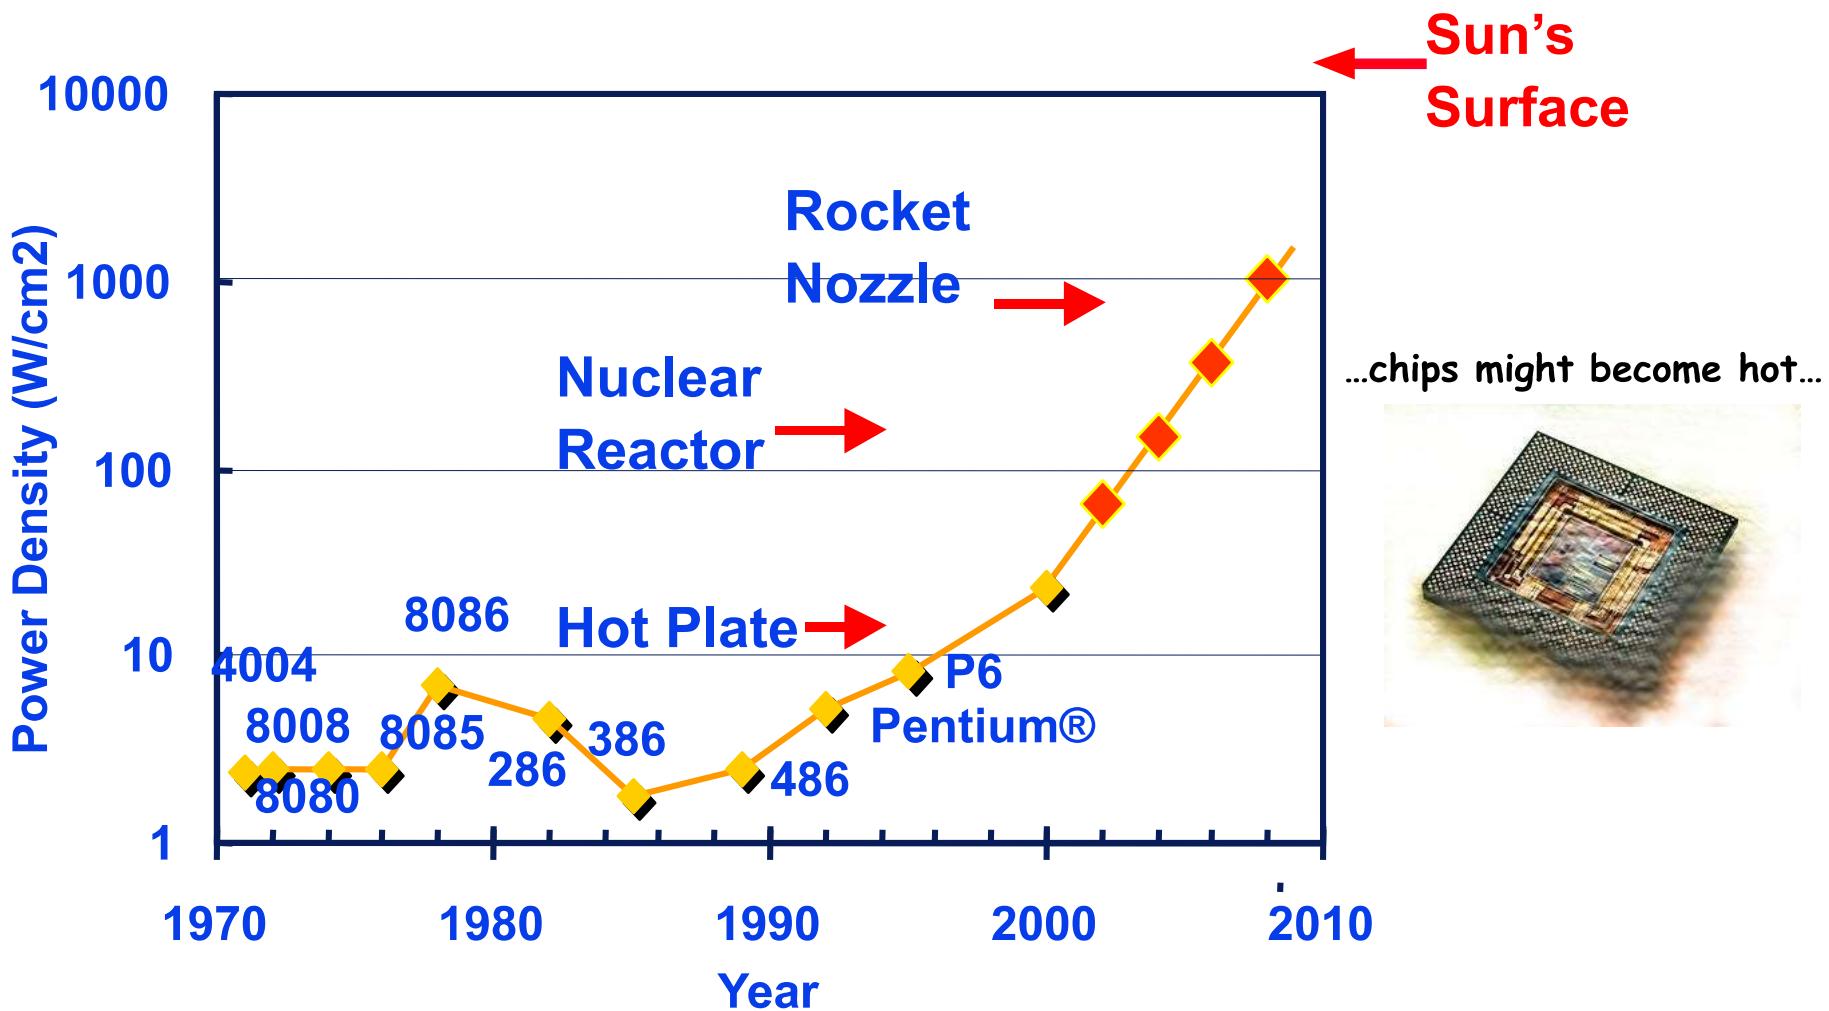
Lead microprocessors power continues to increase



Power delivery and dissipation will be prohibitive

Source: Borkar, De Intel®

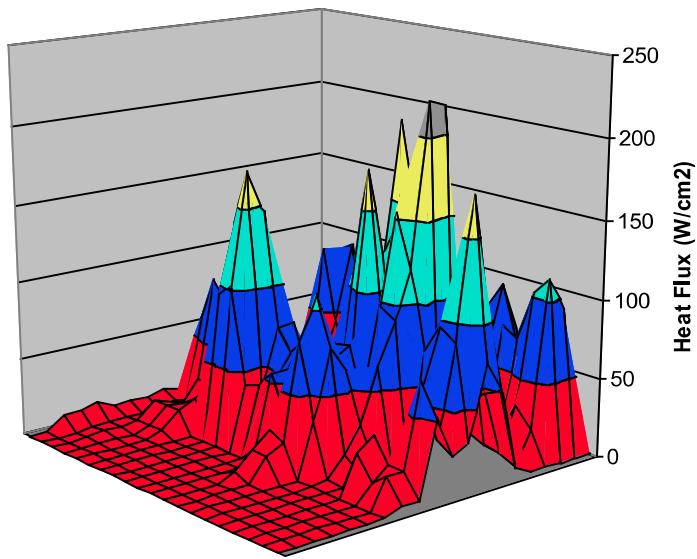
Why worry about power? -- Chip Power Density



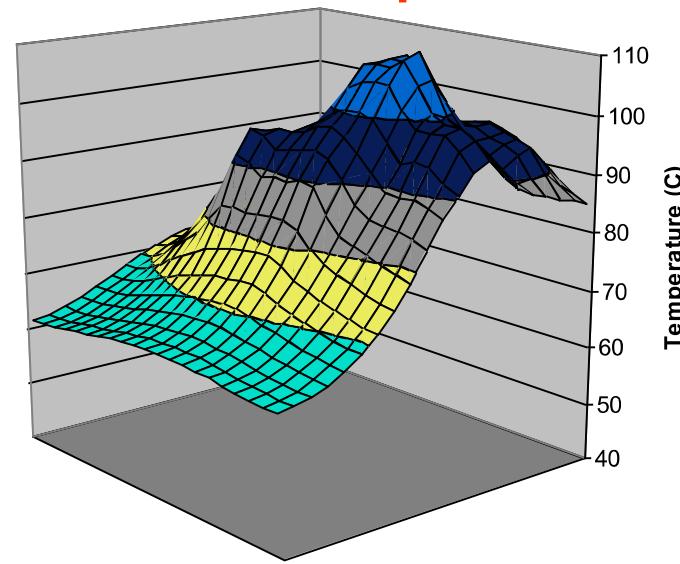
Source: Borkar, De Intel®

Chip Power Density Distribution

Power Map

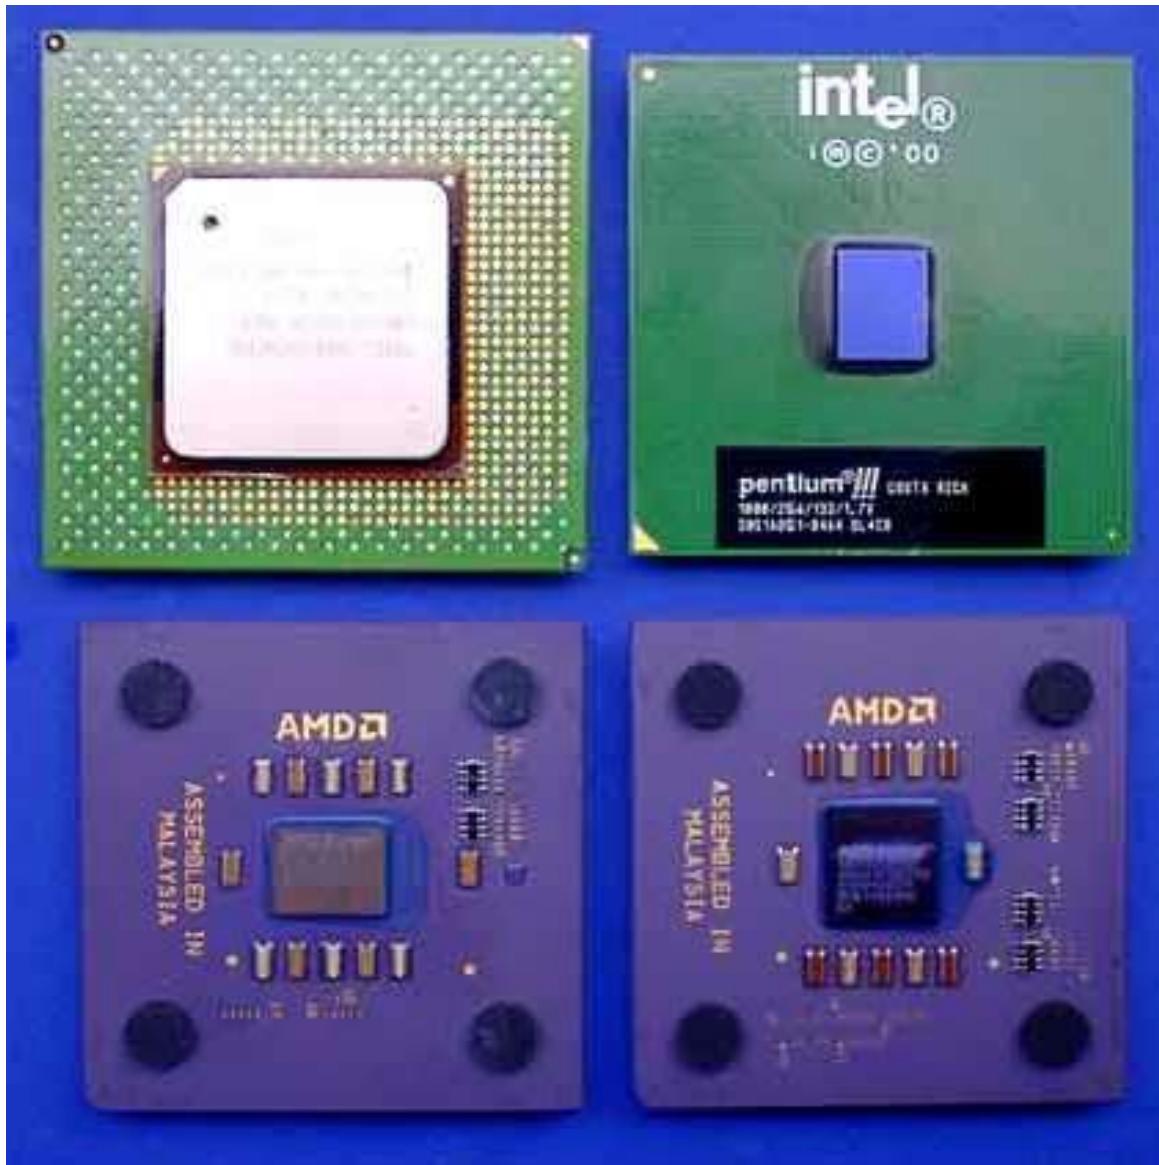


On-Die Temperature



- ❑ Power density is not uniformly distributed across the chip
- ❑ Silicon is not a good heat conductor
- ❑ Max junction temperature is determined by hot-spots
 - Impact on packaging, w.r.t. cooling

Problem Illustration (1/2)

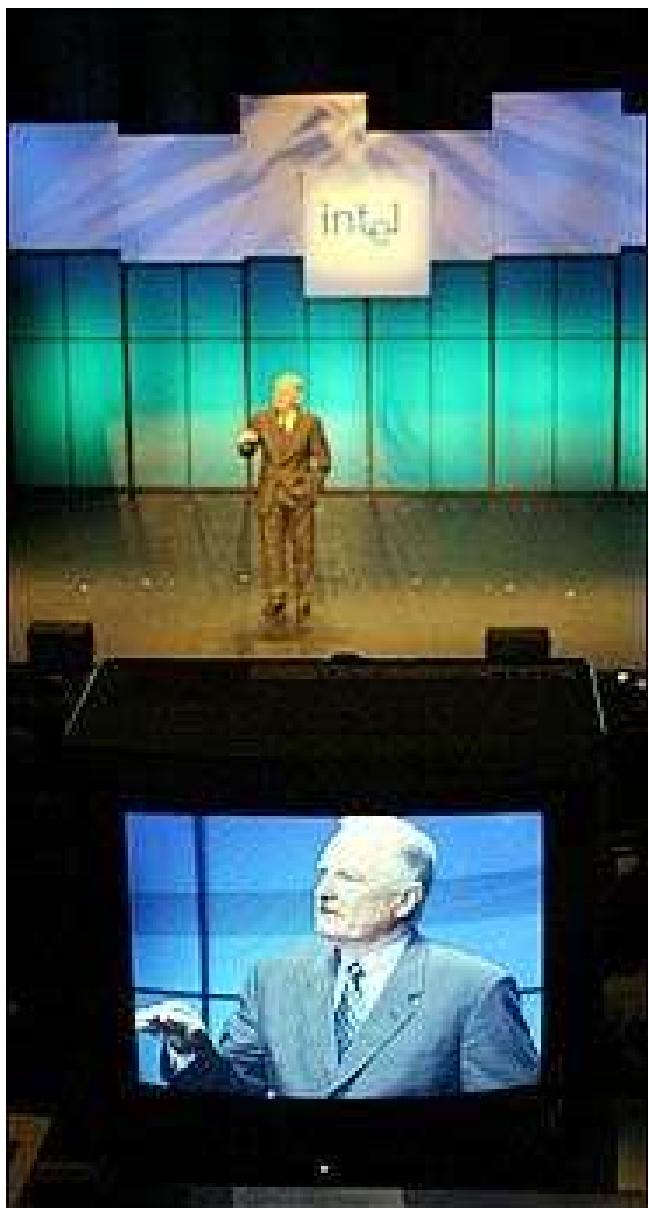


Problem Illustration (2/2)



过热
过热的CPU
弯曲

The Intel's Tejas project



Craig R. Barrett, the chief executive of Intel, told analysts that the company would move down a "**parallel track.**"

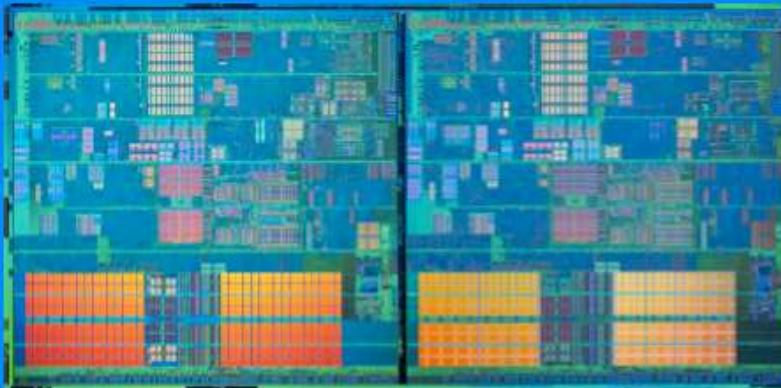
Intel Corporation's newest microprocessor (tejas) was running slower and hotter than its predecessor.

Obtaining more computing power by stamping multiple processors on a single chip rather than straining to increase the speed of a single processor

New York Times, May 17, 2004

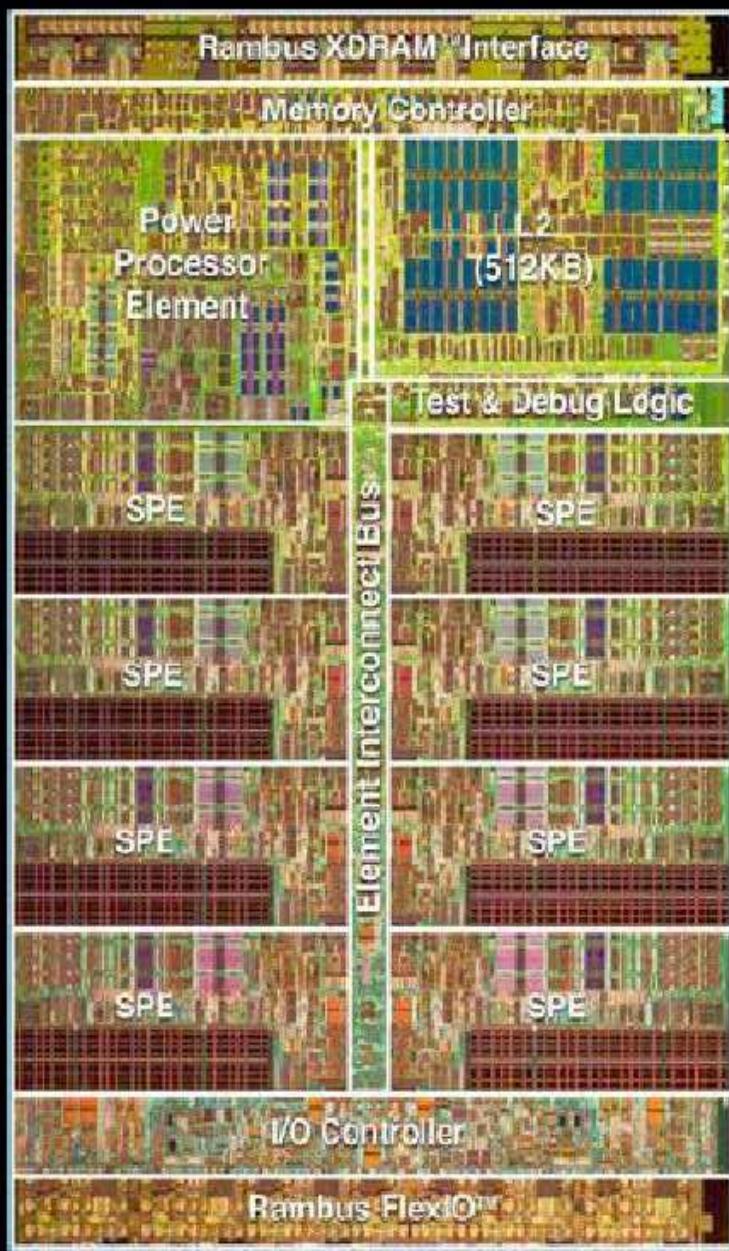
Pentium4 processor

- Dual-Core/Multi-Threaded Pentium®4 Processor on 90nm process
 - 2-1M caches, speeds to 3.2Ghz, support for over clocking, up to 4 threads.
- Shared 800Mhz quad-pumped FSB.
 - Independent bus tuning per agent
- Enhanced auto-halt and 2-state speed step power management
 - Independent events supported per core.



Highlights (3.2 GHz)

- 241M transistors
- 235mm²
- 9 cores, 10 threads
- >200 GFlops (SP)
- >20 GFlops (DP)
- Up to 25 GB/s memory B/W
- Up to 75 GB/s I/O B/W
- >300 GB/s EIB
- Top frequency >4GHz
(observed in lab)



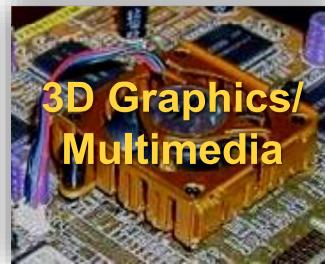
The Performance vs. Power Dilemma

Maintain Battery Life



Wireless/
Handheld
Lowest leakage
and/or dynamic power

Increase Performance



3D Graphics/
Multimedia
Thermal management
Packaging, cooling,
cost



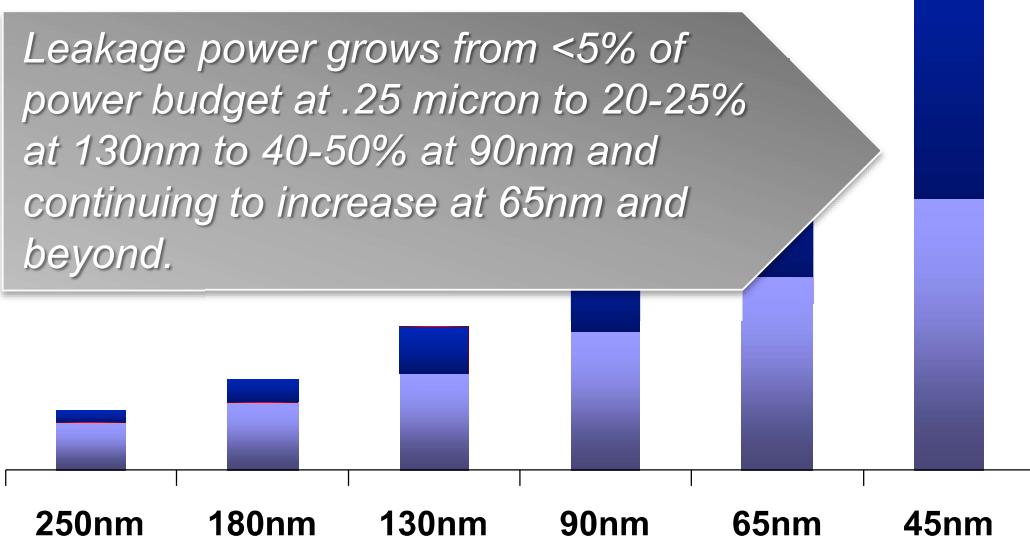
130nm,
90nm, 65nm
Technology

Lower Cost
Increased leakage
IR-drop
Electromigration

Power Management Is Challenging

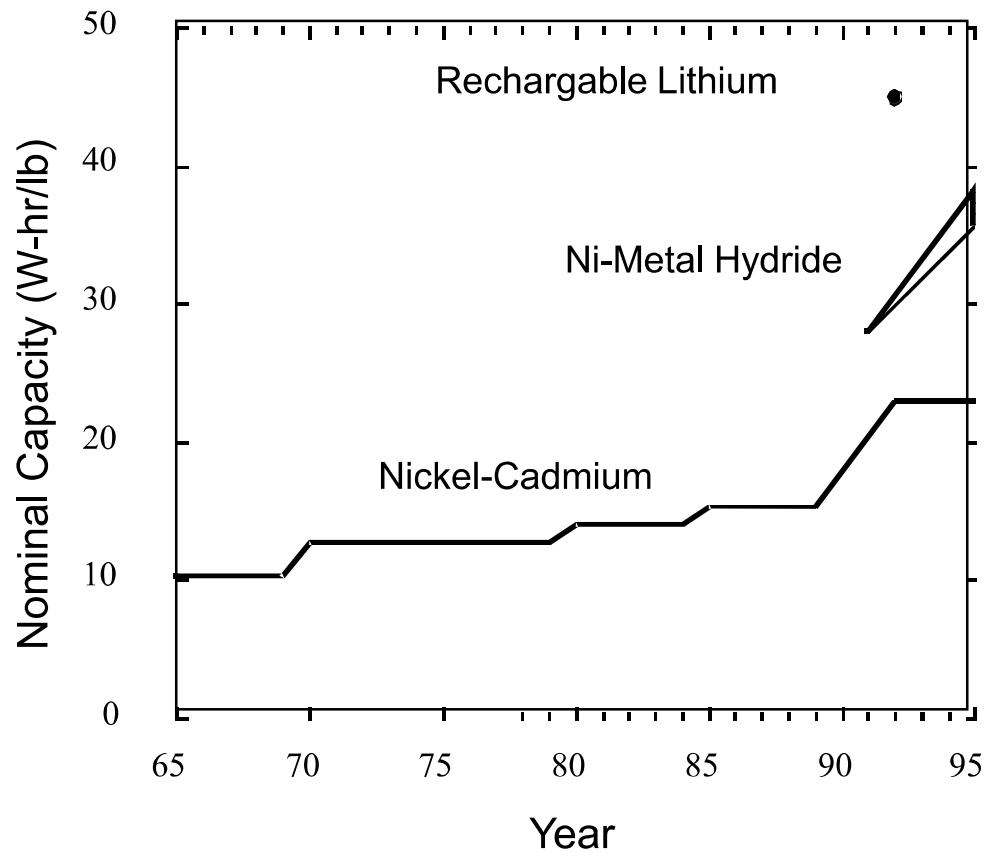
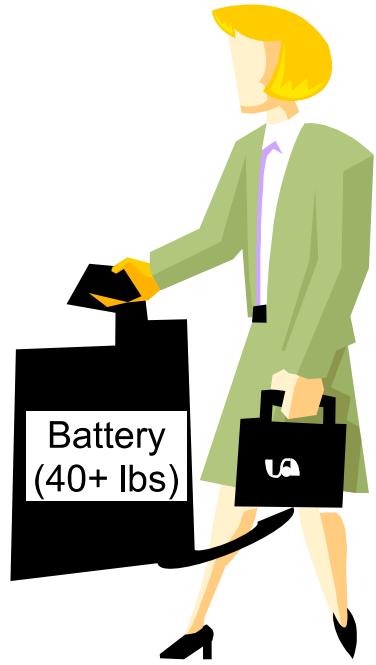
- Leakage Power
- Dynamic Power

Leakage power grows from <5% of power budget at .25 micron to 20-25% at 130nm to 40-50% at 90nm and continuing to increase at 65nm and beyond.



Leakage power begins to dominate at advanced process geometries

Why worry about power ? -- Battery Size/Weight



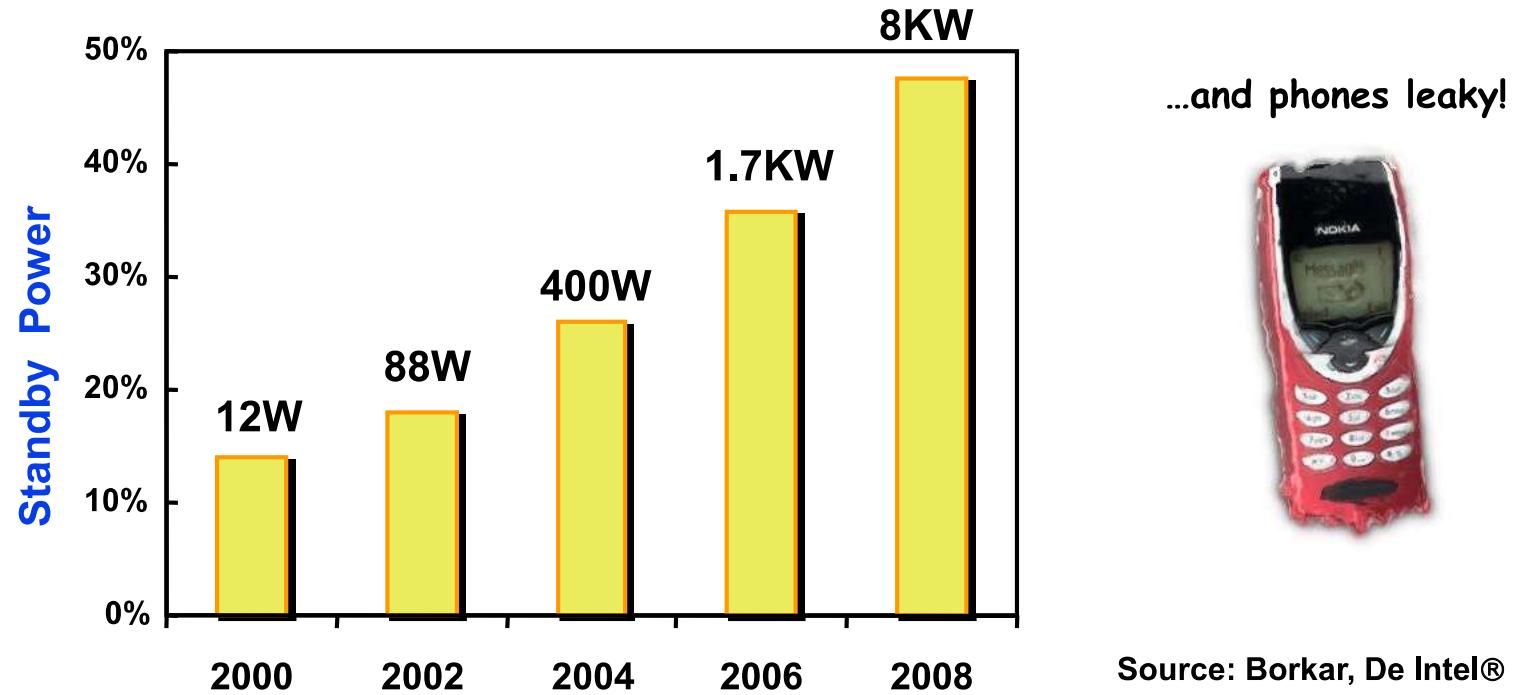
Expected battery lifetime increase
over the next 5 years: **30 to 40%**

From Rabaey, 1995

Why worry about power? -- Standby Power

Year	2002	2005	2008	2011	2014
Power supply V_{dd} (V)	1.5	1.2	0.9	0.7	0.6
Threshold V_T (V)	0.4	0.4	0.35	0.3	0.25

- Drain leakage will increase as V_T decreases to maintain noise margins and meet frequency demands, leading to excessive **battery draining standby** power consumption.



Low power design challenge

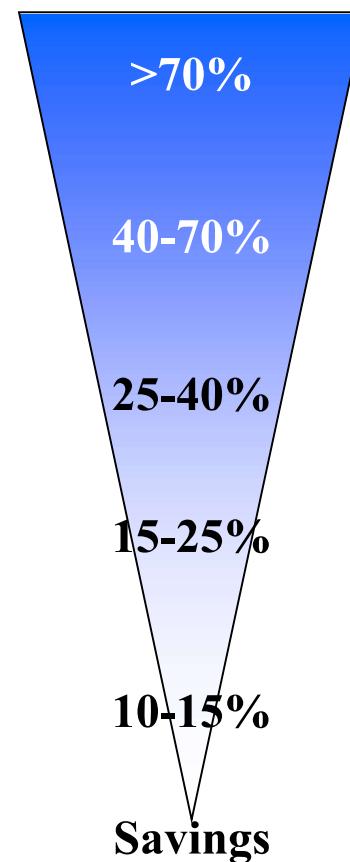
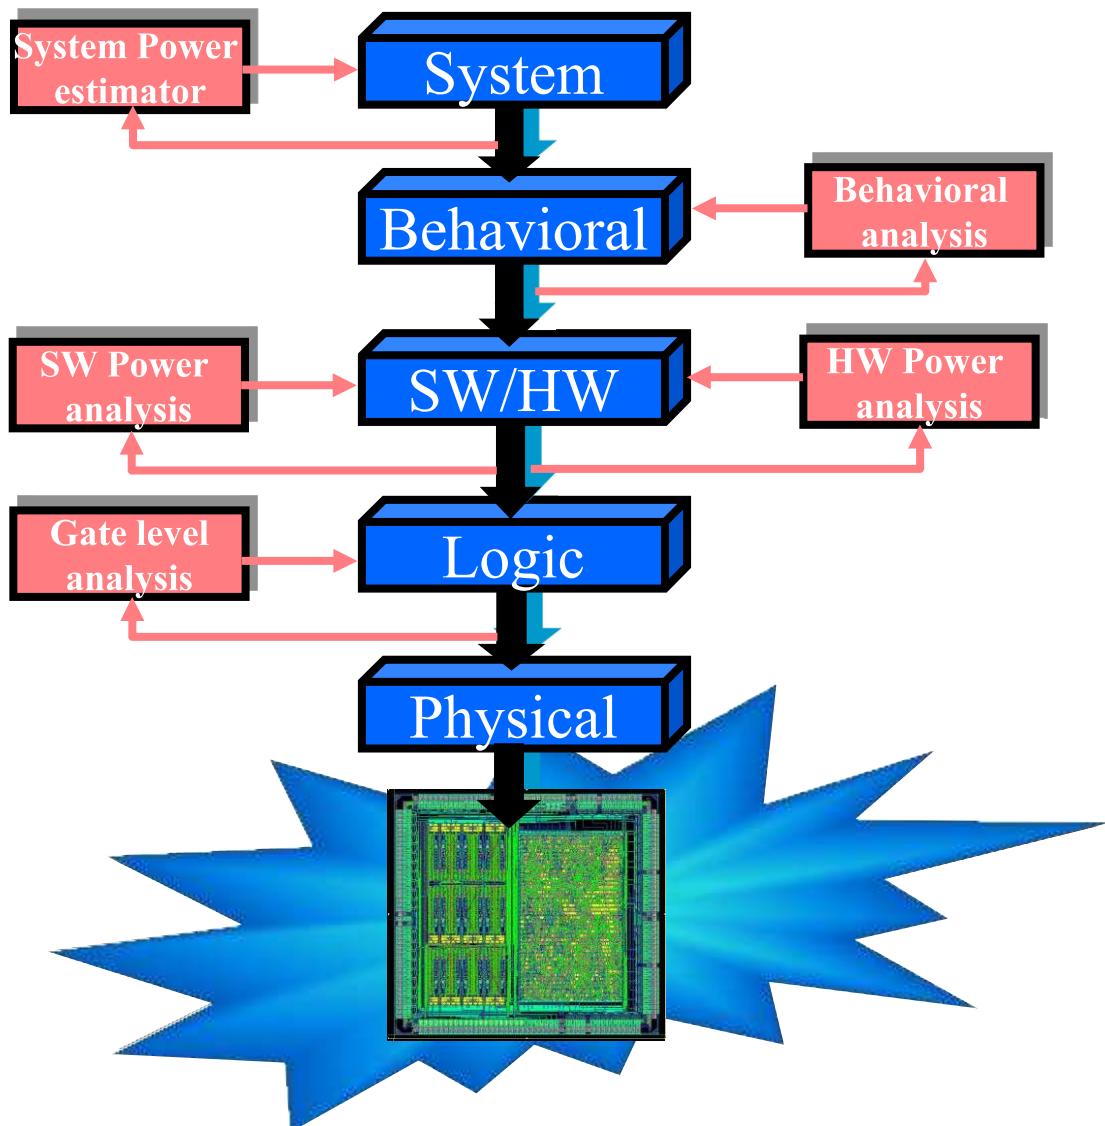
The challenge

“To design an embedded system (HW *and* SW) that provides the target functionality with minimum power consumption”

The solution

From the **system concept** down to the **implementation phase**, adopt a **design style** that includes **power consumption as a figure of merit**, and exploit all the opportunities and techniques available at each design level to reduce it

Power Saving Opportunities



CMOS Energy & Power Equations

$$E = C_L V_{DD}^2 P_{0 \rightarrow 1} + t_{sc} V_{DD} I_{peak} P_{0 \rightarrow 1} + V_{DD} I_{leakage}$$

$$f_{0 \rightarrow 1} = P_{0 \rightarrow 1} * f_{clock}$$

$$P = C_L V_{DD}^2 f_{0 \rightarrow 1} + t_{sc} V_{DD} I_{peak} f_{0 \rightarrow 1} +$$

Dynamic power

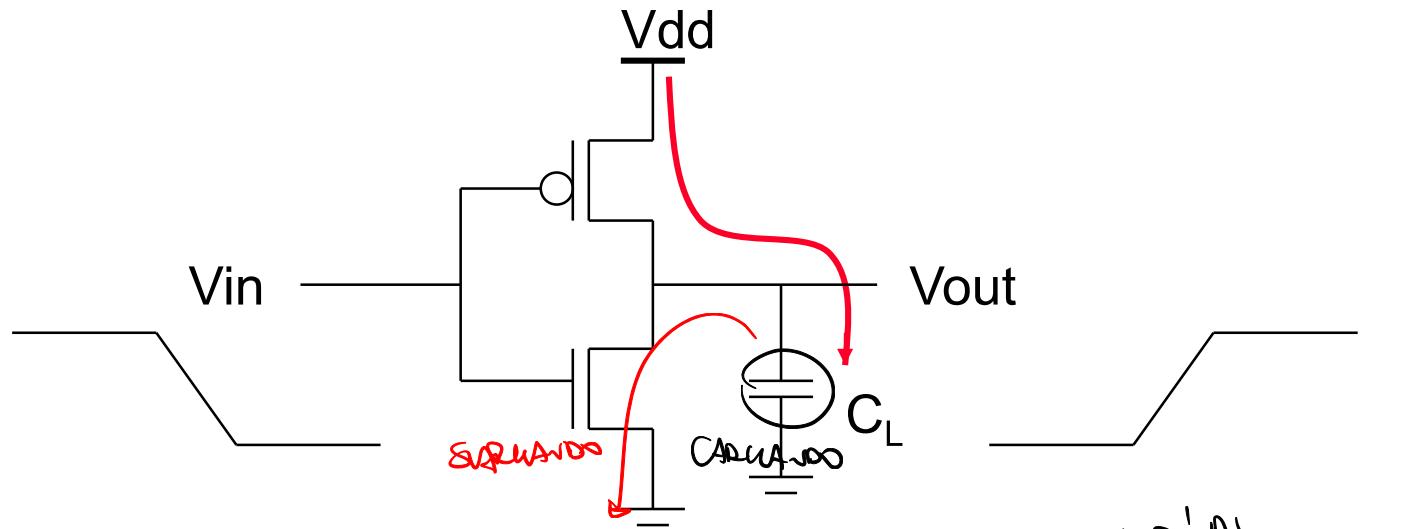
Short-circuit power

$V_{DD} I_{leakage}$

Leakage power

DARIDORRE IL PD
POSSIBLUE.

Dynamic Power Consumption



$$\text{Energy/transition} = C_L * V_{DD}^2 * P_{0 \rightarrow 1}$$

$$P_{dyn} = \text{Energy/transition} * f = C_L * V_{DD}^2 * P_{0 \rightarrow 1} * f$$

$$P_{dyn} = C_{EFF} * V_{DD}^2 * f \quad \text{where } C_{EFF} = P_{0 \rightarrow 1} C_L$$

∅ Switching activity → ∅ Dynamic power

Not a function of transistor sizes!

Data dependent - a function of **switching activity**!

Power = $\frac{1}{2} C_{eff} V_{dd}^2 f$
Switching activity = $f_{0 \rightarrow 1}$

Lowering Dynamic Power

Capacitance:

Function of fan-out,
wire length, transistor
sizes

RISURRE CA
CAPACITÀ

Activity factor:

How often, on average,
do wires switch?

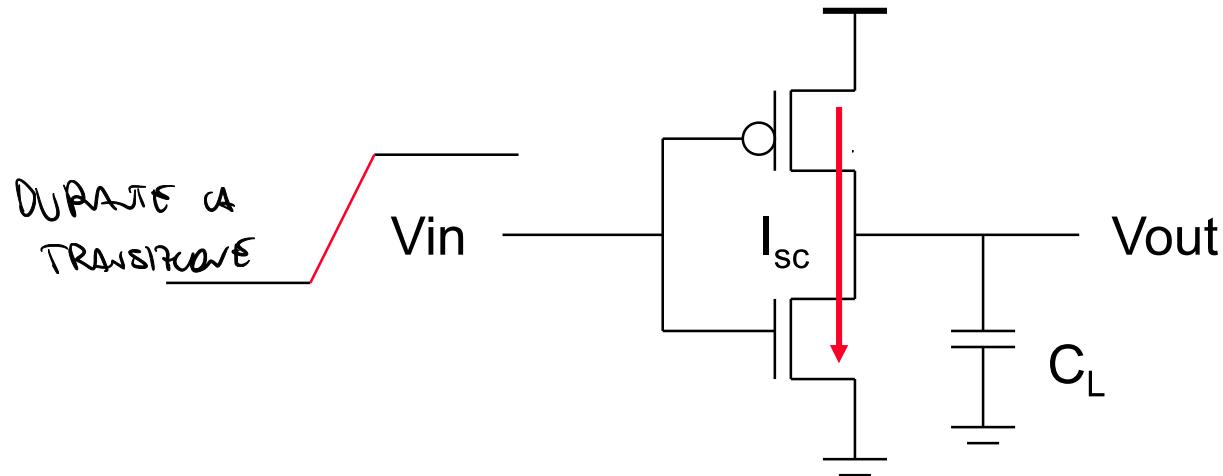
$$P_{\text{dyn}} = C_L V_{DD}^2 P_{0 \rightarrow 1} f$$

RIVOLCO SUPPLY VOLTAGE, PROBLEMI
CON PERFORMANCE $\propto \frac{1}{V_{DD} - V_{t,m}}$ AVERAGI

Supply Voltage:
Has been dropping
with successive
generations

Clock frequency:
Increasing...

Short Circuit Power Consumption



Finite slope of the input signal causes a direct current path between V_{DD} and GND for a short period of time during switching when both the NMOS and PMOS transistors are conducting.

Short Circuit Currents Determinants

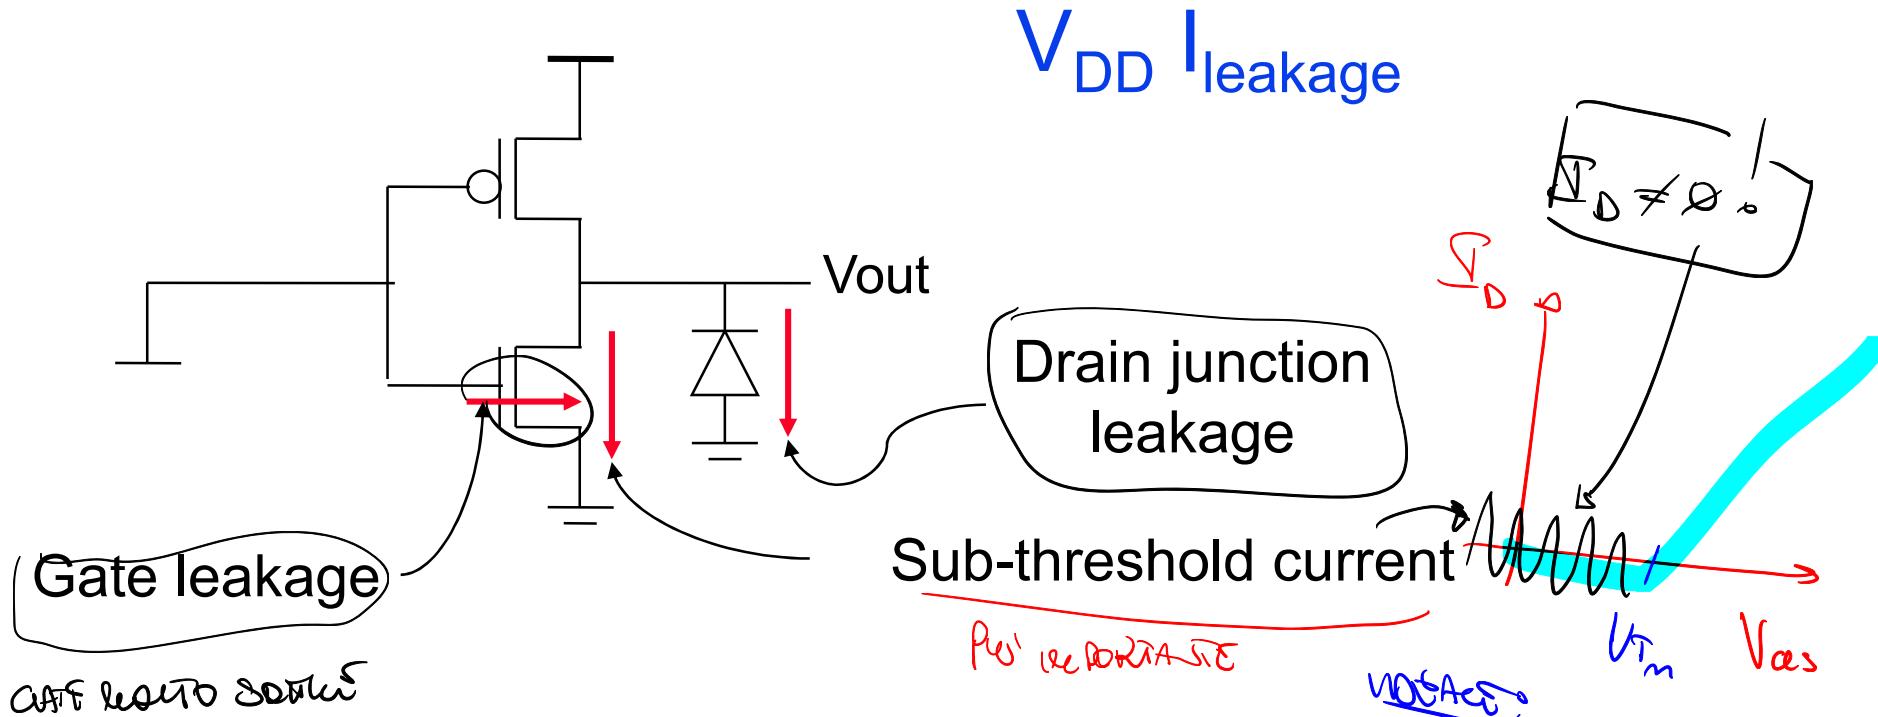
$$E_{sc} = t_{sc} V_{DD} I_{peak} P_{0 \rightarrow 1}$$

$$P_{sc} = t_{sc} V_{DD} I_{peak} f_{0 \rightarrow 1}$$

- Duration and slope of the input signal, t_{sc}
- I_{peak} determined by
 - the saturation current of the P and N transistors which depend on their **sizes**, process technology, temperature, etc.
 - strong function of the ratio between input and output slopes
 - a function of C_L

A sono Arcate
TECHNICAL ASK
AUDREY
US NEW STAR

Leakage (Static) Power Consumption



Sub-threshold current is the dominant factor.

All increase **exponentially** with temperature!

Aumenta con la temperatura

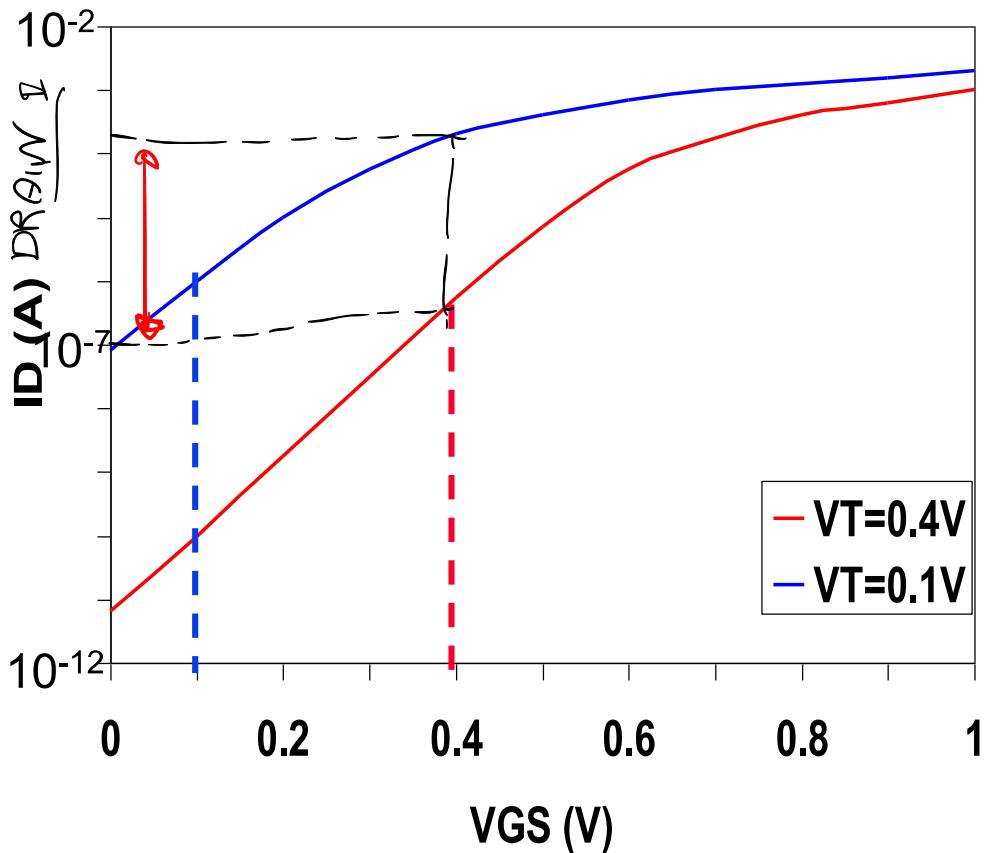
relacionada con



Aumenta con V_T (Al aumentar el V_T)

Leakage as a Function of V_T

- Continued scaling of supply voltage and the subsequent scaling of threshold voltage will make subthreshold conduction a dominate component of power dissipation.



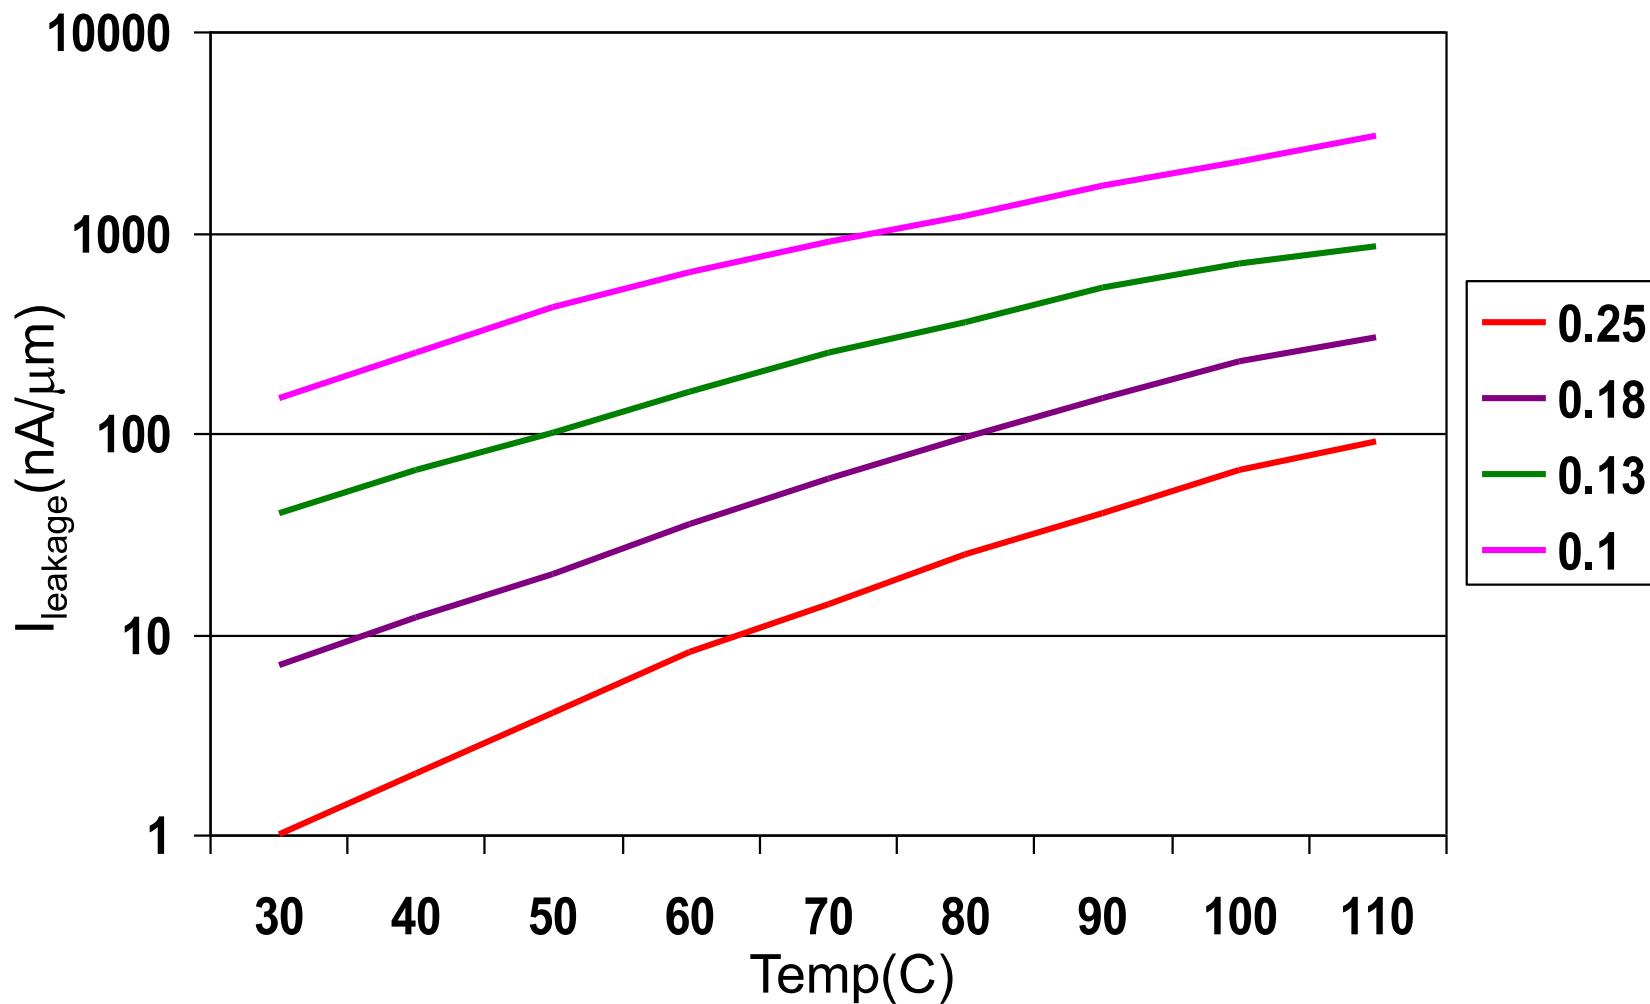
- An 90mV/decade V_T roll-off - so each 255mV increase in V_T gives 3 orders of magnitude reduction in leakage (but adversely affects performance)

TSMC Processes Leakage and V_T

	CL018 G	CL018 LP	CL018 ULP	CL018 HS	CL015 HS	CL013 HS
V_{dd}	1.8 V	1.8 V	1.8 V	2 V	1.5 V	1.2 V
T_{ox} (effective)	42 Å	42 Å	42 Å	42 Å	29 Å	24 Å
L_{gate}	0.16 μm	0.16 μm	0.18 μm	0.13 μm	0.11 μm	0.08 μm
I_{DSat} (n/p) (μA/μm)	600/260	500/180	320/130	780/360	860/370	920/400
I_{off} (leakage) (pA/μm)	20	1.60	0.15	300	1,800	13,000
V_{Tn}	0.42 V	0.63 V	0.73 V	0.40 V	0.29 V	0.25 V
FET Perf. (GHz)	30	22	14	43	52	80

From MPR, 2000

Exponential Increase in Leakage Currents



From De, 1999

Review: Energy & Power Equations

$$E = C_L V_{DD}^2 P_{0 \rightarrow 1} + t_{sc} V_{DD} I_{peak} P_{0 \rightarrow 1} + V_{DD} I_{leakage}$$

$$f_{0 \rightarrow 1} = P_{0 \rightarrow 1} * f_{clock}$$

$$P = C_L V_{DD}^2 f_{0 \rightarrow 1} + t_{sc} V_{DD} I_{peak} f_{0 \rightarrow 1} + V_{DD} I_{leakage}$$

Dynamic power
(% decreasing
relatively with
deep submicron)

Short-circuit
power

Leakage power
(% increasing
with deep
submicron)

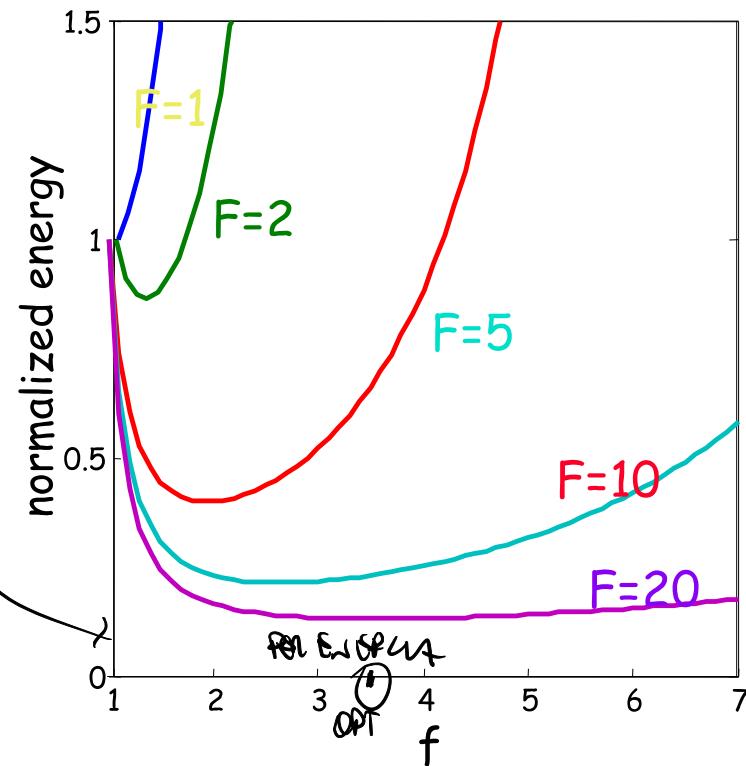
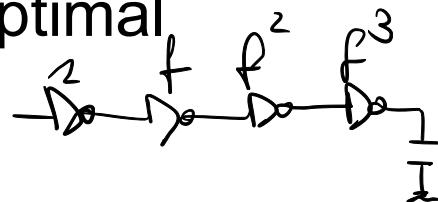
Leakage power grows from <5% of power budget at .25 micron to 20-25% at 130nm to 40-50% at 90nm and continuing to increase at 65nm and beyond.

Power and Energy Design Space

	Constant Throughput/Latency	Variable Throughput/Latency	
Energy	Design Time	Non-active Modules	Run Time
Active	Logic Design Reduced V_{dd}  Sizing Multi- V_{dd}	Clock Gating	DFS, DVS (Dynamic Freq, Voltage Scaling)
Leakage	+ Multi- V_T	Sleep Transistors Multi- V_{dd} Variable V_T	+ Variable V_T

Dynamic Power as a Function of Device Size

- ❑ Device sizing affects dynamic energy consumption
 - gain is largest for networks with large overall effective fan-outs ($F = C_L/C_{g,1}$) *Decreases as F increases*
- ❑ The optimal gate sizing factor (f) for dynamic energy is smaller than the one for performance, especially for large F 's
 - e.g., for $F=20$,
 $f_{\text{opt}}(\text{energy}) = 3.53$ while
 $f_{\text{opt}}(\text{performance}) = 4.47$
- ❑ If energy is a concern avoid oversizing beyond the optimal



From Nikolic, UCB

Standard-Cell Technology Library

Austriamicrosystems, 0,35 um CMOS



0.35µm CMOS

29

Digital Standard Cell Databook

Conditions for characterization library c35_CORELIBD_BC, corner c35_CORELIBD_BC_best: Vdd= 3.63V, Tj= -50.0 deg. C .
 Output transition is defined from 20% to 80% (rising) and from 80% to 20% (falling) output voltage.
 Propagation delay is measured from 50% (input rise) or 50% (input fall) to 50% (output rise) or 50% (output fall).

Strength	1
Cell Area	29.120 μm^2
Equation	$Q = \text{!}A$
Type	Combinational
Input	A
Output	Q



State Table	
A	Q
L	H
H	L

Capacitance [fF]	
A	2.8210

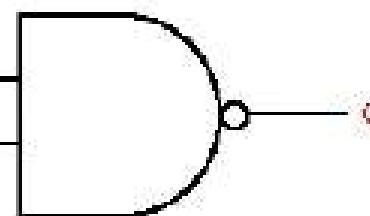
Propagation Delay [ns]				
Input Transition [ns]		0.01 4.00		
Load Capacitance [fF]	5.00	100.00	5.00	100.00
A to Q	fall	0.04	0.41	-0.36
	rise	0.06	0.67	0.83
		0.63	1.88	

Output Transition [ns]					
Input Transition [ns]		0.01		4.00	
Load Capacitance [fF]	5.00	100.00	5.00	100.00	
A to Q	fall	0.04	0.54	0.70	1.43
	rise	0.08	1.03	0.62	1.58

Dynamic Power Consumption [nW/MHz]					
Input Transition [ns]		0.01		4.00	
Load Capacitance [fF]	5.00	100.00	5.00	100.00	
A to Q	fall	1.93	2.35	516.63	510.09
	rise	38.82	40.07	813.89	712.95

Leakage [pW]	
0.26	

Strength	1
Cell Area	43.680 μm^2
Equation	$Q = \text{!}(A \& B)$
Type	Combinational
Input	A, B
Output	Q



State Table		
A	B	Q
L	-	H
H	H	L
-	L	H

Propagation Delay [ns]				
Input Transition [ns]		0.01		4.00
Load Capacitance [fF]	5.00	100.00	5.00	100.00
A to Q	fall	0.06	0.58	-0.13
	rise	0.07	0.68	0.70
B to Q	fall	0.06	0.58	-0.37
	rise	0.08	0.69	0.91
				1.77
				0.60
				1.87

Output Transition [ns]				
Input Transition [ns]		0.01		4.00
Load Capacitance [fF]	5.00	100.00	5.00	100.00
A to Q	fall	0.07	0.83	0.77
	rise	0.10	1.05	0.75
B to Q	fall	0.07	0.83	0.76
	rise	0.11	1.06	0.82
				1.64
				1.71
				1.53
				1.72

Capacitance [fF]	
A	2.7240
B	3.0190

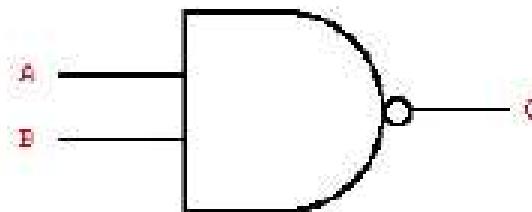
Leakage [pW]	
	0.28

Dynamic Power Consumption [nW/MHz]				
Input Transition [ns]		0.01		4.00
Load Capacitance [fF]	5.00	100.00	5.00	100.00
A to Q	fall	8.00	8.49	555.73
	rise	54.46	55.63	732.49
B to Q	fall	7.37	7.81	616.29
	rise	64.63	64.73	837.28
				471.83
				629.72
				518.84
				703.47

NAND2X1

am
www.ams.com

Strength	2
Cell Area	43.680 μm^2
Equation	$Q = \neg(A \& B)$
Type	Combinational
Input	A, B
Output	Q



State Table		
A	B	Q
L	-	H
H	H	L
-	L	H

Propagation Delay [ns]

Input Transition [ns]		0.01		4.00	
Load Capacitance [fF]		10.00	200.00	10.00	200.00
A to Q	fall	0.06	0.66	-0.09	1.13
	rise	0.06	0.75	0.61	1.81
B to Q	fall	0.06	0.66	-0.27	0.79
	rise	0.06	0.69	0.75	1.78

Output Transition [ns]

Input Transition [ns]		0.01		4.00	
Load Capacitance [fF]		10.00	200.00	10.00	200.00
A to Q	fall	0.07	0.92	0.75	1.73
	rise	0.08	1.16	0.71	1.80
B to Q	fall	0.07	0.92	0.74	1.61
	rise	0.08	1.07	0.80	1.75

Capacitance [fF]	
A	3.8420
B	4.4020

Leakage [pW]	
A	0.31

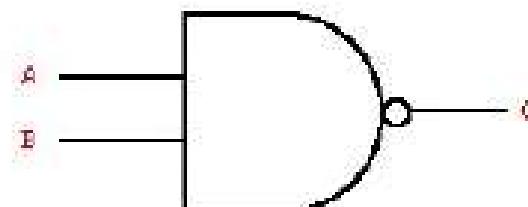
Dynamic Power Consumption [nW/MHz]

Input Transition [ns]		0.01		4.00	
Load Capacitance [fF]		10.00	200.00	10.00	200.00
A to Q	fall	5.99	6.95	974.79	808.95
	rise	64.40	64.67	1272.31	1068.90
B to Q	fall	4.28	5.14	1162.42	945.27
	rise	77.18	78.81	1504.40	1252.61

NAND2X2

am
www.ams.com

Strength	6
Cell Area	72.800 μm^2
Equation	$Q = \neg(A \& B)$
Type	Combinational
Input	A, B
Output	Q



State Table		
A	B	Q
L	-	H
H	H	L
-	L	H

Propagation Delay [ns]

Input Transition [ns]		0.01		4.00	
Load Capacitance [fF]		30.00	600.00	30.00	600.00
A to Q	fall	0.05	0.65	-0.04	1.16
	rise	0.05	0.68	0.51	1.68
B to Q	fall	0.05	0.65	-0.28	0.80
	rise	0.06	0.69	0.73	1.77

Output Transition [ns]

Input Transition [ns]		0.01		4.00	
Load Capacitance [fF]		30.00	600.00	30.00	600.00
A to Q	fall	0.06	0.92	0.72	1.70
	rise	0.07	1.06	0.70	1.73
B to Q	fall	0.06	0.92	0.72	1.59
	rise	0.08	1.07	0.78	1.74

Capacitance [fF]

Leakage [pW]

A	10.1390
B	11.0760

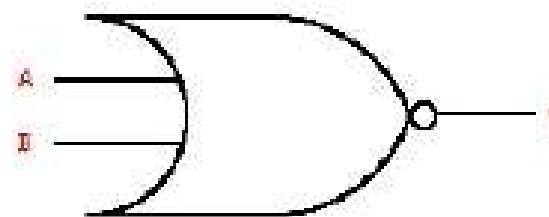
Dynamic Power Consumption [nW/MHz]

Input Transition [ns]		0.01		4.00	
Load Capacitance [fF]		30.00	600.00	30.00	600.00
A to Q	fall	5.73	8.27	3025.58	2459.24
	rise	155.47	156.48	3806.99	3246.79
B to Q	fall	5.87	8.76	3407.23	2766.61
	rise	196.60	199.91	4408.72	3647.73

NAND2X6

am
www.ams.com

Strength	1
Cell Area	43.680 μm^2
Equation	$Q = \text{!}(A \mid B)$
Type	Combinational
Input	A, B
Output	Q



State Table		
A	B	Q
L	L	H
H	-	L
-	H	L

Propagation Delay [ns]

Input Transition [ns]		0.01		4.00	
Load Capacitance [fF]		5.00	100.00	5.00	100.00
A to Q	fall	0.04	0.41	-0.64	0.48
	rise	0.11	1.17	1.28	2.61
B to Q	fall	0.05	0.42	-0.45	0.56
	rise	0.11	1.18	1.00	2.25

Output Transition [ns]

Input Transition [ns]		0.01		4.00	
Load Capacitance [fF]		5.00	100.00	5.00	100.00
A to Q	fall	0.04	0.55	0.73	1.46
	rise	0.15	1.81	0.62	2.12
B to Q	fall	0.05	0.55	0.94	1.55
	rise	0.15	1.81	0.80	2.12

Capacitance [fF]	
A	2.6610
B	2.9400

Leakage [pW]	
0.26	

Dynamic Power Consumption [nW/MHz]

Input Transition [ns]		0.01		4.00	
Load Capacitance [fF]		5.00	100.00	5.00	100.00
A to Q	fall	4.63	5.12	326.27	363.32
	rise	45.80	47.19	661.91	550.88
B to Q	fall	7.48	7.98	453.09	450.46
	rise	55.57	56.53	796.35	658.80

NOR2X1

am
www.ams.com

Dynamic Power Consumption is Data Dependent

- Switching activity, $P_{0 \rightarrow 1}$, has two components
 - A static component – function of the logic topology ~~slow Q0E0~~
 - A dynamic component – function of the timing behavior (glitching) ~~diffuse us~~

2-input NOR Gate

A	B	Out
0	0	1
0	1	0
1	0	0
1	1	0

Static transition probability

$$P_{0 \rightarrow 1} = P_{\text{out}=0} \times P_{\text{out}=1} \quad \frac{\text{Probabilistic}}{\text{of switching}}$$
$$= P_0 \times (1-P_0)$$

With input signal probabilities

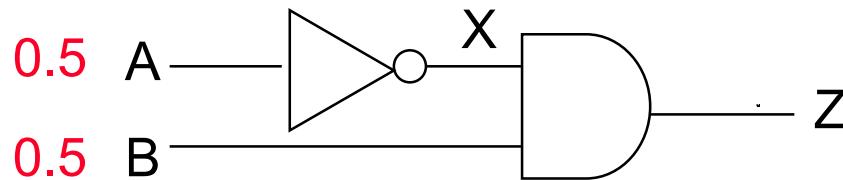
$$P_{A=1} = 1/2$$
$$P_{B=1} = 1/2$$

NOR static transition probability

$$= 3/4 \times 1/4 = \underline{\underline{3/16}}$$

Transition Probabilities for Some Basic Gates

	$P_{0 \rightarrow 1} = P_{\text{out}=0} \times P_{\text{out}=1}$
NOR	$(1 - (1 - P_A)(1 - P_B)) \times (1 - P_A)(1 - P_B)$
OR	$(1 - P_A)(1 - P_B) \times (1 - (1 - P_A)(1 - P_B))$
NAND	$P_A P_B \times (1 - P_A P_B)$
AND	$(1 - P_A P_B) \times P_A P_B$
XOR	$(1 - (P_A + P_B - 2P_A P_B)) \times (P_A + P_B - 2P_A P_B)$



$$\begin{aligned}\text{For } X: P_{0 \rightarrow 1} &= P_0 \times P_1 = (1 - P_A) P_A \\ &= 0.5 \times 0.5 = 0.25\end{aligned}$$

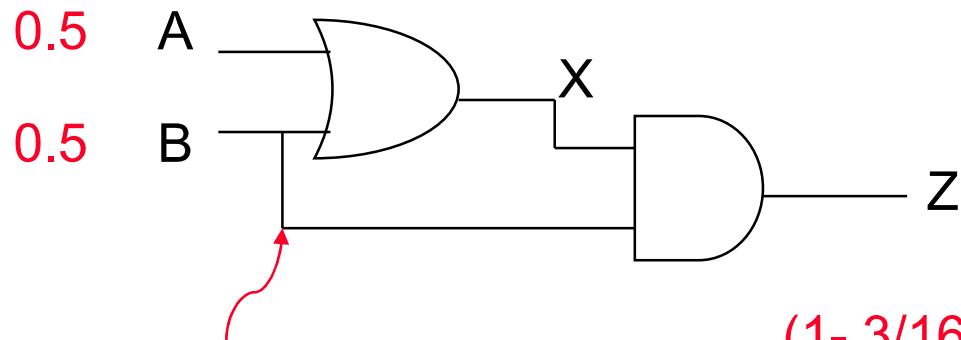
$$\begin{aligned}\text{For } Z: P_{0 \rightarrow 1} &= P_0 \times P_1 = (1 - P_X P_B) P_X P_B \\ &= (1 - (0.5 \times 0.5)) \times (0.5 \times 0.5) = 3/16\end{aligned}$$

ZATOPOWOCIA DLA WFWENTY! QUESTO RUMBO!

Inter-signal Correlations

- Determining switching activity is complicated by the fact that signals exhibit correlation in space and time
 - reconvergent fan-out

$$(1-0.5)(1-0.5) \times (1-(1-0.5)(1-0.5)) = 3/16$$



Reconvergent

$$(1 - 3/16 \times 0.5) \times (3/16 \times 0.5) = 0.085$$

$$P(Z=1) = P(B=1) \& P(A=1 | B=1)$$

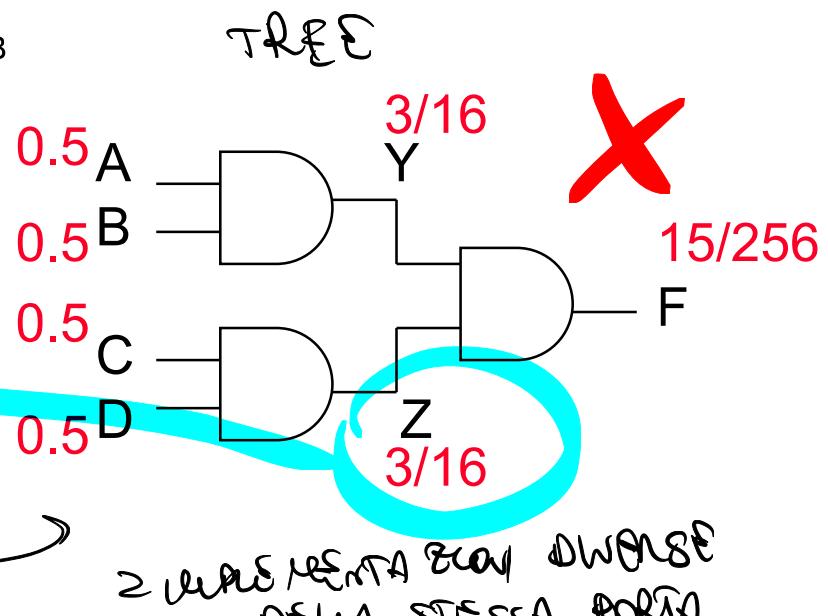
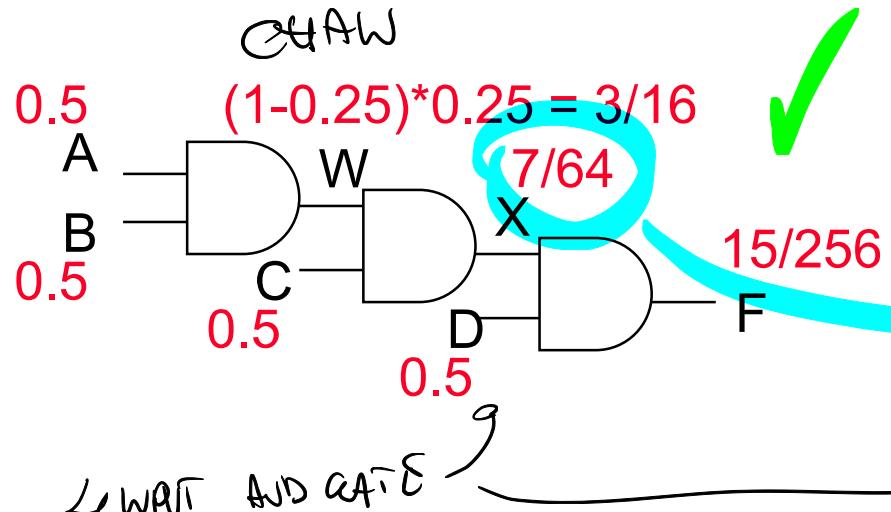


- Have to use conditional probabilities

Logic Restructuring

- Logic restructuring: changing the topology of a logic network to reduce transitions

$$\text{AND: } P_{0 \rightarrow 1} = P_0 \times P_1 = (1 - P_A P_B) \times P_A P_B$$

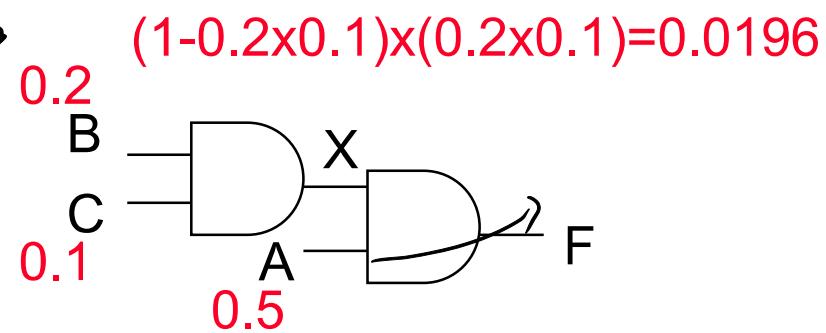
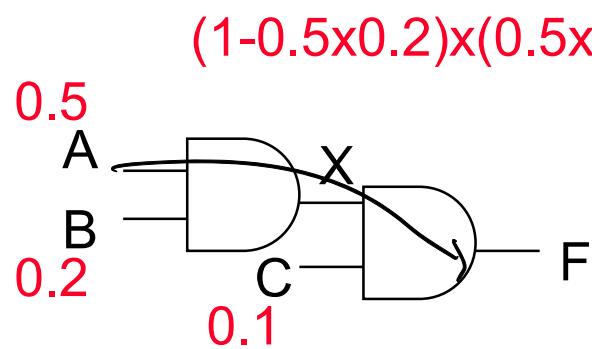


Chain implementation has a lower overall switching activity than the tree implementation for random inputs

Ignores glitching effects

Input Ordering

Pode ser feito automaticamente por ferramentas.

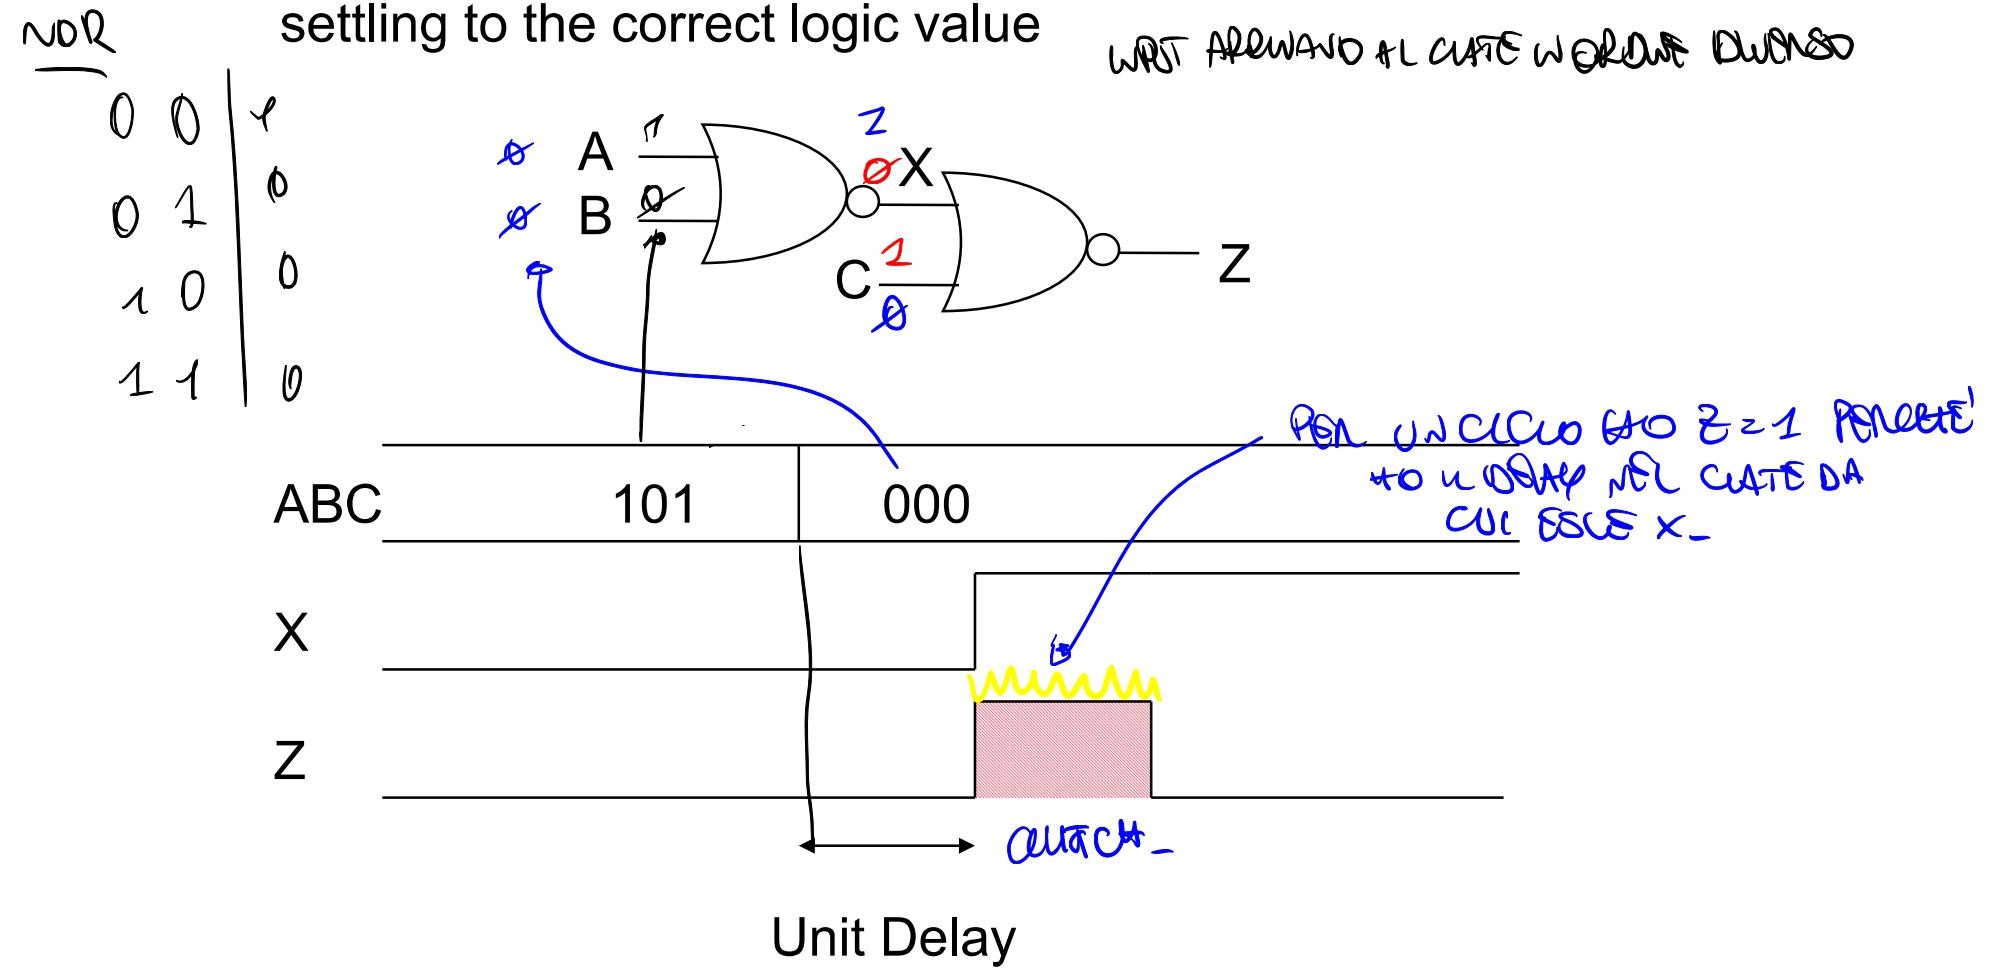


A INPUTA SÓ 2 GATES

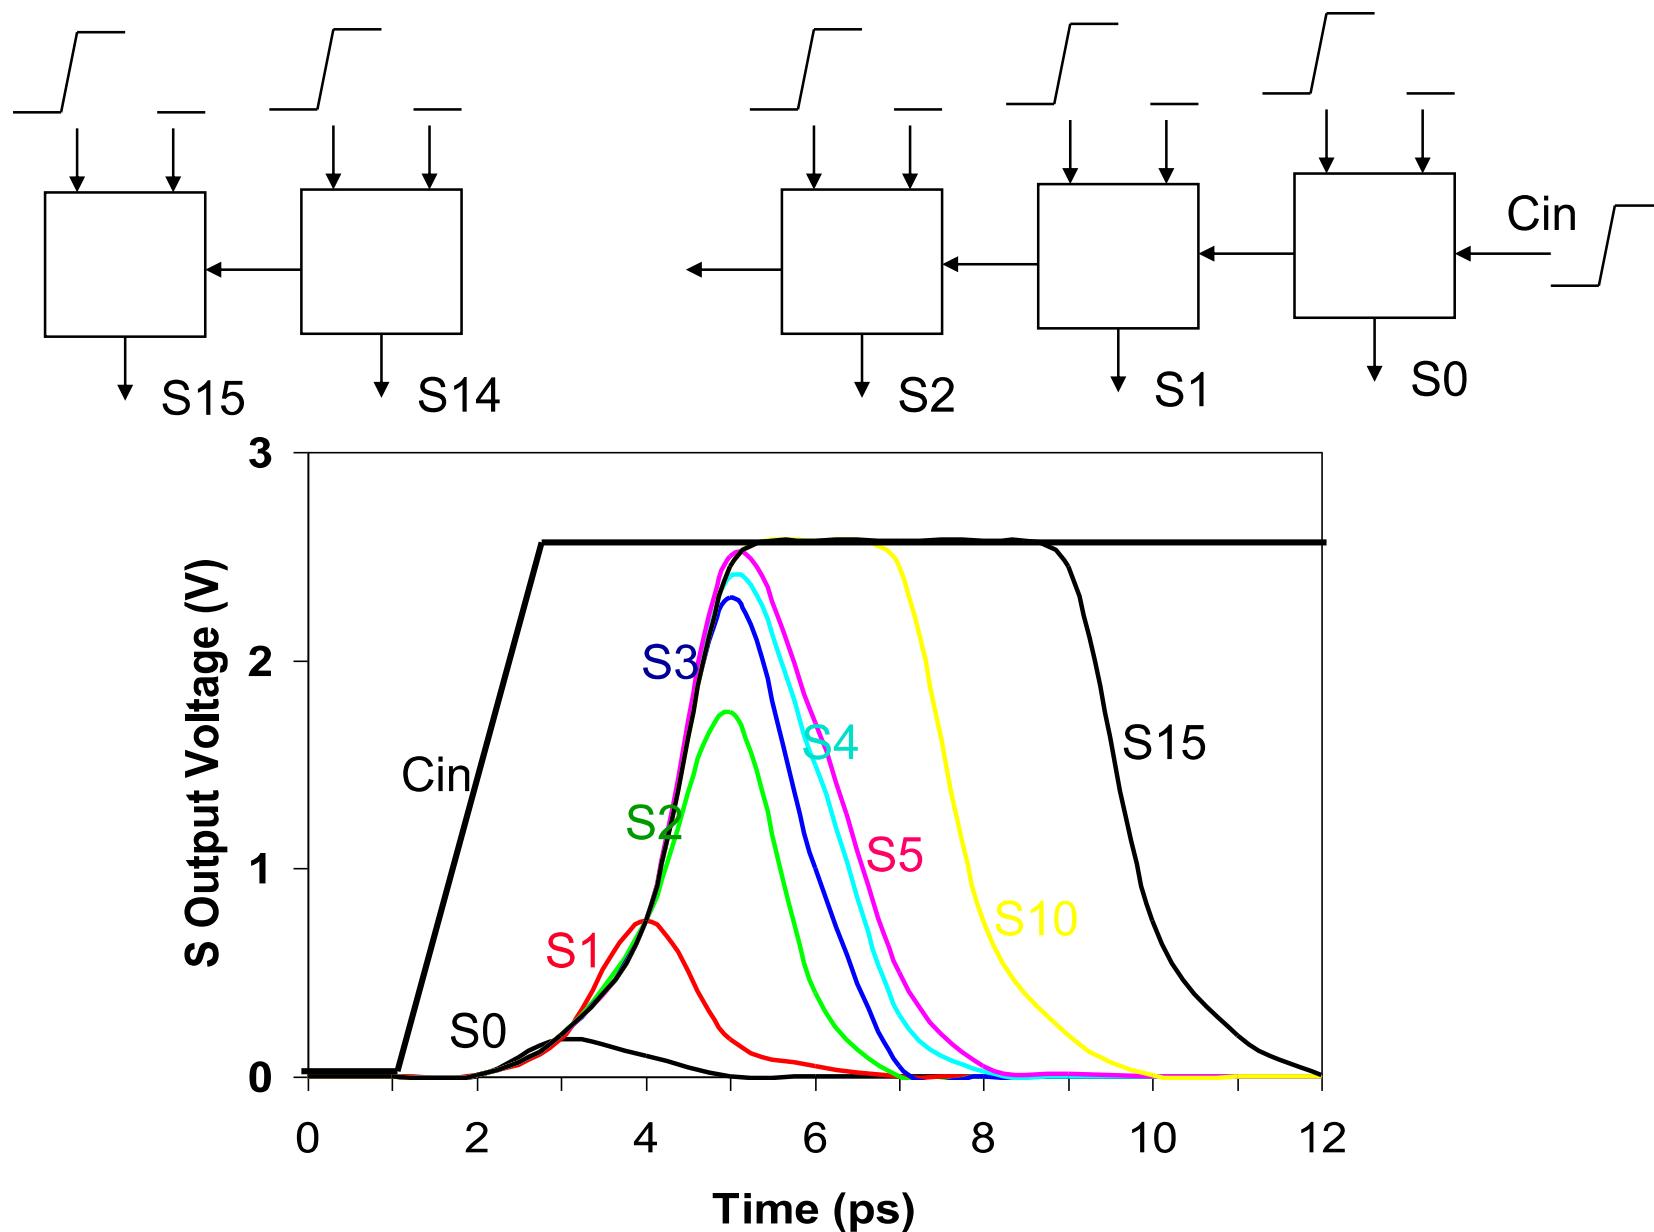
Beneficial to postpone the introduction of signals with a **high** transition rate (signals with signal probability close to 0.5)

Glitching in Static CMOS Networks

- Gates have a nonzero propagation delay resulting in spurious transitions or **glitches** (dynamic hazards)
 - glitch: node exhibits multiple transitions in a single cycle before settling to the correct logic value

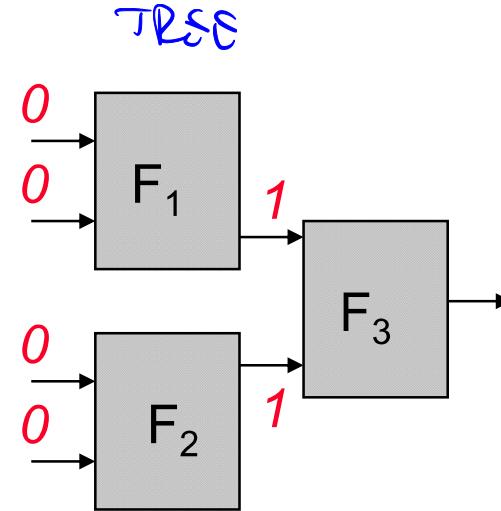
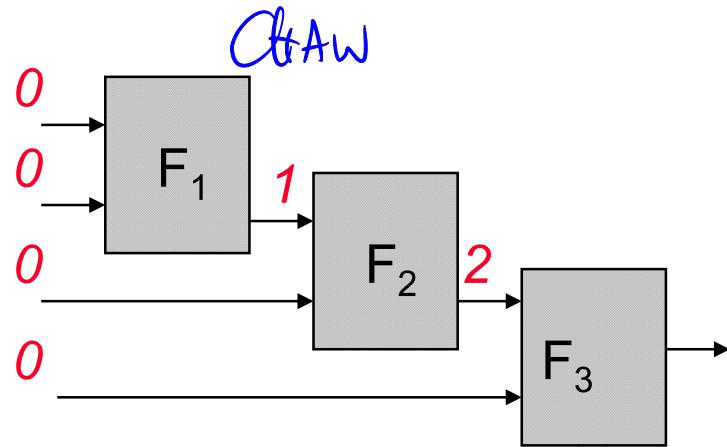


Glitching in an RCA



Balanced Delay Paths to Reduce Glitching

- Glitching is due to a mismatch in the path lengths in the logic network; if all input signals of a gate change simultaneously, no glitching occurs



So equalize the lengths of timing paths through logic

Power and Energy Design Space

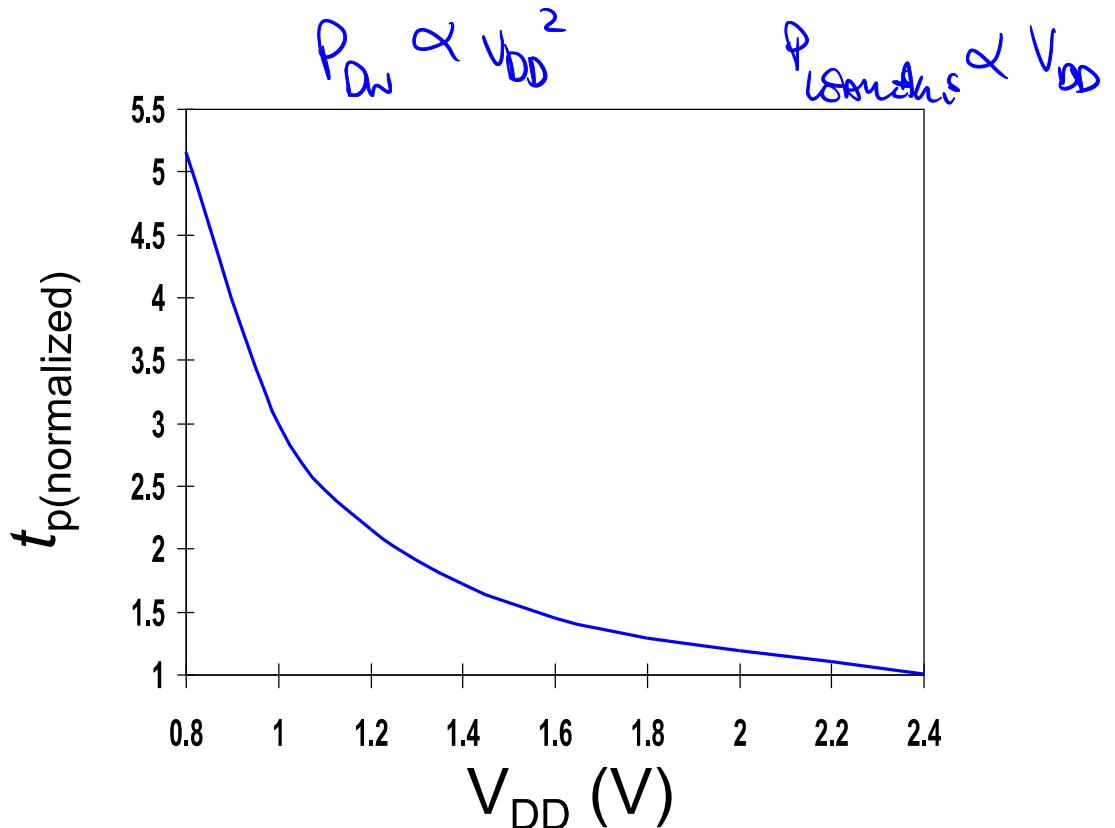
	Constant Throughput/Latency	Variable Throughput/Latency	
Energy	Design Time	Non-active Modules	Run Time
Active	Logic Design Reduced V_{dd} Sizing Multi-V_{dd}	Clock Gating	DFS, DVS (Dynamic Freq, Voltage Scaling)
Leakage	+ Multi- V_T	Sleep Transistors Multi- V_{dd} Variable V_T	+ Variable V_T

Dynamic Power as a Function of V_{DD}

- Decreasing the V_{DD} **decreases** dynamic energy consumption (quadratically)
- But, **increases** gate delay (decreases performance)

$$t_{p_{HL}} \propto \frac{K \cdot C}{\beta_n \cdot (V_{DD} - V_{Tn})}$$

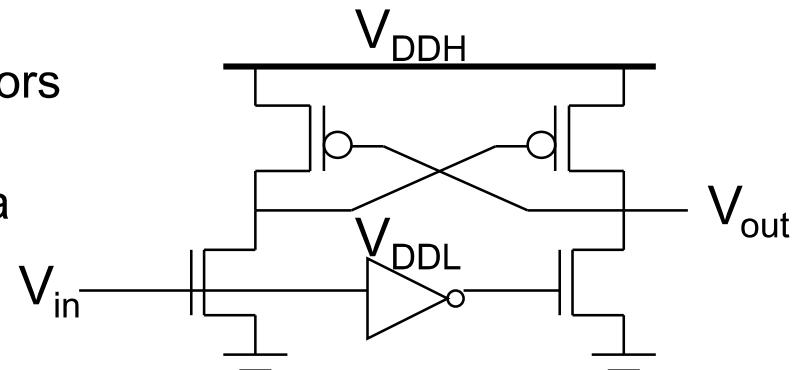
$$t_{p_{LH}} \propto \frac{K \cdot C}{\beta_p \cdot (V_{DD} + V_{Tp})}$$



- Determine the critical path(s) at **design time** and use high V_{DD} for the transistors on those paths for speed. Use a lower V_{DD} on the other gates, especially those that drive large capacitances (as this yields the largest energy benefits).

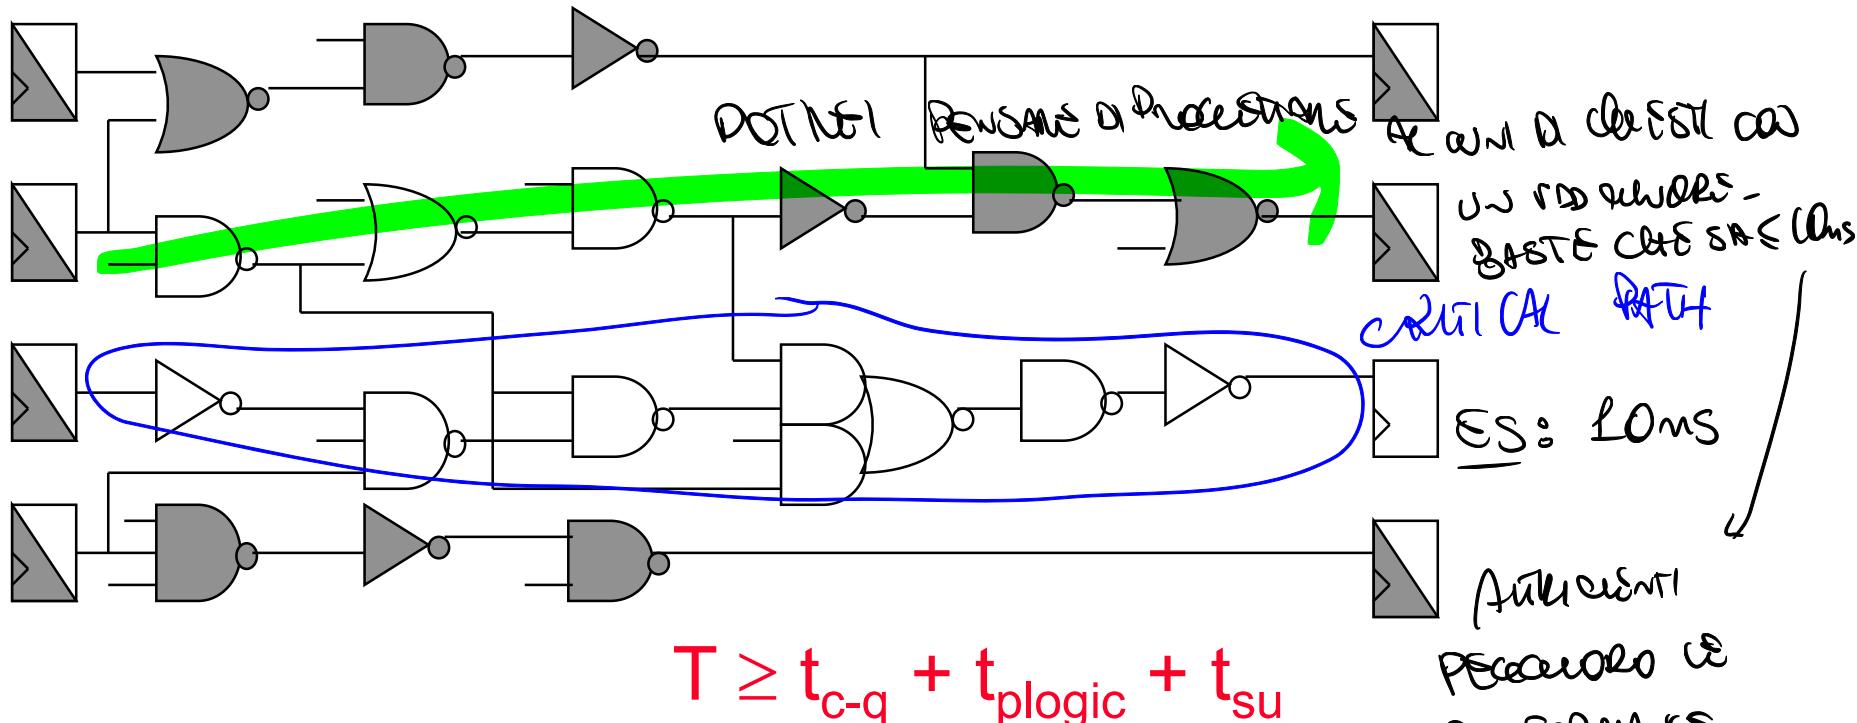
Multiple V_{DD} Considerations

- How many V_{DD} ? – Two is becoming common
 - Many chips already have two supplies (one for core and one for I/O)
- When combining multiple supplies, **level converters** are required whenever a module at the lower supply drives a gate at the higher supply (step-up)
 - If a gate supplied with V_{DDL} drives a gate at V_{DDH} , the PMOS never turns off
 - The cross-coupled PMOS transistors do the level conversion
 - The NMOS transistor operate on a reduced supply
 - Level converters are not needed for a step-down change in voltage
 - Overhead of level converters can be mitigated by doing conversions at register boundaries and embedding the level conversion inside the flipflop (see next slide)



Dual-Supply Inside a Logic Block

- Minimum energy consumption is achieved if **all** logic paths are critical (have the same delay)
- Clustered voltage-scaling
 - Each path starts with V_{DDH} and switches to V_{DDL} (gray logic gates) when delay **slack** is available
 - Level conversion is done in the flipflops at the end of the paths

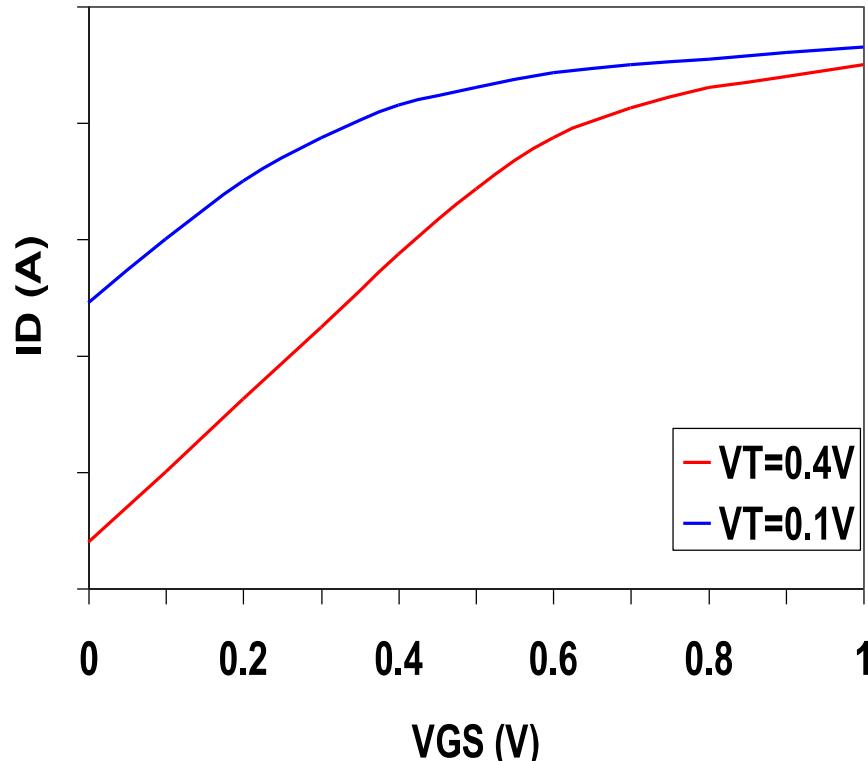


Power and Energy Design Space

	Constant Throughput/Latency	Variable Throughput/Latency	
Energy	Design Time	Non-active Modules	Run Time
Active	Logic Design Reduced V_{dd} Sizing Multi- V_{dd}	Clock Gating	DFS, DVS (Dynamic Freq, Voltage Scaling)
Leakage	+ Multi- V_T	Sleep Transistors Multi- V_{dd} Variable V_T	+ Variable V_T

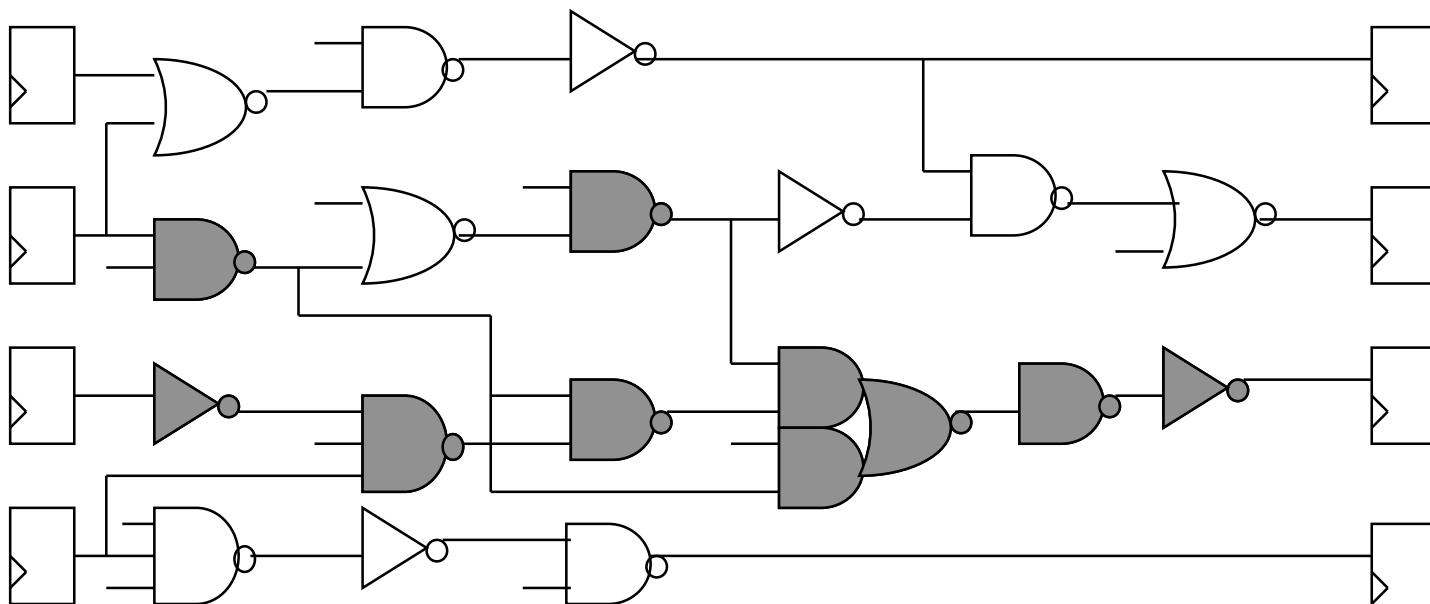
Leakage as a Function of Design Time V_T

- ❑ Reducing the V_T increases the sub-threshold leakage current (exponentially)
 - 90mV reduction in V_T increases leakage by an order of magnitude
- ❑ But, reducing V_T decreases gate delay (increases performance)
- ❑ Determine the critical path(s) at design time and use low V_T devices on the transistors on those paths for speed. Use a high V_T on the other logic for leakage control.
 - A careful assignment of V_T 's can reduce the leakage by as much as 80%



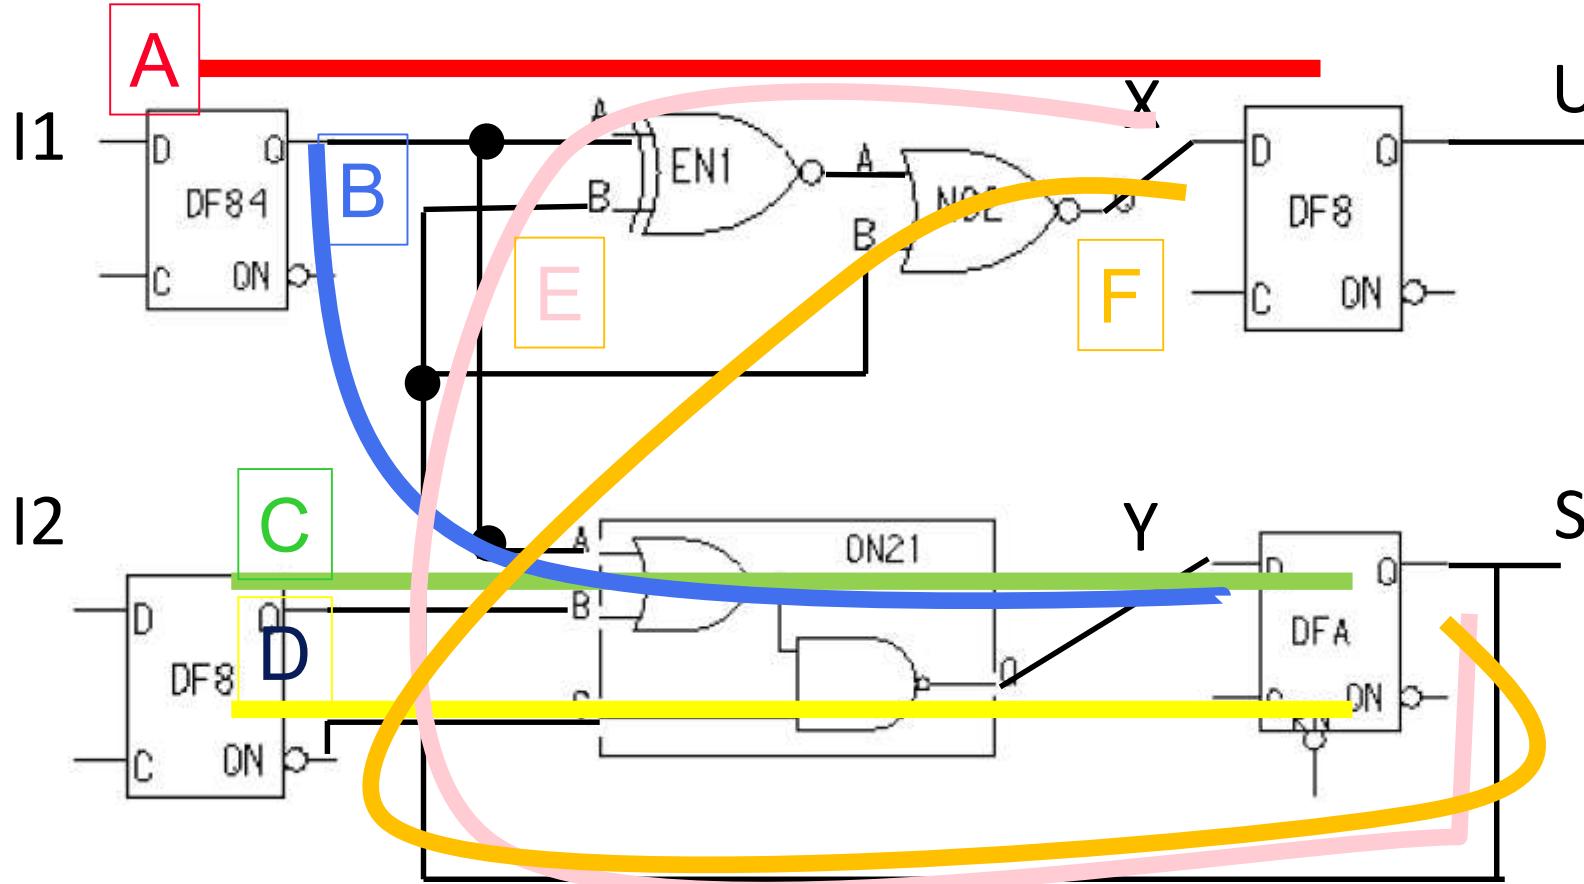
Dual-Thresholds Inside a Logic Block

- ❑ Minimum energy consumption is achieved if **all** logic paths are critical (have the same delay)
- ❑ Use lower threshold on timing-critical paths
 - Assignment can be done on a per gate or transistor basis; no clustering of the logic is needed
 - No level converters are needed



$$T \geq t_{c-q} + t_{p\text{logic}} + t_{su}$$

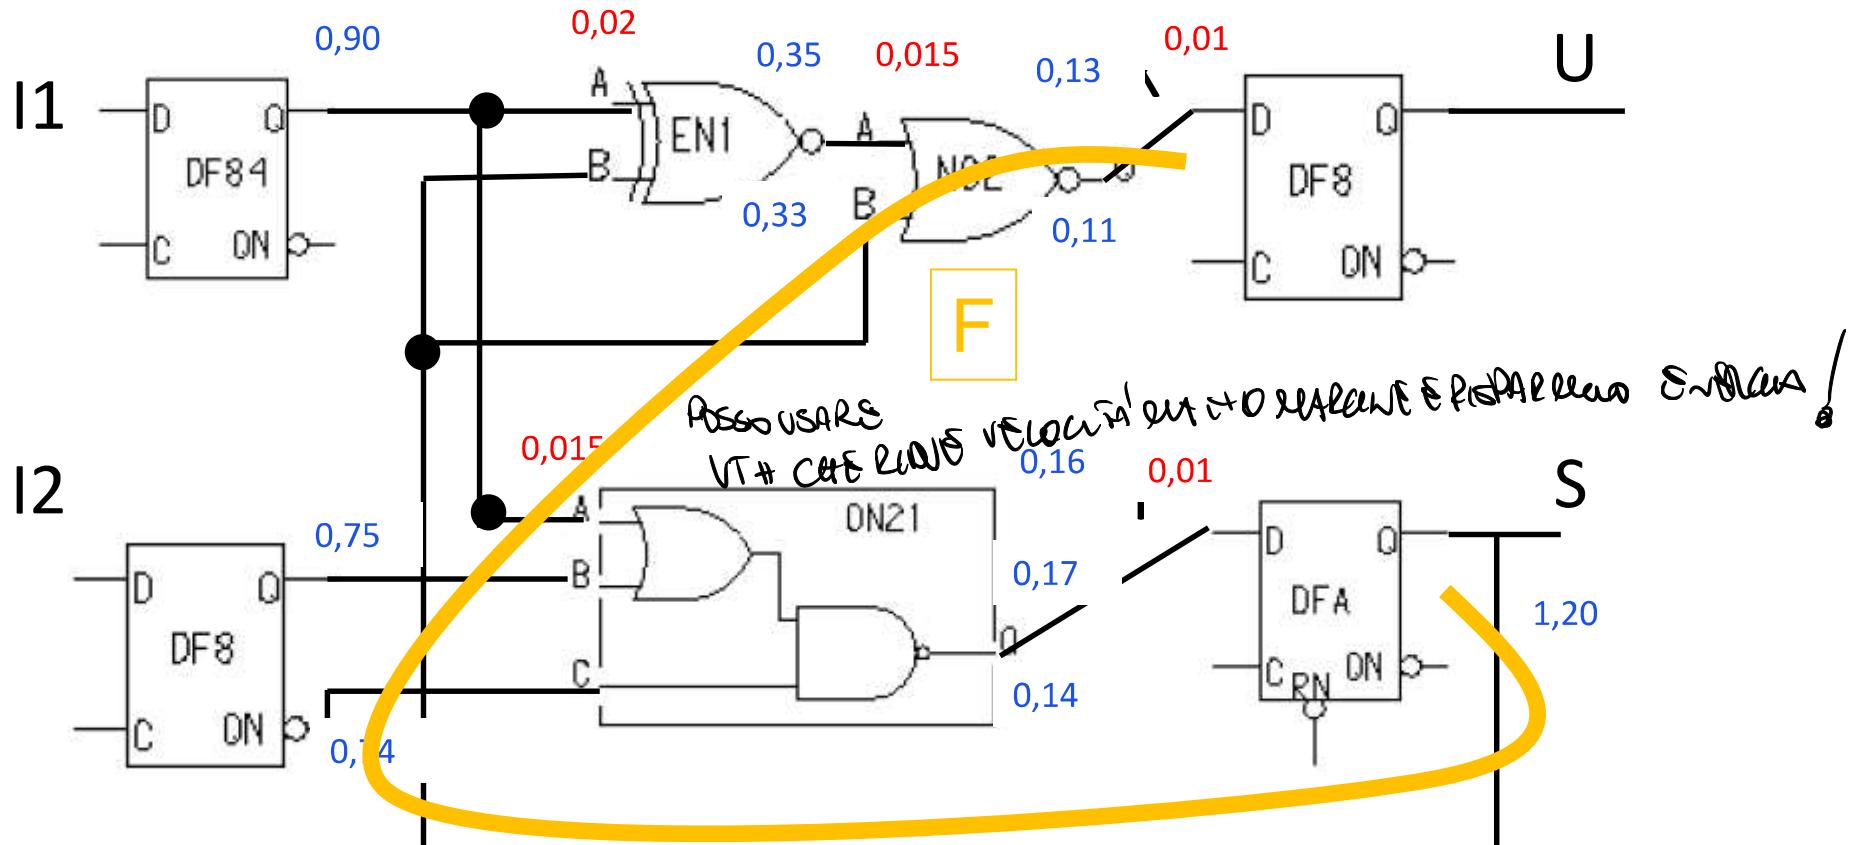
Example for evaluating minimum Clock Period



T (clock period)
↓
[Clock Pulse Diagram]

$$T \geq t_{c-q} + t_{p\text{logic}} + t_{su}$$

Example for evaluating minimum Clock Period



A	0,9	0,35	0,13	0,17	1,55
B	0,9	0,16		0,16	1,22
C	0,75	0,17		0,16	1,08
D	0,74	0,14		0,16	1,04
E	1,2	0,11		0,17	1,48
F	1,2	0,33	0,13	0,17	1,83

$1,83 - 1,22 = \Delta$
 (Δ Vario possa
 prececerre)
 546,45 MHz

T (clock period)



$$T \geq t_{c-q} + t_{\text{plogic}} + t_{\text{su}}$$

Low Power Techniques in Microarchitectures and Memories

Review: Energy & Power Equations

$$E = C_L V_{DD}^2 P_{0 \rightarrow 1} + t_{sc} V_{DD} I_{peak} P_{0 \rightarrow 1} + V_{DD} I_{leakage}$$

$$f_{0 \rightarrow 1} = P_{0 \rightarrow 1} * f_{clock}$$

$$P = C_L V_{DD}^2 f_{0 \rightarrow 1} + t_{sc} V_{DD} I_{peak} f_{0 \rightarrow 1} + V_{DD} I_{leakage}$$

Dynamic power
(% decreasing
relatively with
deep submicron)

Short-circuit
power

Leakage power
(% increasing
with deep
submicron)

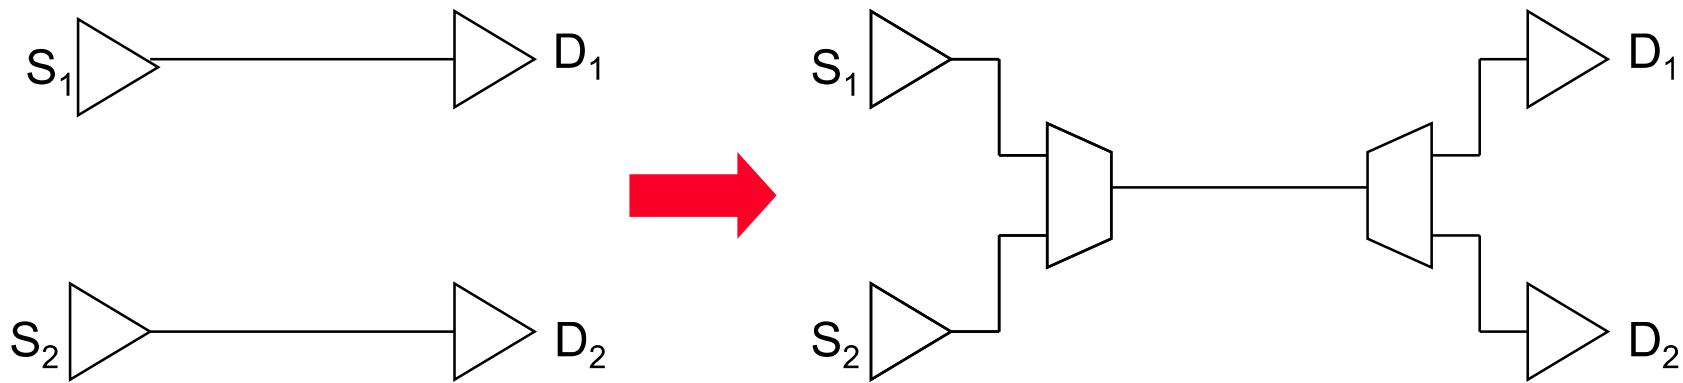
Leakage power grows from <5% of power budget at .25 micron to 20-25% at 130nm to 40-50% at 90nm and continuing to increase at 65nm and beyond.

Power and Energy Design Space

	Constant Throughput/Latency	Variable Throughput/Latency	
Energy	Design Time	Non-active Modules	Run Time
Active	Logic Design Reduced V_{dd} Sizing Multi- V_{dd}	Clock Gating	DFS, DVS (Dynamic Freq, Voltage Scaling)
Leakage	+ Multi- V_T	Sleep Transistors Multi- V_{dd} Variable V_T	+ Variable V_T

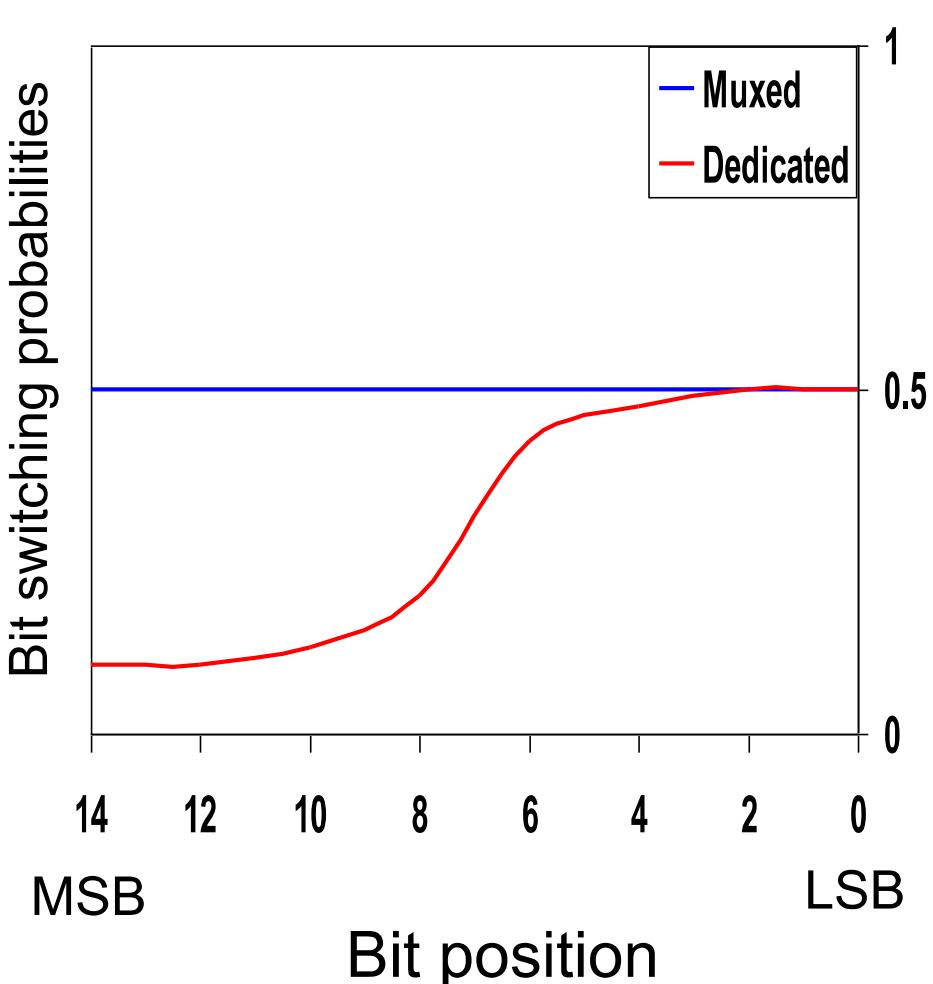
Bus Multiplexing

- Buses are a significant source of power dissipation due to high switching activities and large capacitive loading
 - 15% of total power in Alpha 21064
 - 30% of total power in Intel 80386
- Share long data buses with time multiplexing (S_1 uses even cycles, S_2 odd)



- But what if data samples are correlated (e.g., sign bits)?

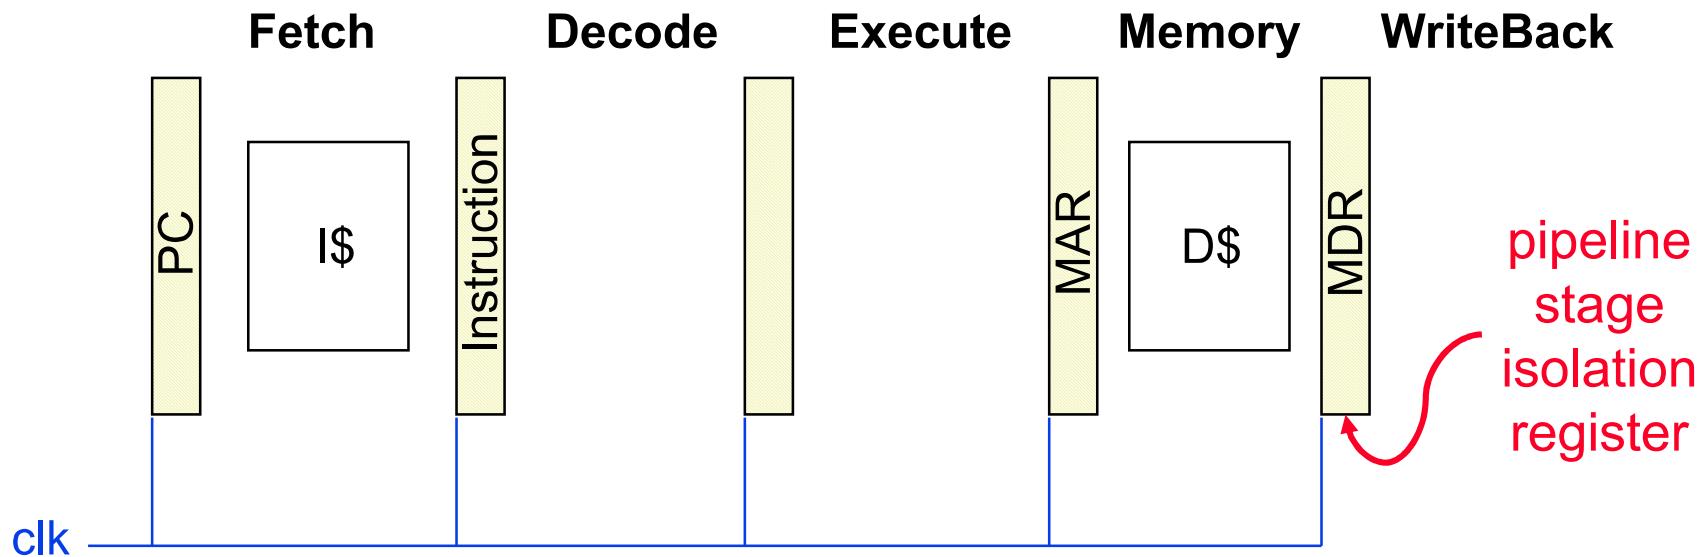
Correlated Data Streams



- ❑ For a shared (multiplexed) bus advantages of data correlation are lost (bus carries samples from two uncorrelated data streams)
 - Bus sharing should **not be used** for **positively** correlated data streams
 - Bus sharing **may prove advantageous** in a **negatively** correlated data stream (where successive samples switch sign bits) - more random switching

Glitch Reduction by Pipelining

- ❑ Glitches depend on the **logic depth** of the circuit - gates deeper in the logic network are more prone to glitching
 - arrival times of the gate inputs are more spread due to delay imbalances
 - usually affected more by primary input switching
- ❑ Reduce logic depth by adding pipeline registers
 - additional energy used by the clock and pipeline registers

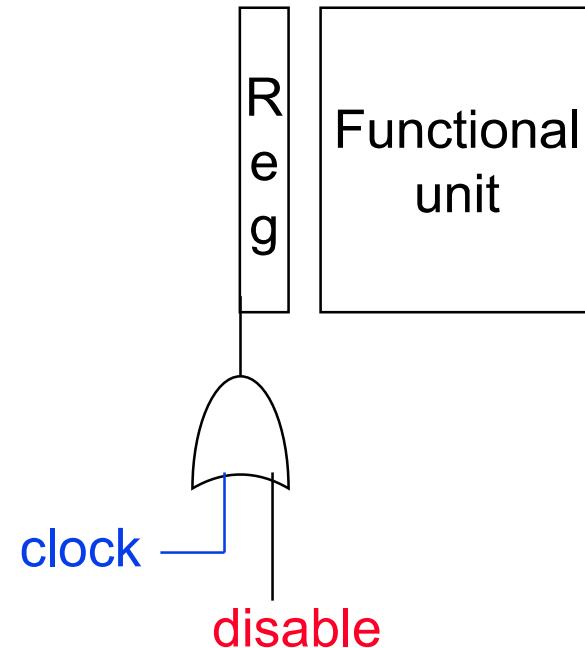


Power and Energy Design Space

	Constant Throughput/Latency	Variable Throughput/Latency	
Energy	Design Time	Non-active Modules	
Active	Logic Design Reduced V_{dd} Sizing Multi- V_{dd}	Clock Gating	DFS, DVS (Dynamic Freq, Voltage Scaling)
Leakage	+ Multi- V_T	Sleep Transistors Multi- V_{dd} Variable V_T	+ Variable V_T

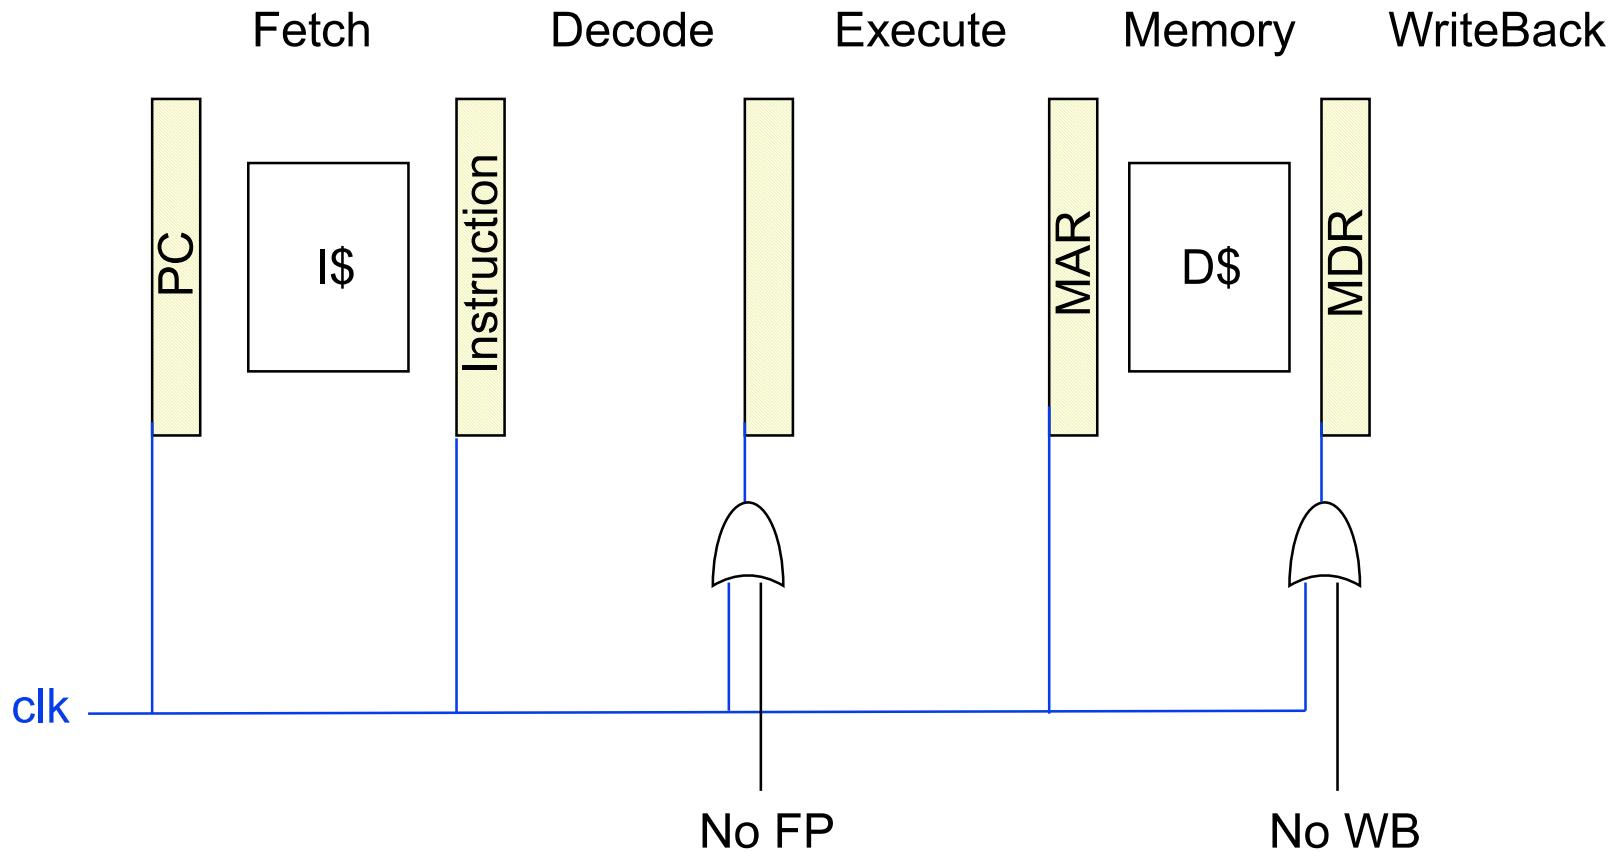
Clock Gating

- ❑ Most popular method for power reduction of clock signals and functional units
- ❑ Gate off clock to idle functional units
 - e.g., floating point units
 - need logic to generate **disable** signal
 - increases complexity of control logic
 - consumes power
 - timing critical to avoid clock glitches at OR gate output
 - additional gate delay on clock signal
 - gating OR gate can replace a buffer in the clock distribution tree



Clock Gating in a Pipelined Datapath

- For idle units (e.g., floating point units in Exec stage, WB stage for instructions with no write back operation)

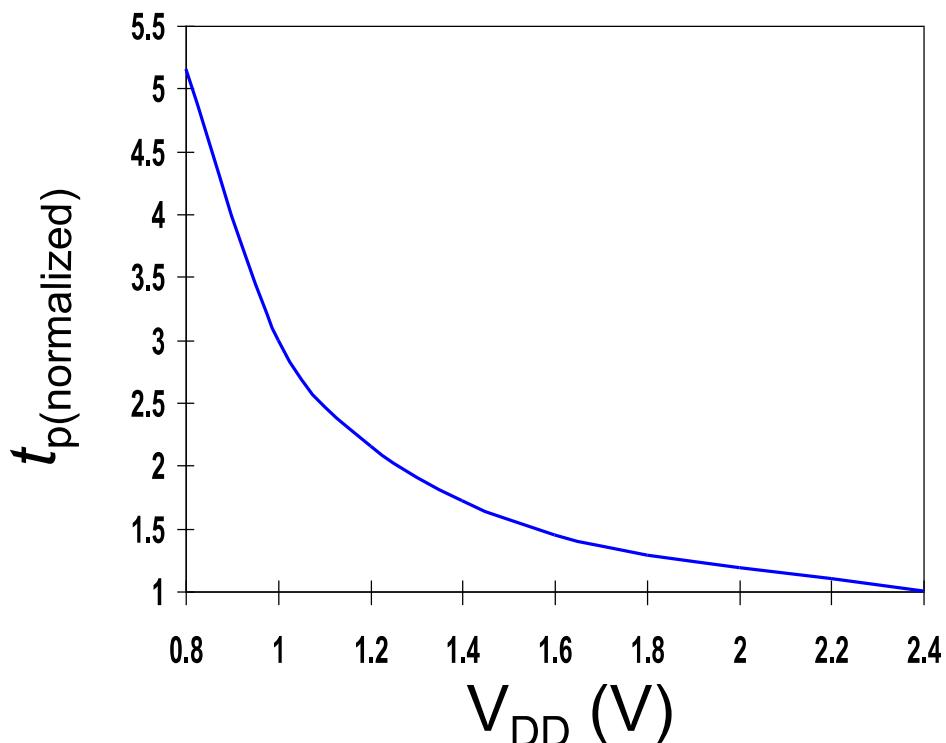


Power and Energy Design Space

	Constant Throughput/Latency	Variable Throughput/Latency	
Energy	Design Time	Non-active Modules	Run Time
Active	Logic Design Reduced V_{dd} Sizing Multi-V_{dd}	Clock Gating	DFS, DVS (Dynamic Freq, Voltage Scaling)
Leakage	+ Multi- V_T	Sleep Transistors Multi- V_{dd} Variable V_T	+ Variable V_T

Review: Dynamic Power as a Function of V_{DD}

- Decreasing the V_{DD} **decreases** dynamic energy consumption (quadratically)
- But, **increases** gate delay (decreases performance)



- Determine the critical path(s) at **design time** and use high V_{DD} for the transistors on those paths for speed. Use a lower V_{DD} on the other gates, especially those that drive large capacitances (as this yields the largest energy benefits).

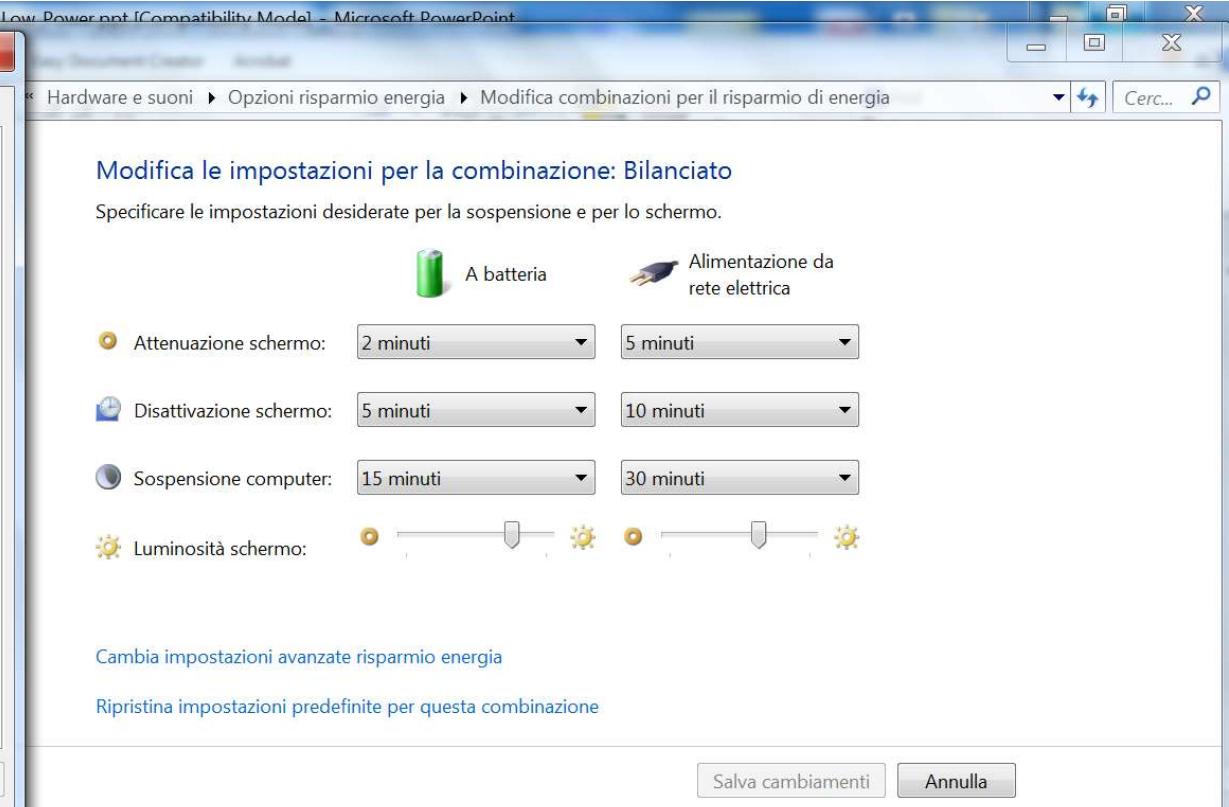
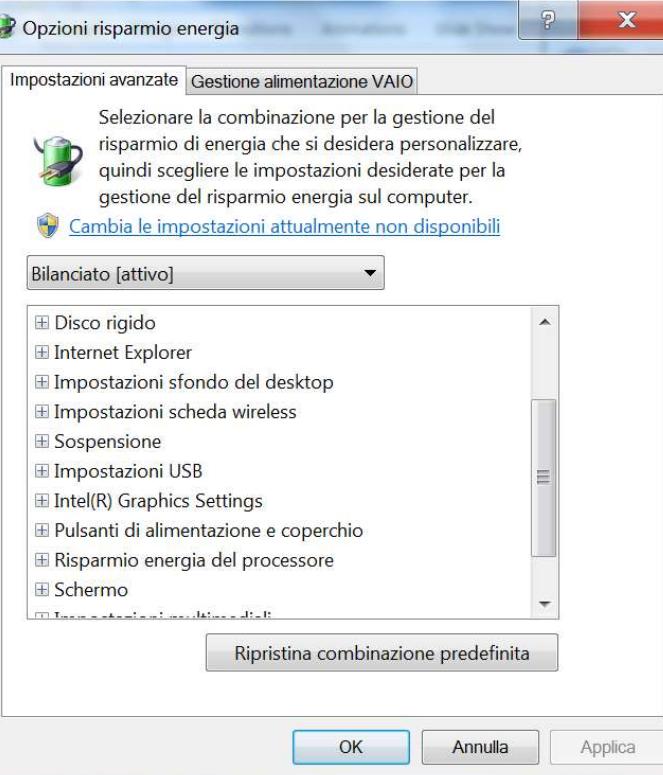
Dynamic Frequency and Voltage Scaling

❑ Intel's SpeedStep

- Hardware that steps down the clock frequency (dynamic frequency scaling – DFS) when the user unplugs from AC power
 - PLL from 650MHz → 500MHz
- CPU stalls during SpeedStep adjustment

❑ Transmeta LongRun

- Hardware that applies **both** DFS and DVS (dynamic supply voltage scaling)
 - 32 levels of V_{DD} from 1.1V to 1.6V
 - PLL from 200MHz → 700MHz in increments of 33MHz
- Triggered when CPU load change is detected by software
 - heavier load → ramp up V_{DD} , when stable speed up clock
 - lighter load → slow down clock, when PLL locks onto new rate, ramp down V_{DD}
- CPU stalls only during PLL relock (< 20 microsec)



Dynamic Thermal Management (DTM)



Trigger Mechanism:
When do we enable
DTM techniques?



Initiation Mechanism:
How do we enable
technique?



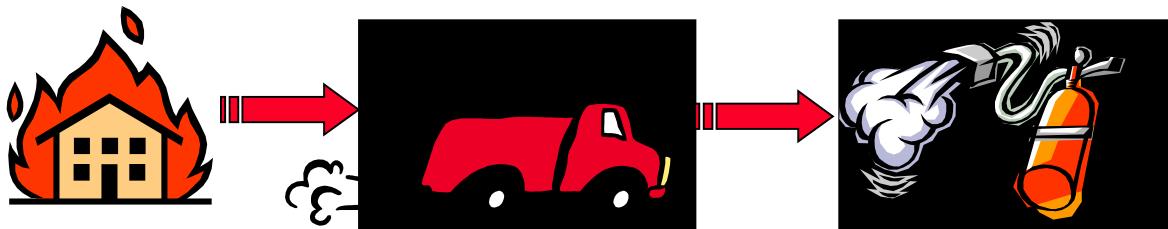
Response Mechanism:
What technique do we
enable?

DTM Trigger Mechanisms



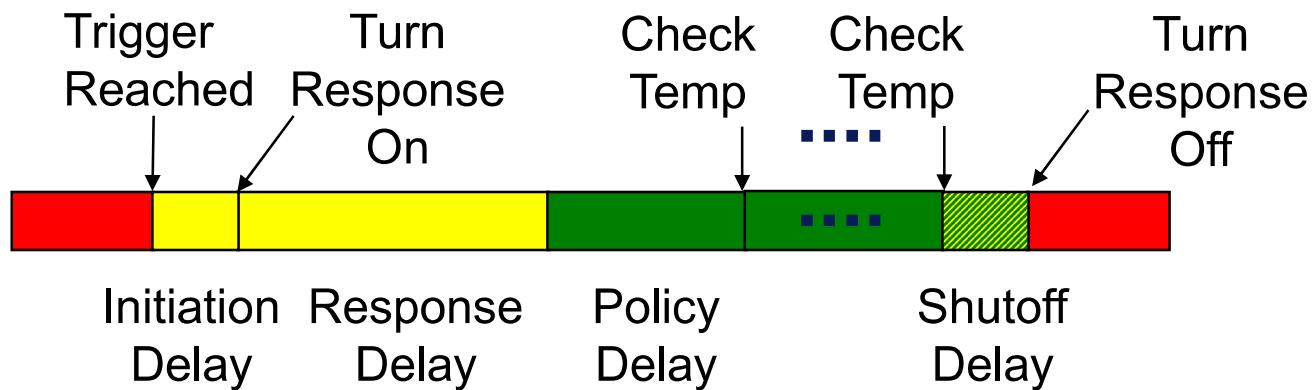
- ❑ Mechanism: How to deduce temperature?
- ❑ Direct approach: on-chip temperature sensors
 - Based on differential voltage change across 2 diodes of different sizes
 - May require >1 sensor
 - Hysteresis and delay are problems
- ❑ Policy: When to begin responding?
 - Trigger level set too high means higher packaging costs
 - Trigger level set too low means frequent triggering and loss in performance
- ❑ Choose trigger level to exploit difference between average and worst case power

DTM Initiation and Response Mechanisms



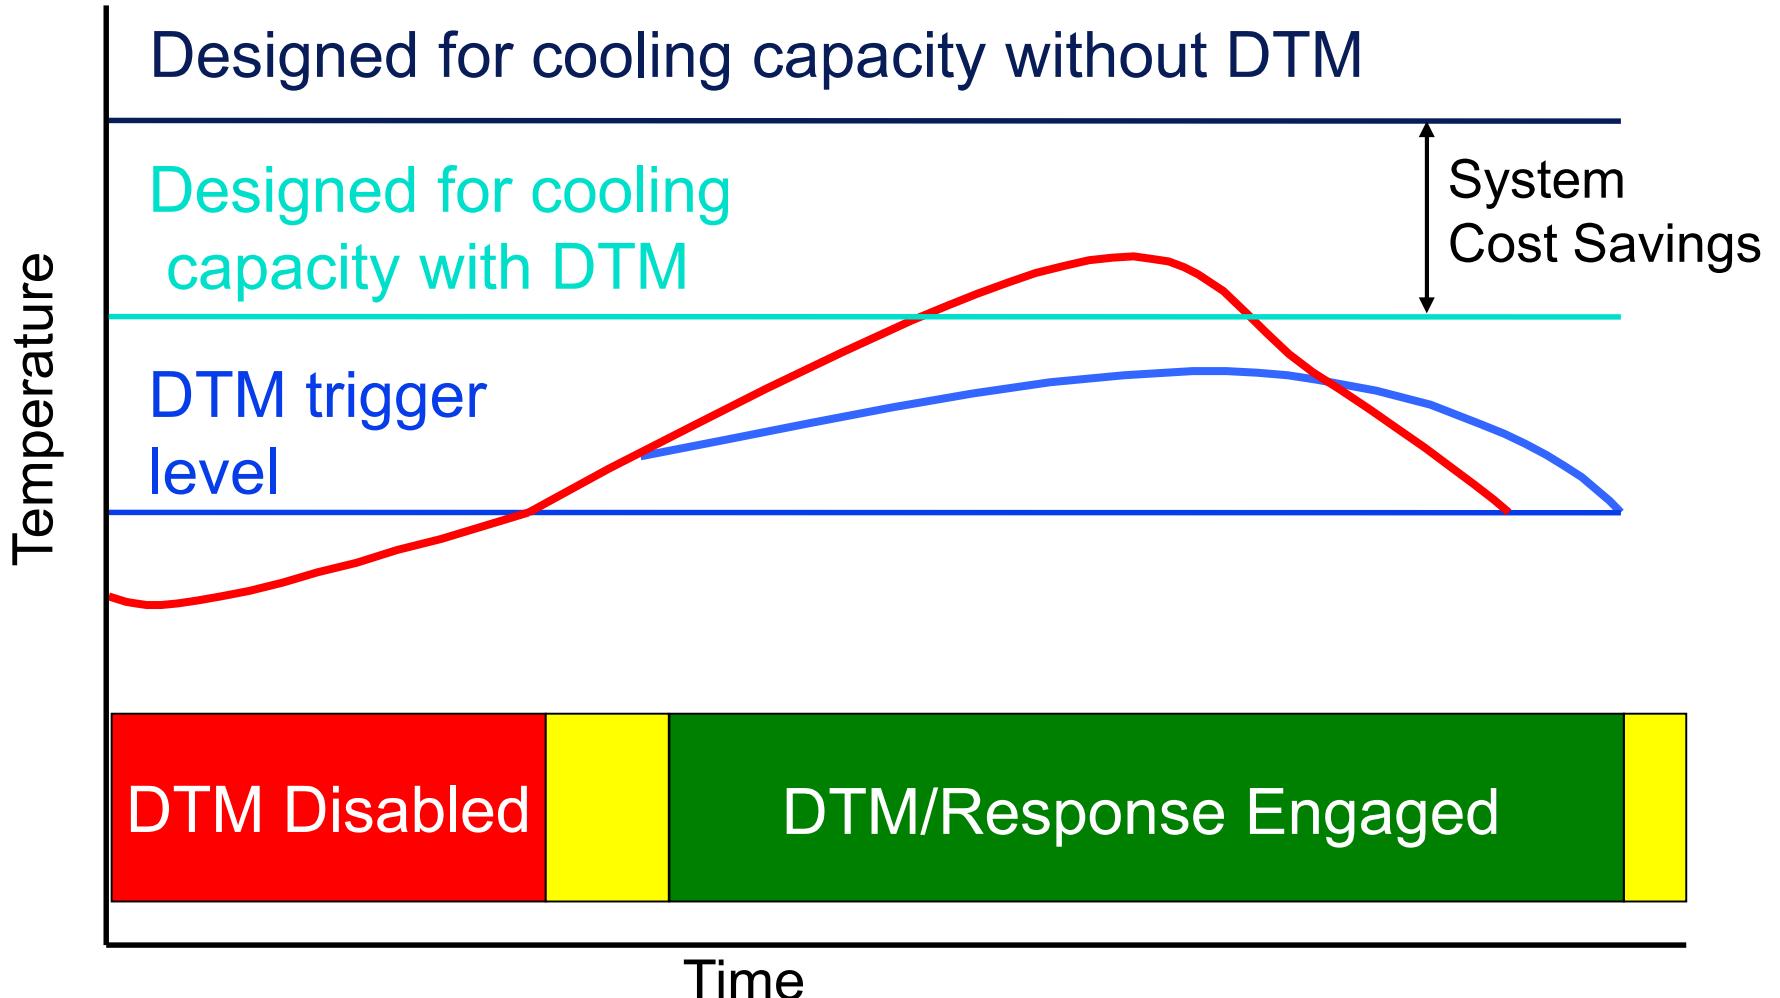
- ❑ Operating system or microarchitectural control?
 - Hardware support can reduce performance penalty by 20-30%
- ❑ Initiation of policy incurs some delay
 - When using DVS and/or DFS, much of the performance penalty can be attributed to enabling/disabling overhead
 - Increasing policy delay reduces overhead; smarter initiation techniques would help as well
- ❑ Thermal window (100Kcycles+)
 - Larger thermal windows “smooth” short thermal spikes

DTM Activation and Deactivation Cycle



- Initiation Delay – OS interrupt/handler
- Response Delay – Invocation time (e.g., adjust clock)
- Policy Delay – Number of cycles engaged
- Shutoff Delay – Disabling time (e.g., re-adjust clock)

DTM Savings Benefits



Power and Energy Design Space

	Constant Throughput/Latency	Variable Throughput/Latency	
Energy	Design Time	Non-active Modules	Run Time
Active	Logic Design Reduced V_{dd} Sizing Multi- V_{dd}	Clock Gating	DFS, DVS (Dynamic Freq, Voltage Scaling)
Leakage	+ Multi- V_T	Sleep Transistors Multi- V_{dd} Variable V_T	+ Variable V_T